

Bibliography

- Abadi, M., A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng. 2015. [TensorFlow: Large-scale machine learning on heterogeneous systems](#). Software available from tensorflow.org.
- Abdallah, N. M. M., H. O. H. Abdalla, M. E. M. Idris, S. A. A. Eldouma, E. A. M. Taha, I. M. H. Rushwan, M. A. M. Elsafi, and M. A. E.-s. Ibrahim. 2025. [Improving crisis communication through translation: A study of u.s. emergency response and disaster management contexts](#). *Journal of Cultural Analysis and Social Change*, 10(3):1132–1146.
- Abney, S. P., R. E. Schapire, and Y. Singer. 1999. [Boosting applied to tagging and PP attachment](#). *EMNLP/NLC*.
- Acevedo, S., A. Laio, and M. Baroni. 2026. [Differential syntactic and semantic encoding in LLMs](#). *ICML*.
- Agarwal, O., S. Subramanian, A. Nenkova, and D. Roth. 2019. [Evaluation of named entity coreference](#). *Workshop on Computational Models of Reference, Anaphora and Coreference*.
- Aggarwal, C. C. and C. Zhai. 2012. A survey of text classification algorithms. In C. C. Aggarwal and C. Zhai, eds, *Mining text data*, 163–222. Springer.
- Agichtein, E. and L. Gravano. 2000. Snowball: Extracting relations from large plain-text collections. *Proceedings of the 5th ACM International Conference on Digital Libraries*.
- Agirre, E., C. Banea, C. Cardie, D. Cer, M. Diab, A. Gonzalez-Agirre, W. Guo, I. Lopez-Gazpio, M. Maritxalar, R. Mihalcea, G. Rigau, L. Uria, and J. Wiebe. 2015. [SemEval-2015 task 2: Semantic textual similarity, English, Spanish and pilot on interpretability](#). *SemEval-15*.
- Agirre, E., M. Diab, D. Cer, and A. Gonzalez-Agirre. 2012. [SemEval-2012 task 6: A pilot on semantic textual similarity](#). *SemEval-12*.
- Agirre, E. and D. Martinez. 2001. [Learning class-to-class selectional preferences](#). *CoNLL*.
- Ahia, O., S. Kumar, H. Gonen, J. Kasai, D. Mortensen, N. Smith, and Y. Tsvetkov. 2023. [Do all languages cost the same? Tokenization in the era of commercial language models](#). *EMNLP*.
- Aho, A. V. and J. D. Ullman. 1972. *The Theory of Parsing, Translation, and Compiling*, volume 1. Prentice Hall.
- Algoet, P. H. and T. M. Cover. 1988. [A sandwich proof of the Shannon-McMillan-Breiman theorem](#). *The Annals of Probability*, 16(2):899–909.
- Allen, J. 1984. [Towards a general theory of action and time](#). *Artificial Intelligence*, 23(2):123–154.
- Allen, J. and M. Core. 1997. Draft of DAMSL: Dialog act markup in several layers. Unpublished manuscript.
- Allen, J., M. S. Hunnicut, and D. H. Klatt. 1987. *From Text to Speech: The MITalk system*. Cambridge University Press.
- Allwood, J. 1995. [An activity-based approach to pragmatics](#). *Gothenburg Papers in Theoretical Linguistics*, 76:1–38.
- Allwood, J., J. Nivre, and E. Ahlsén. 1992. On the semantics and pragmatics of linguistic feedback. *Journal of Semantics*, 9:1–26.
- Althoff, T., C. Danescu-Niculescu-Mizil, and D. Jurafsky. 2014. How to ask for a favor: A case study on the success of altruistic requests. *ICWSM 2014*.
- An, J., H. Kwak, and Y.-Y. Ahn. 2018. [SemAxis: A lightweight framework to characterize domain-specific word semantics beyond sentiment](#). *ACL*.
- Anastasopoulos, A. and G. Neubig. 2020. [Should all cross-lingual embeddings speak English?](#) *ACL*.
- Antoniak, M. and D. Mimno. 2018. [Evaluating the stability of embedding-based word similarities](#). *TACL*, 6:107–119.
- Aone, C. and S. W. Bennett. 1995. [Evaluating automated and manual acquisition of anaphora resolution strategies](#). *ACL*.
- Ardila, R., M. Branson, K. Davis, M. Kohler, J. Meyer, M. Henretty, R. Morais, L. Saunders, F. Tyers, and G. Weber. 2020. [Common voice: A massively-multilingual speech corpus](#). *LREC*.
- Ariel, M. 2001. Accessibility theory: An overview. In T. Sanders, J. Schilperoord, and W. Spooren, eds, *Text Representation: Linguistic and Psycholinguistic Aspects*, 29–87. Benjamins.
- Arkadiev, P. M. 2020. *Morphology in typology: Historical retrospect, state of the art, and prospects*. Oxford.
- Artetxe, M. and H. Schwenk. 2019. [Massively multilingual sentence embeddings for zero-shot cross-lingual transfer and beyond](#). *TACL*, 7:597–610.
- Asher, N. 1993. *Reference to Abstract Objects in Discourse*. Studies in Linguistics and Philosophy (SLAP) 50, Kluwer.
- Asher, N. and A. Lascarides. 2003. *Logics of Conversation*. Cambridge University Press.
- Atal, B. S. and S. Hanauer. 1971. Speech analysis and synthesis by prediction of the speech wave. *JASA*, 50:637–655.
- Austin, J. L. 1962. *How to Do Things with Words*. Harvard University Press.
- Ba, J. L., J. R. Kiros, and G. E. Hinton. 2016. [Layer normalization](#). *NeurIPS workshop*.
- Baayen, R. H. 2001. *Word frequency distributions*. Springer.
- Baayen, R. H., R. Piepenbrock, and L. Gulikers. 1995. *The CELEX Lexical Database (Release 2) [CD-ROM]*. Linguistic Data Consortium, University of Pennsylvania [Distributor].
- Babbage, C. 1864. *Passages from the Life of a Philosopher*. Longman.
- Baccianella, S., A. Esuli, and F. Sebastiani. 2010. [Sentiwordnet 3.0: An enhanced lexical resource for sentiment analysis and opinion mining](#). *LREC*.
- Bach, K. and R. Harnish. 1979. *Linguistic communication and speech acts*. MIT Press.
- Backus, J. W. 1959. The syntax and semantics of the proposed international algebraic language of the Zurich ACM-GAMM Conference. *Information Processing: Proceedings of the International Conference on Information Processing, Paris*. UNESCO.
- Backus, J. W. 1996. Transcript of question and answer session. In R. L. Wexelblat, ed., *History of Programming Languages*, page 162. Academic Press.

- Bada, M., M. Eckert, D. Evans, K. Garcia, K. Shipley, D. Sitnikov, W. A. Baumgartner, K. B. Cohen, K. Verspoor, J. A. Blake, and L. E. Hunter. 2012. Concept annotation in the craft corpus. *BMC bioinformatics*, 13(1):161.
- Baevski, A., Y. Zhou, A. Mohamed, and M. Auli. 2020. wav2vec 2.0: A framework for self-supervised learning of speech representations. *NeurIPS*, volume 33.
- Bagga, A. and B. Baldwin. 1998. Algorithms for scoring coreference chains. *LREC Workshop on Linguistic Coreference*.
- Bahdanau, D., K. H. Cho, and Y. Bengio. 2015. Neural machine translation by jointly learning to align and translate. *ICLR 2015*.
- Bahdanau, D., J. Chorowski, D. Serdyuk, P. Brakel, and Y. Bengio. 2016. End-to-end attention-based large vocabulary speech recognition. *ICASSP*.
- Bahl, L. R. and R. L. Mercer. 1976. Part of speech assignment by a statistical decision algorithm. *Proceedings IEEE International Symposium on Information Theory*.
- Bahl, L. R., F. Jelinek, and R. L. Mercer. 1983. A maximum likelihood approach to continuous speech recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 5(2):179–190.
- Bai, Y., A. Jones, K. Ndousse, A. Askell, A. Chen, N. DasSarma, D. Drain, S. Fort, D. Ganguli, T. Henighan, N. Joseph, S. Kadavath, J. Kernion, T. Conerly, S. El-Showk, N. Elhage, Z. Hatfield-Dodds, D. Hernandez, T. Hume, S. Johnston, S. Kravec, L. Lovitt, N. Nanda, C. Olsson, D. Amodei, T. Brown, J. Clark, S. McCandlish, C. Olah, B. Mann, and J. Kaplan. 2022a. Training a helpful and harmless assistant with reinforcement learning from human feedback.
- Bai, Y., S. Kadavath, S. Kundu, A. Askell, J. Kernion, A. Jones, A. Chen, A. Goldie, A. Mirhoseini, C. McKinnon, C. Chen, C. Olsson, C. Olah, D. Hernandez, D. Drain, D. Ganguli, D. Li, E. Tran-Johnson, E. Perez, J. Kerr, J. Mueller, J. Ladish, J. Landau, K. Ndousse, K. Lukosuite, L. Lovitt, M. Sellitto, N. Elhage, N. Schiefer, N. Mercado, N. DasSarma, R. Lasenby, R. Larson, S. Ringer, S. Johnston, S. Kravec, S. El Showk, S. Fort, T. Lanham, T. Telleen-Lawton, T. Conerly, T. Henighan, T. Hume, S. R. Bowman, Z. Hatfield-Dodds, B. Mann, D. Amodei, N. Joseph, S. McCandlish, T. Brown, and J. Kaplan. 2022b. Constitutional AI: Harmlessness from AI feedback. ArXiv.
- Bajaj, P., D. Campos, N. Craswell, L. Deng, J. G. and Xiaodong Liu, R. Majumder, A. McNamara, B. Mitra, T. Nguye, M. Rosenberg, X. Song, A. Stoica, S. Tiwary, and T. Wang. 2016. MS MARCO: A human generated Machine Reading Comprehension dataset. *NeurIPS*.
- Baker, C. F., C. J. Fillmore, and J. B. Lowe. 1998. The Berkeley FrameNet project. *COLING/ACL*.
- Baker, J. K. 1975a. The DRAGON system – An overview. *IEEE Transactions on ASSP*, ASSP-23(1):24–29.
- Baker, J. K. 1975b. Stochastic modeling for automatic speech understanding. In D. R. Reddy, ed., *Speech Recognition*. Academic Press.
- Baldrige, J., N. Asher, and J. Hunter. 2007. Annotation for and robust parsing of discourse structure on unrestricted texts. *Zeitschrift für Sprachwissenschaft*, 26:213–239.
- Bamman, D., O. Lewke, and A. Mansoor. 2020. An annotated dataset of coreference in English literature. *LREC*.
- Bamman, D., B. O'Connor, and N. A. Smith. 2013. Learning latent personas of film characters. *ACL*.
- Bamman, D., S. Popat, and S. Shen. 2019. An annotated dataset of literary entities. *NAACL HLT*.
- Banerjee, S. and A. Lavie. 2005. ME-TOR: An automatic metric for MT evaluation with improved correlation with human judgments. *Proceedings of ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for MT and/or Summarization*.
- Banko, M., M. Cafarella, S. Soderland, M. Broadhead, and O. Etzioni. 2007. Open information extraction for the web. *IJCAI*.
- Bañón, M., P. Chen, B. Haddow, K. Heafield, H. Hoang, M. Esplà-Gomis, M. L. Forcada, A. Kamran, F. Kirefu, P. Koehn, S. Ortiz Rojas, L. Pla Sempere, G. Ramírez-Sánchez, E. Sarriás, M. Strelec, B. Thompson, W. Waites, D. Wiggins, and J. Zaragoza. 2020. ParaCrawl: Web-scale acquisition of parallel corpora. *ACL*.
- Bar-Hillel, Y. 1960. The present status of automatic translation of languages. In F. Alt, ed., *Advances in Computers 1*, 91–163. Academic Press.
- Barker, C. 2010. Nominals don't provide criteria of identity. In M. Rathert and A. Alexiadou, eds, *The Semantics of Nominalizations across Languages and Frameworks*, 9–24. Mouton.
- Barrett, L. F., B. Mesquita, K. N. Ochsner, and J. J. Gross. 2007. The experience of emotion. *Annual Review of Psychology*, 58:373–403.
- Barzilay, R. and M. Lapata. 2005. Modeling local coherence: An entity-based approach. *ACL*.
- Barzilay, R. and M. Lapata. 2008. Modeling local coherence: An entity-based approach. *Computational Linguistics*, 34(1):1–34.
- Barzilay, R. and L. Lee. 2004. Catching the drift: Probabilistic content models, with applications to generation and summarization. *HLT-NAACL*.
- Baum, L. E. and J. A. Eagon. 1967. An inequality with applications to statistical estimation for probabilistic functions of Markov processes and to a model for ecology. *Bulletin of the American Mathematical Society*, 73(3):360–363.
- Baum, L. E. and T. Petrie. 1966. Statistical inference for probabilistic functions of finite-state Markov chains. *Annals of Mathematical Statistics*, 37(6):1554–1563.
- Bazell, C. E. 1952/1966. The correspondence fallacy in structural linguistics. In E. P. Hamp, F. W. Householder, and R. Austerlitz, eds, *Studies by Members of the English Department, Istanbul University (3)*, reprinted in *Readings in Linguistics II (1966)*, 271–298. University of Chicago Press.
- Bean, D. and E. Riloff. 1999. Corpus-based identification of non-anaphoric noun phrases. *ACL*.
- Bean, D. and E. Riloff. 2004. Unsupervised learning of contextual role knowledge for coreference resolution. *HLT-NAACL*.
- Beckman, M. E. and G. M. Ayers. 1997. Guidelines for ToBI labelling. Unpublished manuscript, Ohio State University, http://www.ling.ohio-state.edu/research/phonetics/E_ToBI/.
- Beckman, M. E. and J. Hirschberg. 1994. The ToBI annotation conventions. Manuscript, Ohio State University.
- Bedi, G., F. Carrillo, G. A. Cecchi, D. F. Slezak, M. Sigman, N. B. Mota, S. Ribeiro, D. C. Javitt, M. Copelli, and C. M. Corcoran. 2015. Automated analysis of free speech predicts psychosis onset in high-risk youths. *npj Schizophrenia*, 1.
- Bejček, E., E. Hajičová, J. Hajič, P. Jínová, V. Kettnerová, V. Kolářová, M. Mikulová,

- J. Mírovský, A. Nedoluzhko, J. Panevová, L. Poláková, M. Ševčíková, J. Štěpánek, and Š. Zikánová. 2013. *Prague dependency treebank 3.0*. Technical report, Institute of Formal and Applied Linguistics, Charles University in Prague. LINDAT/CLARIN digital library at Institute of Formal and Applied Linguistics, Charles University in Prague.
- Belinkov, Y. 2022. *Probing classifiers: Promises, shortcomings, and advances*. *Computational Linguistics*, 48(1):207–219.
- Belinkov, Y., N. Durrani, F. Dalvi, H. Sajjad, and J. Glass. 2017. *What do neural machine translation models learn about morphology?* *ACL*.
- Bellegarda, J. R. 1997. *A latent semantic analysis framework for large-span language modeling*. *EUROSPEECH*.
- Bellegarda, J. R. 2000. *Exploiting latent semantic information in statistical language modeling*. *Proceedings of the IEEE*, 89(8):1279–1296.
- Bellman, R. 1957. *Dynamic Programming*. Princeton University Press.
- Bellman, R. 1984. *Eye of the Hurricane: an autobiography*. World Scientific Singapore.
- Bender, E. M. 2019. *The #BenderRule: On naming the languages we study and why it matters*. Blog post.
- Bender, E. M., B. Friedman, and A. McMillan-Major. 2021a. *A guide for writing data statements for natural language processing*. <http://techpolicylab.uw.edu/data-statements/>.
- Bender, E. M., T. Gebru, A. McMillan-Major, and S. Shmitchell. 2021b. *On the dangers of stochastic parrots: Can language models be too big?* *ACM FAccT*.
- Bender, E. M. and A. Koller. 2020. *Climbing towards NLU: On meaning, form, and understanding in the age of data*. *ACL*.
- Bengio, Y., A. Courville, and P. Vincent. 2013. *Representation learning: A review and new perspectives*. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(8):1798–1828.
- Bengio, Y., R. Ducharme, and P. Vincent. 2000. *A neural probabilistic language model*. *NeurIPS*.
- Bengio, Y., R. Ducharme, P. Vincent, and C. Jauvin. 2003. *A neural probabilistic language model*. *JMLR*, 3:1137–1155.
- Bengio, Y., P. Lamblin, D. Popovici, and H. Larochelle. 2007. *Greedy layer-wise training of deep networks*. *NeurIPS*.
- Bengio, Y., H. Schwenk, J.-S. Senécal, F. Morin, and J.-L. Gauvain. 2006. *Neural probabilistic language models*. In *Innovations in Machine Learning*, 137–186. Springer.
- Bengtson, E. and D. Roth. 2008. *Understanding the value of features for coreference resolution*. *EMNLP*.
- Bennett, R. and E. Elfiner. 2019. *The syntax–prosody interface*. *Annual Review of Linguistics*, 5:151–171.
- Bentivogli, L., M. Cettolo, M. Federico, and C. Federmann. 2018. *Machine translation human evaluation: an investigation of evaluation based on post-editing and its relation with direct assessment*. *ICSLT*.
- Berant, J., A. Chou, R. Frostig, and P. Liang. 2013. *Semantic parsing on freebase from question-answer pairs*. *EMNLP*.
- Berg-Kirkpatrick, T., D. Burkett, and D. Klein. 2012. *An empirical investigation of statistical significance in NLP*. *EMNLP*.
- Berger, A., S. A. Della Pietra, and V. J. Della Pietra. 1996. *A maximum entropy approach to natural language processing*. *Computational Linguistics*, 22(1):39–71.
- Bergsma, S. and D. Lin. 2006. *Bootstrapping path-based pronoun resolution*. *COLING/ACL*.
- Bergsma, S., D. Lin, and R. Goebel. 2008a. *Discriminative learning of selectional preference from unlabeled text*. *EMNLP*.
- Bergsma, S., D. Lin, and R. Goebel. 2008b. *Distributional identification of non-referential pronouns*. *ACL*.
- Bethard, S. 2013. *ClearTK-TimeML: A minimalist approach to TempEval 2013*. *SemEval-13*.
- Bhat, I., R. A. Bhat, M. Shrivastava, and D. Sharma. 2017. *Joining hands: Exploiting monolingual treebanks for parsing of code-mixing data*. *EACL*.
- Bianchi, F., M. Suzgun, G. Attanasio, P. Rottger, D. Jurafsky, T. Hashimoto, and J. Zou. 2024a. *Safety-tuned LLaMAs: Lessons from improving the safety of large language models that follow instructions*. *ICLR*.
- Bianchi, F., M. Suzgun, G. Attanasio, P. Röttger, D. Jurafsky, T. Hashimoto, and J. Y. Zou. 2024b. *Safety-tuned llamas: Lessons from improving the safety of large language models that follow instructions*. *ICLR*, volume 2024.
- Bickel, B. 2003. *Referential density in discourse and syntactic typology*. *Language*, 79(2):708–736.
- Bikel, D. M., S. Miller, R. Schwartz, and R. Weischedel. 1997. *Nymble: A high-performance learning name-finder*. *ANLP*.
- Biran, O. and K. McKeown. 2015. *PDTB discourse parsing as a tagging task: The two taggers approach*. *SIGDIAL*.
- Bird, S. 2020. *Decolonising speech and language technology*. *COLING*.
- Bird, S., E. Klein, and E. Loper. 2009. *Natural Language Processing with Python*. O’Reilly.
- Bisani, M. and H. Ney. 2004. *Bootstrap estimates for confidence intervals in ASR performance evaluation*. *ICASSP*.
- Bishop, C. M. 2006. *Pattern recognition and machine learning*. Springer.
- Bisk, Y., A. Holtzman, J. Thomason, J. Andreas, Y. Bengio, J. Chai, M. Lapata, A. Lazaridou, J. May, A. Nisnevich, N. Pinto, and J. Turian. 2020. *Experience grounds language*. *EMNLP*.
- Bizer, C., J. Lehmann, G. Kobilarov, S. Auer, C. Becker, R. Cyganiak, and S. Hellmann. 2009. *DBpedia—A crystallization point for the Web of Data*. *Web Semantics: science, services and agents on the world wide web*, 7(3):154–165.
- Björkelund, A. and J. Kuhn. 2014. *Learning structured perceptrons for coreference resolution with latent antecedents and non-local features*. *ACL*.
- Black, A. W. and P. Taylor. 1994. *CHATR: A generic speech synthesis system*. *COLING*.
- Black, E., S. P. Abney, D. Flickinger, C. Gdaniec, R. Grishman, P. Harrison, D. Hindle, R. Ingria, F. Jelinek, J. L. Klavans, M. Y. Liberman, M. P. Marcus, S. Roukos, B. Santorini, and T. Strzalkowski. 1991. *A procedure for quantitatively comparing the syntactic coverage of English grammars*. *Speech and Natural Language Workshop*.
- Blei, D. M., A. Y. Ng, and M. I. Jordan. 2003. *Latent Dirichlet allocation*. *JMLR*, 3(5):993–1022.
- Blodgett, S. L., S. Barocas, H. Daumé III, and H. Wallach. 2020. *Language (technology) is power: A critical survey of “bias” in NLP*. *ACL*.
- Blodgett, S. L., L. Green, and B. O’Connor. 2016. *Demographic dialectal variation in social media: A case study of African-American English*. *EMNLP*.
- Blodgett, S. L. and B. O’Connor. 2017. *Racial disparity in natural language processing: A case study of social media African-American English*. *FAT/ML Workshop, KDD*.

- Bloomfield, L. 1914. *An Introduction to the Study of Language*. Henry Holt and Company.
- Bloomfield, L. 1933. *Language*. University of Chicago Press.
- Bobrow, D. G., R. M. Kaplan, M. Kay, D. A. Norman, H. Thompson, and T. Winograd. 1977. GUS, A frame driven dialog system. *Artificial Intelligence*, 8:155–173.
- Bobrow, D. G. and D. A. Norman. 1975. Some principles of memory schemata. In D. G. Bobrow and A. Collins, eds, *Representation and Understanding*. Academic Press.
- Boersma, P. and D. Weenink. 2005. Praat: doing phonetics by computer (version 4.3.14). [Computer program]. Retrieved May 26, 2005, from <http://www.praat.org/>.
- Bojanowski, P., E. Grave, A. Joulin, and T. Mikolov. 2017. [Enriching word vectors with subword information](#). *TACL*, 5:135–146.
- Bollacker, K., C. Evans, P. Paritosh, T. Sturge, and J. Taylor. 2008. Freebase: a collaboratively created graph database for structuring human knowledge. *SIGMOD 2008*.
- Bolukbasi, T., K.-W. Chang, J. Zou, V. Saligrama, and A. T. Kalai. 2016. Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. *NeurIPS*.
- Bommasani, R., D. A. Hudson, E. Adeli, R. Altman, S. Arora, S. von Arx, M. S. Bernstein, J. Bohg, A. Bosselut, E. Brunskill, E. Brynjolfsson, N. S. Chatterji, A. S. Chen, K. A. Creel, J. Davis, D. Demszyk, C. Donahue, M. Doumbouya, E. Durmus, S. Ermon, J. Etchemendy, K. Ethayarajh, L. Fei-Fei, C. Finn, T. Gale, L. E. Gillespie, K. Goel, N. D. Goodman, S. Grossman, N. Guha, T. Hashimoto, P. Henderson, J. Hewitt, D. E. Ho, J. Hong, K. Hsu, J. Huang, T. F. Icard, S. Jain, D. Jurafsky, P. Kalluri, S. Karamcheti, G. Keeling, F. Khani, O. Khat-tab, P. W. Koh, M. S. Krass, R. Krishna, R. Kuditipudi, A. Kumar, F. Ladhak, M. Lee, T. Lee, J. Leskovec, I. Levent, X. L. Li, X. Li, T. Ma, A. Malik, C. D. Manning, S. P. Mirchandani, E. Mitchell, Z. Munyikwa, S. Nair, A. Narayan, D. Narayanan, B. Newman, A. Nie, J. C. Niebles, H. Nilforoshan, J. F. Nyarko, G. Ogut, L. Orr, I. Papadimitriou, J. S. Park, C. Piech, E. Portelance, C. Potts, A. Raghunathan, R. Reich, H. Ren, F. Rong, Y. H. Roohani, C. Ruiz, J. Ryan, C. R’e, D. Sadigh, S. Sagawa, K. Santhanam, A. Shih, K. P. Srinivasan, A. Tamkin, R. Taori, A. W. Thomas, F. Tramèr, R. E. Wang, W. Wang, B. Wu, J. Wu, Y. Wu, S. M. Xie, M. Yasunaga, J. You, M. A. Zaharia, M. Zhang, T. Zhang, X. Zhang, Y. Zhang, L. Zheng, K. Zhou, and P. Liang. 2021. [On the opportunities and risks of foundation models](#). *ArXiv*.
- Booth, T. L. 1969. Probabilistic representation of formal languages. *IEEE Conference Record of the 1969 Tenth Annual Symposium on Switching and Automata Theory*.
- Borges, J. L. 1964. The analytical language of John Wilkins. In *Other inquiries 1937–1952*. University of Texas Press. Trans. Ruth L. C. Simms.
- Bostrom, K. and G. Durrett. 2020. [Byte pair encoding is suboptimal for language model pretraining](#). *EMNLP*.
- Bourlard, H. and N. Morgan. 1994. *Connectionist Speech Recognition: A Hybrid Approach*. Kluwer.
- Bradley, R. A. and M. E. Terry. 1952. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39:324–345.
- Brants, T. 2000. [TnT: A statistical part-of-speech tagger](#). *ANLP*.
- Brants, T., A. C. Popat, P. Xu, F. J. Och, and J. Dean. 2007. [Large language models in machine translation](#). *EMNLP/CoNLL*.
- Braud, C., M. Coavoux, and A. Søgaard. 2017. [Cross-lingual RST discourse parsing](#). *EACL*.
- Bréal, M. 1897. *Essai de Sémantique: Science des significations*. Hachette.
- Brennan, S. E., M. W. Friedman, and C. Pollard. 1987. [A centering approach to pronouns](#). *ACL*.
- Brin, S. 1998. Extracting patterns and relations from the World Wide Web. *Proceedings World Wide Web and Databases International Workshop, Number 1590 in LNCS*. Springer.
- Brockmann, C. and M. Lapata. 2003. [Evaluating and combining approaches to selectional preference acquisition](#). *EACL*.
- Broschart, J. 1997. Why Tongan does it differently. *Linguistic Typology*, 1:123–165.
- Brown, P. F., J. Cocke, S. A. Della Pietra, V. J. Della Pietra, F. Jelinek, J. D. Lafferty, R. L. Mercer, and P. S. Roossin. 1990. [A statistical approach to machine translation](#). *Computational Linguistics*, 16(2):79–85.
- Brown, P. F., S. A. Della Pietra, V. J. Della Pietra, and R. L. Mercer. 1993. [The mathematics of statistical machine translation: Parameter estimation](#). *Computational Linguistics*, 19(2):263–311.
- Brown, T., B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. M. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei. 2020. Language models are few-shot learners. *NeurIPS*, volume 33.
- Brynjolfsson, E. 2022. [The Turing Trap: The promise & peril of human-like artificial intelligence](#). *Daedalus*, 151(2):272–287.
- Brynjolfsson, E., X. Hui, and M. Liu. 2019. [Does machine translation affect international trade? evidence from a large digital platform](#). *Management Science*, 65(12):5449–5460.
- Brysbaert, M., A. B. Warriner, and V. Kuperman. 2014. Concrete-ness ratings for 40 thousand generally known English word lemmas. *Behavior Research Methods*, 46(3):904–911.
- Bu, H., J. Du, X. Na, B. Wu, and H. Zheng. 2017. AISHELL-1: An open-source Mandarin speech corpus and a speech recognition baseline. *O-COCOSDA Proceedings*.
- Buchholz, S. and E. Marsi. 2006. [Conll-x shared task on multilingual dependency parsing](#). *CoNLL*.
- Budanitsky, A. and G. Hirst. 2006. [Evaluating WordNet-based measures of lexical semantic relatedness](#). *Computational Linguistics*, 32(1):13–47.
- Bullinaria, J. A. and J. P. Levy. 2007. Extracting semantic representations from word co-occurrence statistics: A computational study. *Behavior research methods*, 39(3):510–526.
- Bullinaria, J. A. and J. P. Levy. 2012. Extracting semantic representations from word co-occurrence statistics: stop-lists, stemming, and SVD. *Behavior research methods*, 44(3):890–907.
- Caliskan, A., J. J. Bryson, and A. Narayanan. 2017. [Semantics derived automatically from language corpora contain human-like biases](#). *Science*, 356(6334):183–186.
- Callison-Burch, C., M. Osborne, and P. Koehn. 2006. [Re-evaluating the role of BLEU in machine translation research](#). *EACL*.

- Canavan, A., D. Graff, and G. Zipperlen. 1997. CALLHOME American English speech LDC97S42. Linguistic Data Consortium.
- Cao, Y., Y. Kang, C. Wang, and L. Sun. 2024. Instruction mining: Instruction data selection for tuning large language models. *First Conference on Language Modeling*.
- Carbonell, J. R. 1970. AI in CAI: An artificial-intelligence approach to computer-assisted instruction. *IEEE transactions on man-machine systems*, 11(4):190–202.
- Card, D., S. Chang, C. Becker, J. Mendelsohn, R. Voigt, L. Bousttan, R. Abramitzky, and D. Jurafsky. 2022. Computational analysis of 140 years of US political speeches reveals more positive but increasingly polarized framing of immigration. *PNAS*, 119(31):e2120510119.
- Cardie, C. 1993. A case-based approach to knowledge acquisition for domain specific sentence analysis. *AAAI*.
- Cardie, C. 1994. *Domain-Specific Knowledge Acquisition for Conceptual Sentence Analysis*. Ph.D. thesis, University of Massachusetts, Amherst, MA. Available as CMP-SCI Technical Report 94-74.
- Cardie, C. and K. Wagstaff. 1999. Noun phrase coreference as clustering. *EMNLP/VLC*.
- Carletta, J., N. Dahlbäck, N. Reithinger, and M. A. Walker. 1997. Standards for dialogue coding in natural language processing. Technical Report 167, Dagstuhl Seminars. Report from Dagstuhl seminar number 9706.
- Carlson, G. N. 1977. *Reference to kinds in English*. Ph.D. thesis, University of Massachusetts, Amherst. Forward.
- Carlson, L. and D. Marcu. 2001. Discourse tagging manual. Technical Report ISI-TR-545, ISI.
- Carlson, L., D. Marcu, and M. E. Okunowski. 2001. *Building a discourse-tagged corpus in the framework of rhetorical structure theory*. *SIGDIAL*.
- Carreras, X. and L. Màrquez. 2005. Introduction to the CoNLL-2005 shared task: Semantic role labeling. *CoNLL*.
- Chafe, W. L. 1976. Givenness, contrastiveness, definiteness, subjects, topics, and point of view. In C. N. Li, ed., *Subject and Topic*, 25–55. Academic Press.
- Chambers, N. 2013. *NavyTime: Event and time ordering from raw text*. *SemEval-13*.
- Chambers, N., T. Cassidy, B. McDowell, and S. Bethard. 2014. Dense event ordering with a multi-pass architecture. *TACL*, 2:273–284.
- Chambers, N. and D. Jurafsky. 2010. Improving the use of pseudo-words for evaluating selectional preferences. *ACL*.
- Chambers, N. and D. Jurafsky. 2011. Template-based information extraction without the templates. *ACL*.
- Chan, W., N. Jaitly, Q. Le, and O. Vinyals. 2016. Listen, attend and spell: A neural network for large vocabulary conversational speech recognition. *ICASSP*.
- Chandioux, J. 1976. MÉTÉO: un système opérationnel pour la traduction automatique des bulletins météorologiques destinés au grand public. *Meta*, 21:127–133.
- Chang, A. X. and C. D. Manning. 2012. *SUTime: A library for recognizing and normalizing time expressions*. *LREC*.
- Chang, K.-W., R. Samdani, and D. Roth. 2013. A constrained latent variable model for coreference resolution. *EMNLP*.
- Chang, K.-W., R. Samdani, A. Rozovskaya, M. Sammons, and D. Roth. 2012. *Illinois-Coref: The UI system in the CoNLL-2012 shared task*. *CoNLL*.
- Chaplot, D. S. and R. Salakhutdinov. 2018. Knowledge-based word sense disambiguation using topic models. *AAAI*.
- Charniak, E. 1997. Statistical parsing with a context-free grammar and word statistics. *AAAI*.
- Charniak, E., C. Hendrickson, N. Jacobson, and M. Perkowski. 1993. Equations for part-of-speech tagging. *AAAI*.
- Che, W., Z. Li, Y. Li, Y. Guo, B. Qin, and T. Liu. 2009. Multilingual dependency-based syntactic and semantic parsing. *CoNLL*.
- Chen, C. and V. Ng. 2013. Linguistically aware coreference evaluation metrics. *IJCNLP*.
- Chen, D., A. Fisch, J. Weston, and A. Bordes. 2017a. Reading Wikipedia to answer open-domain questions. *ACL*.
- Chen, D. and C. Manning. 2014. A fast and accurate dependency parser using neural networks. *EMNLP*.
- Chen, E., B. Snyder, and R. Barzilay. 2007. Incremental text structuring with online hierarchical ranking. *EMNLP/CoNLL*.
- Chen, S., C. Wang, Y. Wu, Z. Zhang, L. Zhou, S. Liu, Z. Chen, Y. Liu, H. Wang, J. Li, L. He, S. Zhao, and F. Wei. 2025. Neural codec language models are zero-shot text to speech synthesizers. *IEEE Trans. on TASLP*, 33:705–718.
- Chen, S. F. and J. Goodman. 1999. An empirical study of smoothing techniques for language modeling. *Computer Speech and Language*, 13:359–394.
- Chen, X., Z. Shi, X. Qiu, and X. Huang. 2017b. Adversarial multi-criteria learning for Chinese word segmentation. *ACL*.
- Cheng, J., L. Dong, and M. Lapata. 2016. Long short-term memory-networks for machine reading. *EMNLP*.
- Cheng, M., E. Durmus, and D. Jurafsky. 2023. Marked personas: Using natural language prompts to measure stereotypes in language models. *ACL*.
- Cheng, M., C. Lee, P. Khadpe, S. Yu, D. Han, and D. Jurafsky. 2026. Sycophantic AI decreases prosocial intentions and promotes dependence. *Science*, 391(6792):eac8352.
- Chiang, D. 2005. A hierarchical phrase-based model for statistical machine translation. *ACL*.
- Chinchor, N., L. Hirschman, and D. L. Lewis. 1993. Evaluating Message Understanding systems: An analysis of the third Message Understanding Conference. *Computational Linguistics*, 19(3):409–449.
- Chiticariu, L., M. Danilevsky, Y. Li, F. Reiss, and H. Zhu. 2018. *SystemT: Declarative text understanding for enterprise*. *NAACL HLT*, volume 3.
- Chiticariu, L., Y. Li, and F. R. Reiss. 2013. Rule-Based Information Extraction is Dead! Long Live Rule-Based Information Extraction Systems! *EMNLP*.
- Chiu, J. P. C. and E. Nichols. 2016. Named entity recognition with bidirectional LSTM-CNNs. *TACL*, 4:357–370.
- Cho, K., B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio. 2014. Learning phrase representations using RNN encoder-decoder for statistical machine translation. *EMNLP*.
- Choe, D. K. and E. Charniak. 2016. Parsing as language modeling. *EMNLP*.
- Choi, J. D. and M. Palmer. 2011a. Getting the most out of transition-based dependency parsing. *ACL*.
- Choi, J. D. and M. Palmer. 2011b. Transition-based semantic role labeling using predicate argument

- clustering. *Proceedings of the ACL 2011 Workshop on Relational Models of Semantics*.
- Choi, J. D., J. Tetreault, and A. Stent. 2015. It depends: Dependency parser comparison using a web-based evaluation tool. *ACL*.
- Chomsky, N. 1956. Three models for the description of language. *IRE Transactions on Information Theory*, 2(3):113–124.
- Chomsky, N. 1956/1975. *The Logical Structure of Linguistic Theory*. Plenum.
- Chomsky, N. 1957. *Syntactic Structures*. Mouton.
- Chomsky, N. 1963. Formal properties of grammars. In R. D. Luce, R. Bush, and E. Galanter, eds, *Handbook of Mathematical Psychology*, volume 2, 323–418. Wiley.
- Chomsky, N. 1981. *Lectures on Government and Binding*. Foris.
- Chorowski, J., D. Bahdanau, K. Cho, and Y. Bengio. 2014. End-to-end continuous speech recognition using attention-based recurrent NN: First results. *NeurIPS Deep Learning and Representation Learning Workshop*.
- Chou, W., C.-H. Lee, and B. H. Juang. 1993. Minimum error rate training based on n -best string models. *ICASSP*.
- Christodoulopoulos, C., S. Goldwater, and M. Steedman. 2010. Two decades of unsupervised POS induction: How far have we come? *EMNLP*.
- Chu, Y.-J. and T.-H. Liu. 1965. On the shortest arborescence of a directed graph. *Science Sinica*, 14:1396–1400.
- Chu-Carroll, J. and S. Carberry. 1998. Collaborative response generation in planning dialogues. *Computational Linguistics*, 24(3):355–400.
- Church, K. W. 1988. A stochastic parts program and noun phrase parser for unrestricted text. *ANLP*.
- Church, K. W. 1989. A stochastic parts program and noun phrase parser for unrestricted text. *ICASSP*.
- Church, K. W. 1994. Unix for Poets. Slides from 2nd ELSNET Summer School and unpublished paper ms.
- Church, K. W. and W. A. Gale. 1991. A comparison of the enhanced Good-Turing and deleted estimation methods for estimating probabilities of English bigrams. *Computer Speech and Language*, 5:19–54.
- Cialdini, R. B. 1984. *Influence: The psychology of persuasion*. Morrow.
- Cieri, C., D. Miller, and K. Walker. 2004. The Fisher corpus: A resource for the next generations of speech-to-text. *LREC*.
- Clark, E. 1987. The principle of contrast: A constraint on language acquisition. In B. MacWhinney, ed., *Mechanisms of language acquisition*, 1–33. LEA.
- Clark, H. H. 1996. *Using Language*. Cambridge University Press.
- Clark, H. H. and J. E. Fox Tree. 2002. Using uh and um in spontaneous speaking. *Cognition*, 84:73–111.
- Clark, H. H. and E. F. Schaefer. 1989. Contributing to discourse. *Cognitive Science*, 13:259–294.
- Clark, J. and C. Yallop. 1995. *An Introduction to Phonetics and Phonology*, 2nd edition. Blackwell.
- Clark, J. H., E. Choi, M. Collins, D. Garrette, T. Kwiatkowski, V. Nikolaev, and J. Palomaki. 2020a. TyDi QA: A benchmark for information-seeking question answering in typologically diverse languages. *TACL*, 8:454–470.
- Clark, K., M.-T. Luong, Q. V. Le, and C. D. Manning. 2020b. Electra: Pre-training text encoders as discriminators rather than generators. *ICLR*.
- Clark, K. and C. D. Manning. 2015. Entity-centric coreference resolution with model stacking. *ACL*.
- Clark, K. and C. D. Manning. 2016a. Deep reinforcement learning for mention-ranking coreference models. *EMNLP*.
- Clark, K. and C. D. Manning. 2016b. Improving coreference resolution by learning entity-level distributed representations. *ACL*.
- Clark, S., J. R. Curran, and M. Osborne. 2003. Bootstrapping POS-taggers using unlabelled data. *CoNLL*.
- CMU. 1993. The Carnegie Mellon Pronouncing Dictionary v0.1. Carnegie Mellon University.
- Cobbe, K., V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano, C. Hesse, and J. Schulman. 2021a. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Cobbe, K., V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano, C. Hesse, and J. Schulman. 2021b. Training verifiers to solve math word problems. *ArXiv preprint*.
- Coccaro, N. and D. Jurafsky. 1998. Towards better integration of semantic predictors in statistical language modeling. *ICSLP*.
- Coenen, A., E. Reif, A. Yuan, B. Kim, A. Pearce, F. Viégas, and M. Wattenberg. 2019. Visualizing and measuring the geometry of bert. *NeurIPS*.
- Coleman, J. 2005. *Introducing Speech and Language Processing*. Cambridge University Press.
- Collante, A., S. Abedu, S. Khatounabadi, A. Abdellatif, E. Alor, and E. Shihab. 2025. The impact of large language models (LLMs) on code review process. *ArXiv preprint*.
- Collins, M. 1999. *Head-Driven Statistical Models for Natural Language Parsing*. Ph.D. thesis, University of Pennsylvania, Philadelphia.
- Collobert, R. and J. Weston. 2007. Fast semantic extraction using a novel neural network architecture. *ACL*.
- Collobert, R. and J. Weston. 2008. A unified architecture for natural language processing: Deep neural networks with multitask learning. *ICML*.
- Collobert, R., J. Weston, L. Bottou, M. Karlen, K. Kavukcuoglu, and P. Kuksa. 2011. Natural language processing (almost) from scratch. *JMLR*, 12:2493–2537.
- Comrie, B. 1989. *Language Universals and Linguistic Typology*, 2nd edition. Blackwell.
- Conneau, A., K. Khandelwal, N. Goyal, V. Chaudhary, G. Wenzek, F. Guzmán, E. Grave, M. Ott, L. Zettlemoyer, and V. Stoyanov. 2020. Unsupervised cross-lingual representation learning at scale. *ACL*.
- Conneau, A., M. Ma, S. Khanuja, Y. Zhang, V. Axelrod, S. Dalmia, J. Riesa, C. Rivera, and A. Bapna. 2023. Fleurs: Few-shot learning evaluation of universal representations of speech. *IEEE SLT*.
- Connolly, D., J. D. Burger, and D. S. Day. 1994. A machine learning approach to anaphoric reference. *Proceedings of the International Conference on New Methods in Language Processing (NeMLaP)*.
- Cooley, J. W. and J. W. Tukey. 1965. An algorithm for the machine calculation of complex Fourier series. *Mathematics of Computation*, 19(90):297–301.
- Cooper, F. S., A. M. Liberman, and J. M. Borst. 1951. The interconversion of audible and visible patterns as a basis for research in the perception of speech. *Proceedings of the National Academy of Sciences*, 37(5):318–325.
- Cordier, B. 1965. Factor-analysis of correspondences. *COLING 1965*.
- Core, M., M. Ishizaki, J. D. Moore, C. Nakatani, N. Reithinger, D. R. Traum, and S. Tutiya. 1999. The report of the 3rd workshop of the Discourse Resource Initiative. Technical Report No.3 CC-TR-99-1, Chiba Corpus Project, Chiba, Japan.

- Costa-jussà, M. R., J. Cross, O. Çelebi, M. Elbayad, K. Heafield, K. Hefernan, E. Kalbassi, J. Lam, D. Licht, J. Maillard, A. Sun, S. Wang, G. Wenzek, A. Youngblood, B. Akula, L. Barrault, G. M. Gonzalez, P. Hansanti, J. Hoffman, S. Jarrett, K. R. Sadagopan, D. Rowe, S. Spruit, C. Tran, P. Andrews, N. F. Ayan, S. Bhosale, S. Edunov, A. Fan, C. Gao, V. Goswami, F. Guzmán, P. Koehn, A. Mourachko, C. Ropers, S. Saleem, H. Schwenk, J. Wang, and NLLB Team. 2022. [No language left behind: Scaling human-centered machine translation](#). ArXiv.
- Cover, T. M. and J. A. Thomas. 1991. *Elements of Information Theory*. Wiley.
- Covington, M. 2001. A fundamental algorithm for dependency parsing. *Proceedings of the 39th Annual ACM Southeast Conference*.
- Cox, D. 1969. *Analysis of Binary Data*. Chapman and Hall, London.
- Crane, L. D. and P. E. Soto. 2026. [Ai and coder employment: Compiling the evidence](#). FEDS Working Paper.
- Craven, M. and J. Kumlien. 1999. Constructing biological knowledge bases by extracting information from text sources. *ISMB-99*.
- Crawford, K. 2017. The trouble with bias. Keynote at NeurIPS.
- Croft, W. 1990. *Typology and Universals*. Cambridge University Press.
- Crosbie, J. and E. Shutova. 2022. [Induction heads as an essential mechanism for pattern matching in in-context learning](#). ArXiv preprint.
- Cross, J. and L. Huang. 2016. [Span-based constituency parsing with a structure-label system and provably optimal dynamic oracles](#). *EMNLP*.
- Cruse, D. A. 2004. *Meaning in Language: an Introduction to Semantics and Pragmatics*. Oxford University Press. Second edition.
- Cucerzan, S. 2007. [Large-scale named entity disambiguation based on Wikipedia data](#). *EMNLP/CoNLL*.
- Cui, G., L. Yuan, N. Ding, G. Yao, B. He, W. Zhu, Y. Ni, G. Xie, R. Xie, Y. Lin, Z. Liu, and M. Sun. 2024. Ultrafeedback: boosting language models with scaled ai feedback. *ICML 2024*.
- Dahl, G. E., T. N. Sainath, and G. E. Hinton. 2013. Improving deep neural networks for LVCSR using rectified linear units and dropout. *ICASSP*.
- Dahl, G. E., D. Yu, L. Deng, and A. Acero. 2012. Context-dependent pre-trained deep neural networks for large-vocabulary speech recognition. *IEEE Transactions on audio, speech, and language processing*, 20(1):30–42.
- Dahl, M., V. Magesh, M. Suzgun, and D. E. Ho. 2024. [Large legal fictions: Profiling legal hallucinations in large language models](#). *Journal of Legal Analysis*, 16:64–93.
- Dai, A. M. and Q. V. Le. 2015. Semi-supervised sequence learning. *NeurIPS*.
- Das, S. R. and M. Y. Chen. 2001. Yahoo! for Amazon: Sentiment parsing from small talk on the web. EFA 2001 Barcelona Meetings. <http://ssrn.com/abstract=276189>.
- Dauphin, Y. N., A. Fan, M. Auli, and D. Grangier. 2017. [Language modeling with gated convolutional networks](#). *ICML*, volume 70.
- David, Jr., E. E. and O. G. Selfridge. 1962. Eyes and ears for computers. *Proceedings of the IRE (Institute of Radio Engineers)*, 50:1093–1101.
- Davidson, T., D. Bhattacharya, and I. Weber. 2019. [Racial bias in hate speech and abusive language detection datasets](#). *Third Workshop on Abusive Language Online*.
- Davies, M. 2012. Expanding horizons in historical linguistics with the 400-million word Corpus of Historical American English. *Corpora*, 7(2):121–157.
- Davies, M. 2015. The Wikipedia Corpus: 4.6 million articles, 1.9 billion words. Adapted from Wikipedia. <https://www.english-corpora.org/wiki/>.
- Davies, M. 2020. The Corpus of Contemporary American English (COCA): One billion words, 1990–2019. <https://www.english-corpora.org/coca/>.
- Davis, E., L. Morgenstern, and C. L. Ortiz. 2017. The first Winograd schema challenge at IJCAI-16. *AI Magazine*, 38(3):97–98.
- Davis, K. H., R. Biddulph, and S. Balashek. 1952. Automatic recognition of spoken digits. *JASA*, 24(6):637–642.
- Davis, S. and P. Mermelstein. 1980. Comparison of parametric representations for monosyllabic word recognition in continuously spoken sentences. *IEEE Transactions on ASSP*, 28(4):357–366.
- Deerwester, S. C., S. T. Dumais, G. W. Furnas, R. A. Harshman, T. K. Landauer, K. E. Lochbaum, and L. Streeter. 1988. Computer information retrieval using latent semantic structure: US Patent 4,839,853.
- Deerwester, S. C., S. T. Dumais, T. K. Landauer, G. W. Furnas, and R. A. Harshman. 1990. Indexing by latent semantics analysis. *JASIS*, 41(6):391–407.
- Défossiez, A., J. Copet, G. Synnaeve, and Y. Adi. 2023. [High fidelity neural audio compression](#). *TMLR*.
- DeJong, G. F. 1982. An overview of the FRUMP system. In W. G. Lehnert and M. H. Ringle, eds, *Strategies for Natural Language Processing*, 149–176. LEA.
- Denes, P. 1959. The design and operation of the mechanical speech recognizer at University College London. *Journal of the British Institution of Radio Engineers*, 19(4):219–234. Appears together with companion paper (Fry 1959).
- Deng, L., G. Hinton, and B. Kingsbury. 2013. New types of deep neural network learning for speech recognition and related applications: An overview. *ICASSP*.
- Deng, Y. and W. Byrne. 2005. [HMM word and phrase alignment for statistical machine translation](#). *HLT-EMNLP*.
- Denis, P. and J. Baldridge. 2007. [Joint determination of anaphoricity and coreference resolution using integer programming](#). *NAACL-HLT*.
- Denis, P. and J. Baldridge. 2008. [Specialized models and ranking for coreference resolution](#). *EMNLP*.
- Denis, P. and J. Baldridge. 2009. Global joint models for coreference resolution and named entity classification. *Procesamiento del Lenguaje Natural*, 42.
- Dennett, D. C. 2009. Intentional systems theory. In A. Beckermann, B. P. McLaughlin, and S. Walter, eds, *The Oxford Handbook of Philosophy of Mind*, 339–350. Oxford University Press.
- DeRose, S. J. 1988. [Grammatical category disambiguation by statistical optimization](#). *Computational Linguistics*, 14:31–39.
- Devlin, J., M.-W. Chang, K. Lee, and K. Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). *NAACL HLT*.
- Di Eugenio, B. 1990. [Centering theory and the Italian pronominal system](#). *COLING*.
- Di Eugenio, B. 1996. [The discourse functions of Italian subjects: A centering approach](#). *COLING*.
- Dias Oliva, T., D. Antonialli, and A. Gomes. 2021. [Fighting hate speech, silencing drag queens? artificial intelligence in content moderation and risks to lgbtq voices online](#). *Sexuality & Culture*, 25:700–732.

- Ditman, T. and G. R. Kuperberg. 2010. Building coherence: A framework for exploring the breakdown of links across clause boundaries in schizophrenia. *Journal of neurolinguistics*, 23(3):254–269.
- Dixon, L., J. Li, J. Sorensen, N. Thain, and L. Vasserman. 2018. Measuring and mitigating unintended bias in text classification. *2018 AAAI/ACM Conference on AI, Ethics, and Society*.
- Dixon, N. and H. Maxey. 1968. Terminal analog synthesis of continuous speech using the diphone method of segment assembly. *IEEE Transactions on Audio and Electroacoustics*, 16(1):40–50.
- Do, Q. N. T., S. Bethard, and M.-F. Moens. 2017. Improving implicit semantic role labeling by predicting semantic frame arguments. *IJCNLP*.
- Doddington, G. 2002. Automatic evaluation of machine translation quality using n-gram co-occurrence statistics. *HLT*.
- Dodge, J., S. Gururangan, D. Card, R. Schwartz, and N. A. Smith. 2019. Show your work: Improved reporting of experimental results. *EMNLP*.
- Dong, L. and M. Lapata. 2016. Language to logical form with neural attention. *ACL*.
- Dorr, B. 1994. Machine translation divergences: A formal description and proposed solution. *Computational Linguistics*, 20(4):597–633.
- Dostert, L. 1955. The Georgetown-IBM experiment. In *Machine Translation of Languages: Fourteen Essays*, 124–135. MIT Press.
- Doumbouya, M. K. B., D. Jurafsky, and C. D. Manning. 2025. Tversky neural networks: Psychologically plausible deep learning with differentiable tversky similarity. ArXiv preprint.
- Dowty, D. R. 1979. *Word Meaning and Montague Grammar*. D. Reidel.
- Dozat, T. and C. D. Manning. 2017. Deep biaffine attention for neural dependency parsing. *ICLR*.
- Dozat, T. and C. D. Manning. 2018. Simpler but more accurate semantic dependency parsing. *ACL*.
- Dozat, T., P. Qi, and C. D. Manning. 2017. Stanford’s graph-based neural dependency parser at the CoNLL 2017 shared task. *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*.
- Dror, R., G. Baumer, M. Bogomolov, and R. Reichart. 2017. Replicability analysis for natural language processing: Testing significance with multiple datasets. *TACL*, 5:471–486.
- Dror, R., L. Peled-Cohen, S. Shlomov, and R. Reichart. 2020. *Statistical Significance Testing for Natural Language Processing*, volume 45 of *Synthesis Lectures on Human Language Technologies*. Morgan & Claypool.
- Dryer, M. S. and M. Haspelmath, eds. 2013. *The World Atlas of Language Structures Online*. Max Planck Institute for Evolutionary Anthropology, Leipzig. Available online at <http://wals.info>.
- Durrett, G. and D. Klein. 2013. Easy victories and uphill battles in coreference resolution. *EMNLP*.
- Durrett, G. and D. Klein. 2014. A joint model for entity analysis: Coreference, typing, and linking. *TACL*, 2:477–490.
- Earley, J. 1968. *An Efficient Context-Free Parsing Algorithm*. Ph.D. thesis, Carnegie Mellon University, Pittsburgh, PA.
- Earley, J. 1970. An efficient context-free parsing algorithm. *CACM*, 6(8):451–455.
- Edmonds, J. 1967. Optimum branchings. *Journal of Research of the National Bureau of Standards B*, 71(4):233–240.
- Edunov, S., M. Ott, M. Auli, and D. Grangier. 2018. Understanding back-translation at scale. *EMNLP*.
- Efron, B. and R. J. Tibshirani. 1993. *An introduction to the bootstrap*. CRC press.
- Egghe, L. 2007. Untangling Herdan’s law and Heaps’ law: Mathematical and informetric arguments. *JASIST*, 58(5):702–709.
- Eisner, J. 1996. Three new probabilistic models for dependency parsing: An exploration. *COLING*.
- Ekman, P. 1999. Basic emotions. In T. Dalgleish and M. J. Power, eds, *Handbook of Cognition and Emotion*, 45–60. Wiley.
- Elhage, N., N. Nanda, C. Olsson, T. Henighan, N. Joseph, B. Mann, A. Askell, Y. Bai, A. Chen, T. Conerly, N. DasSarma, D. Drain, D. Ganguli, Z. Hatfield-Dodds, D. Hernandez, A. Jones, J. Kernion, L. Lovitt, K. Ndousse, D. Amodei, T. Brown, J. Clark, J. Kaplan, S. McCandlish, and C. Olah. 2021. A mathematical framework for transformer circuits. White paper.
- Elman, J. L. 1990. Finding structure in time. *Cognitive science*, 14(2):179–211.
- Elsner, M., J. Austerweil, and E. Charniak. 2007. A unified local and global model for discourse coherence. *NAACL-HLT*.
- Elsner, M. and E. Charniak. 2008. Coreference-inspired coherence modeling. *ACL*.
- Elsner, M. and E. Charniak. 2011. Extending the entity grid with entity-specific features. *ACL*.
- Elvevåg, B., P. W. Foltz, D. R. Weinberger, and T. E. Goldberg. 2007. Quantifying incoherence in speech: an automated methodology and novel application to schizophrenia. *Schizophrenia research*, 93(1-3):304–316.
- Emami, A. and F. Jelinek. 2005. A neural syntactic language model. *Machine learning*, 60(1):195–227.
- Emami, A., P. Trichelair, A. Trischler, K. Suleman, H. Schulz, and J. C. K. Cheung. 2019. The KNOWREF coreference corpus: Removing gender and number cues for difficult pronominal anaphora resolution. *ACL*.
- Erk, K. 2007. A simple, similarity-based model for selectional preferences. *ACL*.
- Ethayarajh, K. 2019. How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings. *EMNLP*.
- Ethayarajh, K., D. Duvenaud, and G. Hirst. 2019a. Towards understanding linear word analogies. *ACL*.
- Ethayarajh, K., D. Duvenaud, and G. Hirst. 2019b. Understanding undesirable word embedding associations. *ACL*.
- Ethayarajh, K. and D. Jurafsky. 2020. Utility is in the eye of the user: A critique of NLP leaderboards. *EMNLP*.
- Ethayarajh, K., H. C. Zhang, and S. Behzad. 2022. Stanford human preferences dataset v2 (shp-2).
- Etzioni, O., M. Cafarella, D. Downey, A.-M. Popescu, T. Shaked, S. Soderland, D. S. Weld, and A. Yates. 2005. Unsupervised named-entity extraction from the web: An experimental study. *Artificial Intelligence*, 165(1):91–134.
- Evans, N. 2000. Word classes in the world’s languages. In G. Booij, C. Lehmann, and J. Mugdan, eds, *Morphology: A Handbook on Inflection and Word Formation*, 708–732. Mouton.
- Fader, A., S. Soderland, and O. Etzioni. 2011. Identifying relations for open information extraction. *EMNLP*.
- Fan, A., S. Bhosale, H. Schwenk, Z. Ma, A. El-Kishky, S. Goyal,

- M. Baines, O. Celebi, G. Wenzek, V. Chaudhary, N. Goyal, T. Birch, V. Liptchinsky, S. Edunov, M. Auli, and A. Joulin. 2021. [Beyond english-centric multilingual machine translation](#). *JMLR*, 22(107):1–48.
- Fang, C. M., A. R. Liu, V. Danry, E. Lee, S. W. Chan, P. Pataranutaporn, P. Maes, J. Phang, M. Lampe, L. Ahmad, and S. Agarwal. 2025. [How AI and human behaviors shape psychosocial effects of extended chatbot use: A longitudinal randomized controlled study](#). ArXiv preprint.
- Fant, G. M. 1951. Speech communication research. *Ing. Vetenskaps Akad. Stockholm, Sweden*, 24:331–337.
- Fant, G. M. 1960. *Acoustic Theory of Speech Production*. Mouton.
- Fant, G. M. 1986. Glottal flow: Models and interaction. *Journal of Phonetics*, 14:393–399.
- Fant, G. M. 2004. *Speech Acoustics and Phonetics*. Kluwer.
- Fast, E., B. Chen, and M. S. Bernstein. 2016. Empath: Understanding Topic Signals in Large-Scale Text. *CHI*.
- Fauconnier, G. and M. Turner. 2008. *The way we think: Conceptual blending and the mind's hidden complexities*. Basic Books.
- Feldman, J. A. and D. H. Ballard. 1982. Connectionist models and their properties. *Cognitive Science*, 6:205–254.
- Fellbaum, C., ed. 1998. *WordNet: An Electronic Lexical Database*. MIT Press.
- Feng, S., B. M. Halpern, O. Kudina, and O. Scharenborg. 2024. [Towards inclusive automatic speech recognition](#). *Computer Speech & Language*, 84:101567.
- Feng, V. W. and G. Hirst. 2011. [Classifying arguments by scheme](#). *ACL*.
- Feng, V. W. and G. Hirst. 2014. [A linear-time bottom-up discourse parser with constraints and post-editing](#). *ACL*.
- Feng, V. W., Z. Lin, and G. Hirst. 2014. [The impact of deep hierarchical discourse structures in the evaluation of text coherence](#). *COLING*.
- Fernandes, E. R., C. N. dos Santos, and R. L. Milidiú. 2012. [Latent structure perceptron with feature induction for unrestricted coreference resolution](#). *CoNLL*.
- Ferragina, P. and U. Scaiella. 2011. Fast and accurate annotation of short texts with wikipedia pages. *IEEE Software*, 29(1):70–75.
- Ferro, L., L. Gerber, I. Mani, B. Sundheim, and G. Wilson. 2005. Tides 2005 standard for the annotation of temporal expressions. Technical report, MITRE.
- Ferrucci, D. A. 2012. Introduction to “This is Watson”. *IBM Journal of Research and Development*, 56(3/4):1:1–1:15.
- Field, A. and Y. Tsvetkov. 2019. [Entity-centric contextual affective analysis](#). *ACL*.
- Fillmore, C. J. 1966. A proposal concerning English prepositions. In F. P. Dinneen, ed., *17th annual Round Table*, volume 17 of *Monograph Series on Language and Linguistics*, 19–34. Georgetown University Press.
- Fillmore, C. J. 1968. The case for case. In E. W. Bach and R. T. Harms, eds, *Universals in Linguistic Theory*, 1–88. Holt, Rinehart & Winston.
- Fillmore, C. J. 1985. Frames and the semantics of understanding. *Quaderni di Semantica*, VI(2):222–254.
- Fillmore, C. J. 2003. Valency and semantic roles: the concept of deep structure case. In V. Agel, L. M. Eichinger, H. W. Eroms, P. Hellwig, H. J. Heringer, and H. Lobin, eds, *Dependenz und Valenz: Ein internationales Handbuch der zeitgenössischen Forschung*, chapter 36, 457–475. Walter de Gruyter.
- Fillmore, C. J. 2012. [ACL lifetime achievement award: Encounters with language](#). *Computational Linguistics*, 38(4):701–718.
- Fillmore, C. J. and C. F. Baker. 2009. A frames approach to semantic analysis. In B. Heine and H. Narrog, eds, *The Oxford Handbook of Linguistic Analysis*, 313–340. Oxford University Press.
- Fillmore, C. J., C. R. Johnson, and M. R. L. Petruck. 2003. Background to FrameNet. *International journal of lexicography*, 16(3):235–250.
- Finkelstein, L., E. Gabrilovich, Y. Matias, E. Rivlin, Z. Solan, G. Wolfman, and E. Ruppín. 2002. Placing search in context: The concept revisited. *ACM Transactions on Information Systems*, 20(1):116–131.
- Finlayson, M. A. 2016. Inferring Propp’s functions from semantically annotated text. *The Journal of American Folklore*, 129(511):55–77.
- Firth, J. R. 1957. A synopsis of linguistic theory 1930–1955. In *Studies in Linguistic Analysis*. Philological Society. Reprinted in Palmer, F. (ed.) 1968. *Selected Papers of J. R. Firth*. Longman, Harlow.
- Fitt, S. 2002. Unisyn lexicon. <http://www.cstr.ed.ac.uk/projects/unisyn/>.
- Flanagan, J. L. 1972. *Speech Analysis, Synthesis, and Perception*. Springer.
- Flanagan, J. L., K. Ishizaka, and K. L. Shipley. 1975. Synthesis of speech from a dynamic model of the vocal cords and vocal tract. *The Bell System Technical Journal*, 54(3):485–506.
- Foland, W. and J. H. Martin. 2016. [CU-NLP at SemEval-2016 task 8: AMR parsing using LSTM-based recurrent neural networks](#). *SemEval-2016*.
- Foland, Jr., W. R. and J. H. Martin. 2015. [Dependency-based semantic role labeling using convolutional neural networks](#). **SEM 2015*.
- Foltz, P. W., W. Kintsch, and T. K. Landauer. 1998. The measurement of textual coherence with latent semantic analysis. *Discourse processes*, 25(2-3):285–307.
- Fox, B. A. 1993. *Discourse Structure and Anaphora: Written and Conversational English*. Cambridge.
- Francis, W. N. and H. Kučera. 1982. *Frequency Analysis of English Usage*. Houghton Mifflin, Boston.
- Franz, A. and T. Brants. 2006. All our n-gram are belong to you. <https://research.google/blog/all-our-n-gram-are-belong-to-you/>.
- Friedman, B. and D. G. Hendry. 2019. *Value Sensitive Design: Shaping Technology with Moral Imagination*. MIT Press.
- Friedman, B., D. G. Hendry, and A. Borming. 2017. A survey of value sensitive design methods. *Foundations and Trends in Human-Computer Interaction*, 11(2):63–125.
- Fry, D. B. 1955. Duration and intensity as physical correlates of linguistic stress. *JASA*, 27:765–768.
- Fry, D. B. 1959. Theoretical aspects of mechanical speech recognition. *Journal of the British Institution of Radio Engineers*, 19(4):211–218. Appears together with companion paper (Denes 1959).
- Furnas, G. W., T. K. Landauer, L. M. Gomez, and S. T. Dumais. 1987. The vocabulary problem in human-system communication. *Communications of the ACM*, 30(11):964–971.
- Gabow, H. N., Z. Galil, T. Spencer, and R. E. Tarjan. 1986. [Efficient algorithms for finding minimum spanning trees in undirected and directed graphs](#). *Combinatorica*, 6(2):109–122.
- Gaddy, D., M. Stern, and D. Klein. 2018. [What’s going on in neural constituency parsers? an analysis](#). *NAACL HLT*.

- Gage, P. 1994. A new algorithm for data compression. *The C Users Journal*, 12(2):23–38.
- Gale, W. A. and K. W. Church. 1994. What is wrong with adding one? In N. Oostdijk and P. de Haan, eds, *Corpus-Based Research into Language*, 189–198. Rodopi.
- Gale, W. A. and K. W. Church. 1991. A program for aligning sentences in bilingual corpora. *ACL*.
- Gale, W. A. and K. W. Church. 1993. A program for aligning sentences in bilingual corpora. *Computational Linguistics*, 19:75–102.
- Gale, W. A., K. W. Church, and D. Yarowsky. 1992a. One sense per discourse. *HLL*.
- Gale, W. A., K. W. Church, and D. Yarowsky. 1992b. Work on statistical methods for word sense disambiguation. *AAAI Fall Symposium on Probabilistic Approaches to Natural Language*.
- Gao, T., H. Yen, J. Yu, and D. Chen. 2023. Enabling large language models to generate text with citations. *EMNLP*.
- Garg, N., L. Schiebinger, D. Jurafsky, and J. Zou. 2018. Word embeddings quantify 100 years of gender and ethnic stereotypes. *Proceedings of the National Academy of Sciences*, 115(16):E3635–E3644.
- Garside, R. 1987. The CLAWS word-tagging system. In R. Garside, G. Leech, and G. Sampson, eds, *The Computational Analysis of English*, 30–41. Longman.
- Garside, R., G. Leech, and A. McEnery. 1997. *Corpus Annotation*. Longman.
- Gebru, T., J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford. 2020. Datasheets for datasets. ArXiv.
- Gehman, S., S. Gururangan, M. Sap, Y. Choi, and N. A. Smith. 2020. RealToxicityPrompts: Evaluating neural toxic degeneration in language models. *Findings of EMNLP*.
- Gemmeke, J. F., D. P. Ellis, D. Freedman, A. Jansen, W. Lawrence, R. C. Moore, M. Plakal, and M. Ritter. 2017. Audio Set: An ontology and human-labeled dataset for audio events. *ICASSP*.
- Gerber, M. and J. Y. Chai. 2010. Beyond nombank: A study of implicit arguments for nominal predicates. *ACL*.
- Gerlich, M. 2025. AI tools in society: Impacts on cognitive offloading and the future of critical thinking. *Societies*, 15(1):6.
- Gers, F. A., J. Schmidhuber, and F. Cummins. 2000. Learning to forget: Continual prediction with lstm. *Neural computation*, 12(10):2451–2471.
- Geva, M., R. Schuster, J. Berant, and O. Levy. 2021. Transformer feed-forward layers are key-value memories. *EMNLP*.
- Gil, D. 2000. Syntactic categories, cross-linguistic variation and universal grammar. In P. M. Vogel and B. Comrie, eds, *Approaches to the Typology of Word Classes*, 173–216. Mouton.
- Gildea, D. and D. Jurafsky. 2000. Automatic labeling of semantic roles. *ACL*.
- Gildea, D. and D. Jurafsky. 2002. Automatic labeling of semantic roles. *Computational Linguistics*, 28(3):245–288.
- Gildea, D. and M. Palmer. 2002. The necessity of syntactic parsing for predicate argument recognition. *ACL*.
- Giles, C. L., G. M. Kuhn, and R. J. Williams. 1994. Dynamic recurrent neural networks: Theory and applications. *IEEE Trans. Neural Netw. Learning Syst.*, 5(2):153–156.
- Gillick, L. and S. J. Cox. 1989. Some statistical issues in the comparison of speech recognition algorithms. *ICASSP*.
- Girard, G. 1718. *La justesse de la langue française: ou les différentes significations des mots qui passent pour synonymes*. Laurent d’Houry, Paris.
- Giuliano, V. E. 1965. The interpretation of word associations. *Statistical Association Methods For Mechanized Documentation. Symposium Proceedings. Washington, D.C., USA, March 17, 1964*. <https://nvlpubs.nist.gov/nistpubs/Legacy/MP/nbsmiscellaneouspub269.pdf>.
- Gladkova, A., A. Drozd, and S. Mat-suoka. 2016. Analogy-based detection of morphological and semantic relations with word embeddings: what works and what doesn’t. *NAACL Student Research Workshop*.
- Glenberg, A. M. and D. A. Robertson. 2000. Symbol grounding and meaning: A comparison of high-dimensional and embodied theories of meaning. *Journal of memory and language*, 43(3):379–401.
- Godfrey, J., E. Holliman, and J. McDaniel. 1992. SWITCHBOARD: Telephone speech corpus for research and development. *ICASSP*.
- Goel, V. and W. Byrne. 2000. Minimum bayes-risk automatic speech recognition. *Computer Speech & Language*, 14(2):115–135.
- Goffman, E. 1974. *Frame analysis: An essay on the organization of experience*. Harvard University Press.
- Gonen, H. and Y. Goldberg. 2019. Lipstick on a pig: Debiasing methods cover up systematic gender biases in word embeddings but do not remove them. *NAACL HLT*.
- Goodfellow, I., Y. Bengio, and A. Courville. 2016. *Deep Learning*. MIT Press.
- Goodman, J. 2006. A bit of progress in language modeling: Extended version. Technical Report MSR-TR-2001-72, Machine Learning and Applied Statistics Group, Microsoft Research, Redmond, WA.
- Gould, S. J. 1980. *The Panda’s Thumb*. Penguin Group.
- Goyal, N., C. Gao, V. Chaudhary, P.-J. Chen, G. Wenzek, D. Ju, S. Krishnan, M. Ranzato, F. Guzmán, and A. Fan. 2022. The flores-101 evaluation benchmark for low-resource and multilingual machine translation. *TACL*, 10:522–538.
- Graff, D. 1997. The 1996 Broadcast News speech and language-model corpus. *Proceedings DARPA Speech Recognition Workshop*.
- Graves, A. 2012. Sequence transduction with recurrent neural networks. *ICASSP*.
- Graves, A. 2013. Generating sequences with recurrent neural networks. ArXiv.
- Graves, A., S. Fernández, F. Gomez, and J. Schmidhuber. 2006. Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. *ICML*.
- Graves, A., S. Fernández, M. Liwicki, H. Bunke, and J. Schmidhuber. 2007. Unconstrained on-line handwriting recognition with recurrent neural networks. *NeurIPS*.
- Graves, A. and N. Jaitly. 2014. Towards end-to-end speech recognition with recurrent neural networks. *ICML*.
- Graves, A., A.-r. Mohamed, and G. Hinton. 2013. Speech recognition with deep recurrent neural networks. *ICASSP*.
- Graves, A. and J. Schmidhuber. 2005. Framewise phoneme classification with bidirectional LSTM and other neural network architectures. *Neural Networks*, 18(5-6):602–610.
- Graves, A., G. Wayne, and I. Danihelka. 2014. Neural Turing machines. ArXiv.

- Gray, R. M. 1984. Vector quantization. *IEEE Transactions on ASSP*, ASSP-1(2):4–29.
- Green, B. F., A. K. Wolf, C. Chomsky, and K. Laughery. 1961. [Baseball: An automatic question answerer](#). *Proceedings of the Western Joint Computer Conference* 19.
- Greenberg, J. H. 1960. A quantitative approach to the morphological typology of language. *International journal of American linguistics*, 26(3):178–194.
- Greenberg, S., D. Ellis, and J. Hollenback. 1996. Insights into spoken language gleaned from phonetic transcription of the Switchboard corpus. *ICSLP*.
- Greene, B. B. and G. M. Rubin. 1971. Automatic grammatical tagging of English. Department of Linguistics, Brown University, Providence, Rhode Island.
- Greenwald, A. G., D. E. McGhee, and J. L. K. Schwartz. 1998. Measuring individual differences in implicit cognition: the implicit association test. *Journal of personality and social psychology*, 74(6):1464–1480.
- Grenager, T. and C. D. Manning. 2006. [Unsupervised discovery of a statistical verb lexicon](#). *EMNLP*.
- Grice, H. P. 1975. Logic and conversation. In P. Cole and J. L. Morgan, eds, *Speech Acts: Syntax and Semantics Volume 3*, 41–58. Academic Press.
- Grice, H. P. 1978. Further notes on logic and conversation. In P. Cole, ed., *Pragmatics: Syntax and Semantics Volume 9*, 113–127. Academic Press.
- Grishman, R. and B. Sundheim. 1995. [Design of the MUC-6 evaluation](#). *MUC-6*.
- Grosz, B. J. 1977a. The representation and use of focus in a system for understanding dialogs. *IJCAI-77*. Morgan Kaufmann.
- Grosz, B. J. 1977b. *The Representation and Use of Focus in Dialogue Understanding*. Ph.D. thesis, University of California, Berkeley.
- Grosz, B. J., A. K. Joshi, and S. Weinstein. 1983. [Providing a unified account of definite noun phrases in English](#). *ACL*.
- Grosz, B. J., A. K. Joshi, and S. Weinstein. 1995. [Centering: A framework for modeling the local coherence of discourse](#). *Computational Linguistics*, 21(2):203–225.
- Gruber, J. S. 1965. *Studies in Lexical Relations*. Ph.D. thesis, MIT.
- Grünwald, S., A. Friedrich, and J. Kuhn. 2021. [Applying Occam’s razor to transformer-based dependency parsing: What works, what doesn’t, and what is really necessary](#). *IWPT*.
- Guinaudeau, C. and M. Strube. 2013. [Graph-based local coherence modeling](#). *ACL*.
- Gundel, J. K., N. Hedberg, and R. Zacharski. 1993. Cognitive status and the form of referring expressions in discourse. *Language*, 69(2):274–307.
- Gurnee, W., N. Sofroniew, A. Pearce, M. Piotrowski, I. Kauvar, R. Chen, A. Soligo, P. Bogdan, E. Ong, R. Wang, T. B. Thompson, D. Abrahams, S. Kantamneni, E. Ameisen, J. Batson, and J. Lindsey. 2026. [Verbalizable representations form a global workspace in language models](#). Unpublished white paper.
- Gururangan, S., A. Marasović, S. Swayamdipta, K. Lo, I. Beltagy, D. Downey, and N. A. Smith. 2020. [Don’t stop pretraining: Adapt language models to domains and tasks](#). *ACL*.
- Gusfield, D. 1997. *Algorithms on Strings, Trees, and Sequences*. Cambridge University Press.
- Guyon, I. and A. Elisseeff. 2003. An introduction to variable and feature selection. *JMLR*, 3:1157–1182.
- Haber, J. and M. Poesio. 2020. [Assessing polyseme sense similarity through co-predication acceptability and contextualised embedding distance](#). **SEM*.
- Habernal, I. and I. Gurevych. 2016. [Which argument is more convincing? Analyzing and predicting convincingness of Web arguments using bidirectional LSTM](#). *ACL*.
- Habernal, I. and I. Gurevych. 2017. [Argumentation mining in user-generated web discourse](#). *Computational Linguistics*, 43(1):125–179.
- Haghighi, A. and D. Klein. 2009. [Simple coreference resolution with rich syntactic and semantic features](#). *EMNLP*.
- Hajishirzi, H., L. Zilles, D. S. Weld, and L. Zettlemoyer. 2013. [Joint coreference resolution and named-entity linking with multi-pass sieves](#). *EMNLP*.
- Hajič, J. 1998. *Building a Syntactically Annotated Corpus: The Prague Dependency Treebank*, 106–132. Karolinum.
- Hajič, J. 2000. [Morphological tagging: Data vs. dictionaries](#). *NAACL*.
- Hajič, J., M. Ciaramita, R. Johansson, D. Kawahara, M. A. Martí, L. Márquez, A. Meyers, J. Nivre, S. Padó, J. Štěpánek, P. Stranák, M. Surdeanu, N. Xue, and Y. Zhang. 2009. [The conll-2009 shared task: Syntactic and semantic dependencies in multiple languages](#). *CoNLL*.
- Hakkani-Tür, D., K. Oflazer, and G. Tür. 2002. Statistical morphological disambiguation for agglutinative languages. *Journal of Computers and Humanities*, 36(4):381–410.
- Halliday, M. A. K. and R. Hasan. 1976. *Cohesion in English*. Longman. English Language Series, Title No. 9.
- Hamilton, W. L., K. Clark, J. Leskovec, and D. Jurafsky. 2016a. [Inducing domain-specific sentiment lexicons from unlabeled corpora](#). *EMNLP*.
- Hamilton, W. L., J. Leskovec, and D. Jurafsky. 2016b. [Diachronic word embeddings reveal statistical laws of semantic change](#). *ACL*.
- Han, W., Y. Zhang, Z. Chen, B. Liu, M. Pechenizkiy, M. Fang, and Y. Zheng. 2026. [MuBench: Assessment of multilingual capabilities of large language models across 61 languages](#). *ACL Findings*.
- Hannun, A. 2017. Sequence modeling with CTC. *Distill*, 2(11).
- Hannun, A. Y., A. L. Maas, D. Jurafsky, and A. Y. Ng. 2014. First-pass large vocabulary continuous speech recognition using bi-directional recurrent DNNs. ArXiv preprint arXiv:1408.2873.
- Harris, C. M. 1953. A study of the building blocks in speech. *JASA*, 25(5):962–969.
- Harris, Z. S. 1946. From morpheme to utterance. *Language*, 22(3):161–183.
- Harris, Z. S. 1954. [Distributional structure](#). *Word*, 10:146–162.
- Harris, Z. S. 1962. *String Analysis of Sentence Structure*. Mouton, The Hague.
- Hashimoto, T., M. Srivastava, H. Namkoong, and P. Liang. 2018. Fairness without demographics in repeated loss minimization. *ICML*.
- Hastie, T., R. J. Tibshirani, and J. H. Friedman. 2001. *The Elements of Statistical Learning*. Springer.
- Hatzivassiloglou, V. and K. McKeown. 1997. [Predicting the semantic orientation of adjectives](#). *ACL*.
- Hatzivassiloglou, V. and J. Wiebe. 2000. [Effects of adjective orientation and gradability on sentence subjectivity](#). *COLING*.
- Haviland, S. E. and H. H. Clark. 1974. What’s new? Acquiring new information as a process in comprehension. *Journal of Verbal Learning and Verbal Behaviour*, 13:512–521.

- Hawkins, J. A. 1978. *Definiteness and indefiniteness: a study in reference and grammaticality prediction*. Croom Helm Ltd.
- Hayashi, T., R. Yamamoto, K. Inoue, T. Yoshimura, S. Watanabe, T. Toda, K. Takeda, Y. Zhang, and X. Tan. 2020. ESPnet-TTS: Unified, reproducible, and integratable open source end-to-end text-to-speech toolkit. *ICASSP*.
- He, L., K. Lee, M. Lewis, and L. Zettlemoyer. 2017. *Deep semantic role labeling: What works and what's next*. *ACL*.
- He, W., K. Liu, J. Liu, Y. Lyu, S. Zhao, X. Xiao, Y. Liu, Y. Wang, H. Wu, Q. She, X. Liu, T. Wu, and H. Wang. 2018. *DuReader: a Chinese machine reading comprehension dataset from real-world applications*. *Workshop on Machine Reading for Question Answering*.
- Heafield, K. 2011. *KenLM: Faster and smaller language model queries*. *Workshop on Statistical Machine Translation*.
- Heafield, K., I. Pouzyrevsky, J. H. Clark, and P. Koehn. 2013. *Scalable modified Kneser-Ney language model estimation*. *ACL*.
- Heaps, H. S. 1978. *Information retrieval. Computational and theoretical aspects*. Academic Press.
- Hearst, M. A. 1992a. *Automatic acquisition of hyponyms from large text corpora*. *COLING*.
- Hearst, M. A. 1992b. *Automatic acquisition of hyponyms from large text corpora*. *COLING*.
- Hearst, M. A. 1997. *Texttiling: Segmenting text into multi-paragraph subtopic passages*. *Computational Linguistics*, 23:33–64.
- Hearst, M. A. 1998. Automatic discovery of WordNet relations. In C. Fellbaum, ed., *WordNet: An Electronic Lexical Database*. MIT Press.
- Heim, I. 1982. *The semantics of definite and indefinite noun phrases*. Ph.D. thesis, University of Massachusetts at Amherst.
- Heinz, J. M. and K. N. Stevens. 1961. On the properties of voiceless fricative consonants. *JASA*, 33:589–596.
- Hellrich, J., S. Buechel, and U. Hahn. 2019. *Modeling word emotion in historical language: Quantity beats supposed stability in seed word selection*. *3rd Joint SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature*.
- Hellrich, J. and U. Hahn. 2016. *Bad company—Neighborhoods in neural embedding spaces considered harmful*. *COLING*.
- Henderson, J. 1994. *Description Based Parsing in a Connectionist Network*. Ph.D. thesis, University of Pennsylvania, Philadelphia, PA.
- Henderson, J. 2003. *Inducing history representations for broad coverage statistical parsing*. *HLT-NAACL-03*.
- Henderson, J. 2004. *Discriminative training of a neural network statistical parser*. *ACL*.
- Henderson, P., J. Hu, J. Romoff, E. Brunskill, D. Jurafsky, and J. Pineau. 2020. *Towards the systematic reporting of the energy and carbon footprints of machine learning*. *Journal of Machine Learning Research*, 21(248):1–43.
- Henderson, P., X. Li, D. Jurafsky, T. Hashimoto, M. A. Lemley, and P. Liang. 2023. *Foundation models and fair use*. *JMLR*, 24(400):1–79.
- Hendrickx, I., S. N. Kim, Z. Kozareva, P. Nakov, D. Ó Séaghdha, S. Padó, M. Pennacchiotti, L. Romano, and S. Szpakowicz. 2009. *Semeval-2010 task 8: Multi-way classification of semantic relations between pairs of nominals*. *5th International Workshop on Semantic Evaluation*.
- Hendrix, G. G., C. W. Thompson, and J. Slocum. 1973. Language processing via canonical verbs and semantic models. *Proceedings of IJCAI-73*.
- Hendrycks, D., C. Burns, S. Basart, A. Zou, D. Song, and J. Steinhardt. 2021. *Measuring massive multitask language understanding*. *ICLR*.
- Herdan, G. 1960. *Type-token mathematics*. Mouton.
- Hermann, K. M., T. Kocisky, E. Grefenstette, L. Espeholt, W. Kay, M. Suleyman, and P. Blunsom. 2015. Teaching machines to read and comprehend. *NeurIPS*.
- Hernault, H., H. Prendinger, D. A. duVerle, and M. Ishizuka. 2010. Hilda: A discourse parser using support vector machine classification. *Dialogue & Discourse*, 1(3).
- Hewitt, J. and P. Liang. 2019. *Designing and interpreting probes with control tasks*. *EMNLP*. Association for Computational Linguistics.
- Hewitt, J. and C. D. Manning. 2019. *A structural probe for finding syntax in word representations*. *NAACL HLT*.
- Hidey, C., E. Musi, A. Hwang, S. Muresan, and K. McKeown. 2017. Analyzing the semantic types of claims and premises in an online persuasive forum. *4th Workshop on Argument Mining*.
- Hill, F., R. Reichart, and A. Korhonen. 2015. *Simlex-999: Evaluating semantic models with (genuine) similarity estimation*. *Computational Linguistics*, 41(4):665–695.
- Hinton, G. E. 1986. Learning distributed representations of concepts. *COGSCI*.
- Hinton, G. E., S. Osindero, and Y.-W. Teh. 2006. A fast learning algorithm for deep belief nets. *Neural computation*, 18(7):1527–1554.
- Hinton, G. E., N. Srivastava, A. Krizhevsky, I. Sutskever, and R. R. Salakhutdinov. 2012. Improving neural networks by preventing co-adaptation of feature detectors. *ArXiv preprint arXiv:1207.0580*.
- Hirschman, L., M. Light, E. Breck, and J. D. Burger. 1999. *Deep Read: A reading comprehension system*. *ACL*.
- Hirst, G. 1981. *Anaphora in Natural Language Understanding: A survey*. Number 119 in Lecture notes in computer science. Springer-Verlag.
- Hirst, G. 1987. *Semantic Interpretation and the Resolution of Ambiguity*. Cambridge University Press.
- Hjelmslev, L. 1969. *Prologomena to a Theory of Language*. University of Wisconsin Press. Translated by Francis J. Whitfield; original Danish edition 1943.
- Hobbs, J. R. 1978. Resolving pronoun references. *Lingua*, 44:311–338.
- Hobbs, J. R. 1979. Coherence and coreference. *Cognitive Science*, 3:67–90.
- Hobbs, J. R., D. E. Appelt, J. Bear, D. Israel, M. Kameyama, M. E. Stickel, and M. Tyson. 1997. FAS-TUS: A cascaded finite-state transducer for extracting information from natural-language text. In E. Roche and Y. Schabes, eds, *Finite-State Language Processing*, 383–406. MIT Press.
- Hochreiter, S. and J. Schmidhuber. 1997. Long short-term memory. *Neural Computation*, 9(8):1735–1780.
- Hoffmann, J., S. Borgeaud, A. Mensch, E. Buchatskaya, T. Cai, E. Rutherford, D. de Las Casas, L. A. Hendricks, J. Welbl, A. Clark, T. Hennigan, E. Noland, K. Millican, G. van den Driessche, B. Damoc, A. Guy, S. Osindero, K. Simonyan, E. Elsen, J. W. Rae, O. Vinyals, and L. Sifre. 2022. *Training compute-optimal large language models*. *Advances in Neural Information Processing Systems*, volume 35.
- Hofmann, T. 1999. Probabilistic latent semantic indexing. *SIGIR-99*.
- Hofmann, V., P. R. Kalluri, D. Jurafsky, and S. King. 2024. Ai generates covertly racist decisions about people based on their dialect. *Nature*, 633(8028):147–154.

- Holtzman, A., J. Buys, L. Du, M. Forbes, and Y. Choi. 2020. [The curious case of neural text degeneration](#). *ICLR*.
- Hopcroft, J. E. and J. D. Ullman. 1979. *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley.
- Hou, Y., K. Markert, and M. Strube. 2018. [Unrestricted bridging resolution](#). *Computational Linguistics*, 44(2):237–284.
- Householder, F. W. 1995. Dionysius Thrax, the *technai*, and Sextus Empiricus. In E. F. K. Koerner and R. E. Asher, eds, *Concise History of the Language Sciences*, 99–103. Elsevier Science.
- Hovy, E. H. 1990. [Parsimonious and profligate approaches to the question of discourse structure relations](#). *Proceedings of the 5th International Workshop on Natural Language Generation*.
- Hovy, E. H., M. P. Marcus, M. Palmer, L. A. Ramshaw, and R. Weischedel. 2006. [OntoNotes: The 90% solution](#). *HLT-NAACL*.
- Hsu, W.-N., B. Bolte, Y.-H. H. Tsai, K. Lakhotia, R. Salakhutdinov, and A. Mohamed. 2021. Hubert: Self-supervised speech representation learning by masked prediction of hidden units. *IEEE/ACM TASLP*, 29:3451–3460.
- Hu, E. J., Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. 2022. [LoRA: Low-rank adaptation of large language models](#). *ICLR*.
- Hu, M. and B. Liu. 2004. Mining and summarizing customer reviews. *SIGKDD-04*.
- Huang, E. H., R. Socher, C. D. Manning, and A. Y. Ng. 2012. [Improving word representations via global context and multiple word prototypes](#). *ACL*.
- Huang, X., W. Zhu, H. Hu, C. He, L. Li, S. Huang, and F. Yuan. 2025. [BenchMAX: A comprehensive multilingual evaluation suite for large language models](#). *ACL Findings*.
- Huang, Z., W. Xu, and K. Yu. 2015. Bidirectional LSTM-CRF models for sequence tagging. *arXiv preprint arXiv:1508.01991*.
- Huffman, S. 1996. Learning information extraction patterns from examples. In S. Wertzner, E. Riloff, and G. Scheller, eds, *Connectionist, Statistical, and Symbolic Approaches to Learning Natural Language Processing*, 246–260. Springer.
- Hunt, A. J. and A. W. Black. 1996. Unit selection in a concatenative speech synthesis system using a large speech database. *ICASSP*.
- Hutchins, W. J. 1986. *Machine Translation: Past, Present, Future*. Ellis Horwood, Chichester, England.
- Hutchins, W. J. 1997. [From first conception to first demonstration: The nascent years of machine translation, 1947–1954. A chronology](#). *Machine Translation*, 12:192–252.
- Hutchins, W. J. and H. L. Somers. 1992. *An Introduction to Machine Translation*. Academic Press.
- Hutchinson, B., V. Prabhakaran, E. Denton, K. Webster, Y. Zhong, and S. Denuyl. 2020. [Social biases in NLP models as barriers for persons with disabilities](#). *ACL*.
- Hymes, D. 1974. Ways of speaking. In R. Bauman and J. Sherzer, eds, *Explorations in the ethnography of speaking*, 433–451. Cambridge University Press.
- Ibrahim, L. and M. Cheng. 2026. [Thinking beyond the anthropomorphic paradigm benefits LLM research](#). *ACL*.
- Iida, R., K. Inui, H. Takamura, and Y. Matsumoto. 2003. [Incorporating contextual cues in trainable models for coreference resolution](#). *EACL Workshop on The Computational Treatment of Anaphora*.
- Irsay, O. and C. Cardie. 2014. [Opinion mining with deep recurrent neural networks](#). *EMNLP*.
- Iscenko, Z., S. Strand, Y. Chen, G. Aimard, M. Codreanu, V. Sampathkumar, A. Imas, J. Jacobs, E. Muiruri, J. Mateos-Garcia, J. J. Ng, S. Javed, J. Martin, O. Ajmeri, D. Calin, A. Kim, F. C. Millet, and J. Manyika. 2026. [Google’s AI & economy ATLAS v1.0: Mapping Gemini usage in the economy](#). *ArXiv preprint*.
- ISO8601. 2004. Data elements and interchange formats—information interchange—representation of dates and times. Technical report, International Organization for Standards (ISO).
- Itakura, F. 1975. Minimum prediction residual principle applied to speech recognition. *IEEE Transactions on ASSP*, ASSP-32:67–72.
- Iter, D., K. Guu, L. Lansing, and D. Jurafsky. 2020. [Pretraining with contrastive sentence objectives improves discourse performance of language models](#). *ACL*.
- Iter, D., J. Yoon, and D. Jurafsky. 2018. [Automatic detection of incoherent speech for diagnosing schizophrenia](#). *Fifth Workshop on Computational Linguistics and Clinical Psychology*.
- Iyer, S., I. Konstas, A. Cheung, J. Krishnamurthy, and L. Zettlemoyer. 2017. [Learning a neural semantic parser from user feedback](#). *ACL*.
- Iyer, S., X. V. Lin, R. Pasunuru, T. Mihaylov, D. Simig, P. Yu, K. Shuster, T. Wang, Q. Liu, P. S. Koura, X. Li, B. O’Horo, G. Pereyra, J. Wang, C. Dewan, A. Celikyilmaz, L. Zettlemoyer, and V. Stoyanov. 2022. [Opt-impl: Scaling language model instruction meta learning through the lens of generalization](#). *ArXiv preprint*.
- Izacard, G., P. Lewis, M. Lomeli, L. Hosseini, F. Petroni, T. Schick, J. Dwivedi-Yu, A. Joulin, S. Riedel, and E. Grave. 2022. [Few-shot learning with retrieval augmented language models](#). *ArXiv preprint*.
- Jackendoff, R. 1983. *Semantics and Cognition*. MIT Press.
- Jacobs, P. S. and L. F. Rau. 1990. SCISOR: A system for extracting information from on-line news. *CACM*, 33(11):88–97.
- Jaech, A., G. Mulcaire, S. Hathi, M. Ostendorf, and N. A. Smith. 2016. [Hierarchical character-word models for language identification](#). *ACL Workshop on NLP for Social Media*.
- Jaitly, N., P. Nguyen, A. Senior, and V. Vanhoucke. 2012. Application of pretrained deep neural networks to large vocabulary speech recognition. *INTERSPEECH*.
- Jauhiainen, T., M. Lui, M. Zampieri, T. Baldwin, and K. Lindén. 2019. [Automatic language identification in texts: A survey](#). *JAIR*, 65(1):675–682.
- Jefferson, G. 1972. Side sequences. In D. Sudnow, ed., *Studies in social interaction*, 294–333. Free Press, New York.
- Jefferson, G. 1984. Notes on a systematic deployment of the acknowledged tokens ‘yeah’ and ‘mm hm’. *Papers in Linguistics*, 17(2):197–216.
- Jeffreys, H. 1948. *Theory of Probability*, 2nd edition. Clarendon Press. Section 3.23.
- Jekat, S., A. Klein, E. Maier, I. Maleck, M. Mast, and J. Quantz. 1995. Dialogue acts in verbmobil. *Verbmobil-Report*–65–95.
- Jelinek, F. 1969. A fast sequential decoding algorithm using a stack. *IBM Journal of Research and Development*, 13:675–685.
- Jelinek, F. 1990. Self-organized language modeling for speech recognition. In A. Waibel and K.-F. Lee, eds, *Readings in Speech Recognition*, 450–506. Morgan Kaufmann. Originally distributed as IBM technical report in 1985.

- Jelinek, F. and R. L. Mercer. 1980. Interpolated estimation of Markov source parameters from sparse data. In E. S. Gelsema and L. N. Kanal, eds, *Proceedings, Workshop on Pattern Recognition in Practice*, 381–397. North Holland.
- Jelinek, F., R. L. Mercer, and L. R. Bahl. 1975. Design of a linguistic statistical decoder for the recognition of continuous speech. *IEEE Transactions on Information Theory*, IT-21(3):250–256.
- Ji, H. and R. Grishman. 2011. [Knowledge base population: Successful approaches and challenges](#). *ACL*.
- Ji, H., R. Grishman, and H. T. Dang. 2010. Overview of the tac 2011 knowledge base population track. *TAC-11*.
- Ji, Y. and J. Eisenstein. 2014. [Representation learning for text-level discourse parsing](#). *ACL*.
- Ji, Y. and J. Eisenstein. 2015. [One vector is not enough: Entity-augmented distributed semantics for discourse relations](#). *TACL*, 3:329–344.
- Jia, R. and P. Liang. 2016. [Data recombination for neural semantic parsing](#). *ACL*.
- Jia, S., T. Meng, J. Zhao, and K.-W. Chang. 2020. [Mitigating gender bias amplification in distribution by posterior regularization](#). *ACL*.
- Jiang, C., B. Qi, X. Hong, D. Fu, Y. Cheng, F. Meng, M. Yu, B. Zhou, and J. Zhou. 2024. [On large language models’ hallucination with regard to known facts](#). *NAACL HLT*.
- Johnson, J., M. Douze, and H. Jégou. 2017. Billion-scale similarity search with GPUs. *ArXiv preprint arXiv:1702.08734*.
- Johnson, K. 2003. *Acoustic and Auditory Phonetics*, 2nd edition. Blackwell.
- Johnson, W. E. 1932. Probability: deductive and inductive problems (appendix to). *Mind*, 41(164):421–423.
- Johnson-Laird, P. N. 1983. *Mental Models*. Harvard University Press, Cambridge, MA.
- Jones, M. P. and J. H. Martin. 1997. [Contextual spelling correction using latent semantic analysis](#). *ANLP*.
- Jones, R., A. McCallum, K. Nigam, and E. Riloff. 1999. Bootstrapping for text learning tasks. *IJCAI-99 Workshop on Text Mining: Foundations, Techniques and Applications*.
- Jones, T. 2015. [Toward a description of African American Vernacular English dialect regions using “Black Twitter”](#). *American Speech*, 90(4):403–440.
- Joos, M. 1950. Description of language design. *JASA*, 22:701–708.
- Jordan, M. 1986. Serial order: A parallel distributed processing approach. Technical Report ICS Report 8604, University of California, San Diego.
- Joshi, A. K. and P. Hopely. 1999. A parser from antiquity. In A. Kornai, ed., *Extended Finite State Models of Language*, 6–15. Cambridge University Press.
- Joshi, A. K. and S. Kuhn. 1979. Centered logic: The role of entity centered sentence representation in natural language inferencing. *IJCAI-79*.
- Joshi, A. K. and S. Weinstein. 1981. Control of inference: Role of some aspects of discourse structure – centering. *IJCAI-81*.
- Joshi, M., D. Chen, Y. Liu, D. S. Weld, L. Zettlemoyer, and O. Levy. 2020. [SpanBERT: Improving pre-training by representing and predicting spans](#). *TACL*, 8:64–77.
- Joshi, M., O. Levy, D. S. Weld, and L. Zettlemoyer. 2019. [BERT for coreference resolution: Baselines and analysis](#). *EMNLP*.
- Joty, S., G. Carenini, and R. T. Ng. 2015. [CODRA: A novel discriminative framework for rhetorical analysis](#). *Computational Linguistics*, 41(3):385–435.
- Jurafsky, D. 2014. *The Language of Food*. W. W. Norton, New York.
- Jurafsky, D., V. Chahuneau, B. R. Routledge, and N. A. Smith. 2014. Narrative framing of consumer sentiment in online restaurant reviews. *First Monday*, 19(4).
- Jurafsky, D., C. Wooters, G. Tajchman, J. Segal, A. Stolcke, E. Fosler, and N. Morgan. 1994. The Berkeley restaurant project. *ICSLP*.
- Jurgens, D., S. M. Mohammad, P. Turney, and K. Holyoak. 2012. [SemEval-2012 task 2: Measuring degrees of relational similarity](#). **SEM 2012*.
- Jurgens, D., Y. Tsvetkov, and D. Jurafsky. 2017. [Incorporating dialectal variability for socially equitable language identification](#). *ACL*.
- Justeson, J. S. and S. M. Katz. 1991. [Co-occurrences of antonymous adjectives and their contexts](#). *Computational linguistics*, 17(1):1–19.
- Kalchbrenner, N. and P. Blunsom. 2013. [Recurrent continuous translation models](#). *EMNLP*.
- Kallini, J., I. Papadimitriou, R. Futrell, K. Mahowald, and C. Potts. 2024. [Mission: Impossible language models](#). *ACL*.
- Kameyama, M. 1986. [A property-sharing constraint in centering](#). *ACL*.
- Kamp, H. 1981. A theory of truth and semantic representation. In J. Groenendijk, T. Janssen, and M. Stokhof, eds, *Formal Methods in the Study of Language*, 189–222. Mathematical Centre, Amsterdam.
- Kamphuis, C., A. P. de Vries, L. Boytsov, and J. Lin. 2020. Which BM25 do you mean? a large-scale reproducibility study of scoring variants. *European Conference on Information Retrieval*.
- Kane, S. K., M. R. Morris, A. Paradiso, and J. Campbell. 2017. [“at times avuncular and cantankerous, with the reflexes of a mongoose”: Understanding self-expression through augmentative and alternative communication devices](#). *CSCW*.
- Kaplan, R. M. 1973. A general syntactic processor. In R. Rustin, ed., *Natural Language Processing*, 193–241. Algorithmics Press.
- Karamanis, N., M. Poesio, C. Mellish, and J. Oberlander. 2004. [Evaluating centering-based metrics of coherence for text structuring using a reliably annotated corpus](#). *ACL*.
- Karita, S., N. Chen, T. Hayashi, T. Hori, H. Inaguma, Z. Jiang, M. Someki, N. E. Y. Soplin, R. Yamamoto, X. Wang, S. Watanabe, T. Yoshimura, and W. Zhang. 2019. A comparative study on transformer vs RNN in speech applications. *IEEE ASRU-19*.
- Karlsson, F., A. Voutilainen, J. Heikkilä, and A. Anttila, eds. 1995. *Constraint Grammar: A Language-Independent System for Parsing Unrestricted Text*. Mouton de Gruyter.
- Karpukhin, V., B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih. 2020. [Dense passage retrieval for open-domain question answering](#). *EMNLP*.
- Karttunen, L. 1969. [Discourse referents](#). *COLING*. Preprint No. 70.
- Karttunen, L. 1999. Comments on Joshi. In A. Kornai, ed., *Extended Finite State Models of Language*, 16–18. Cambridge University Press.
- Kasami, T. 1965. An efficient recognition and syntax analysis algorithm for context-free languages. Technical Report AFCRL-65-758, Air Force Cambridge Research Laboratory, Bedford, MA.
- Katz, J. J. and J. A. Fodor. 1963. The structure of a semantic theory. *Language*, 39:170–210.
- Kay, M. 1967. [Experiments with a powerful parser](#). *COLING*.
- Kay, M. 1973. The MIND system. In R. Rustin, ed., *Natural Language Processing*, 155–188. Algorithmics Press.

- Kay, M. 1982. Algorithm schemata and data structures in syntactic processing. In S. Allén, ed., *Text Processing: Text Analysis and Generation, Text Typology and Attribution*, 327–358. Almqvist and Wiksell, Stockholm.
- Kay, M. and M. Röscheisen. 1988. Text-translation alignment. Technical Report P90-00143, Xerox Palo Alto Research Center, Palo Alto, CA.
- Kay, M. and M. Röscheisen. 1993. Text-translation alignment. *Computational Linguistics*, 19:121–142.
- Kehler, A. 1993. The effect of establishing coherence in ellipsis and anaphora resolution. *ACL*.
- Kehler, A. 1994. Temporal relations: Reference or discourse coherence? *ACL*.
- Kehler, A. 1997a. Current theories of centering for pronoun interpretation: A critical evaluation. *Computational Linguistics*, 23(3):467–475.
- Kehler, A. 1997b. Probabilistic coreference in information extraction. *EMNLP*.
- Kehler, A. 2000. *Coherence, Reference, and the Theory of Grammar*. CSLI Publications.
- Kehler, A., D. E. Appelt, L. Taylor, and A. Simma. 2004. The (non)utility of predicate-argument frequencies for pronoun interpretation. *HLT-NAACL*.
- Kehler, A. and H. Rohde. 2013. A probabilistic reconciliation of coherence-driven and centering-driven theories of pronoun interpretation. *Theoretical Linguistics*, 39(1-2):1–37.
- Keller, F. and M. Lapata. 2003. Using the web to obtain frequencies for unseen bigrams. *Computational Linguistics*, 29:459–484.
- Kendall, T. and C. Farrington. 2020. The Corpus of Regional African American Language. Version 2020.05. Eugene, OR: The Online Resources for African American Language Project. <http://oraal.uoregon.edu/coraal>.
- Kennedy, C. and B. K. Boguraev. 1996. Anaphora for everyone: Pronominal anaphora resolution without a parser. *COLING*.
- Khandelwal, U., O. Levy, D. Jurafsky, L. Zettlemoyer, and M. Lewis. 2019. Generalization through memorization: Nearest neighbor language models. *ICLR*.
- Khattab, O., C. Potts, and M. Zaharia. 2021. Relevance-guided supervision for OpenQA with ColBERT. *TACL*, 9:929–944.
- Khattab, O. and M. Zaharia. 2020. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. *SIGIR*.
- Kiela, D. and S. Clark. 2014. A systematic study of semantic vector space model parameters. *EACL 2nd Workshop on Continuous Vector Space Models and their Compositionality (CVSC)*.
- Kim, E. 2019. Optimize computational efficiency of skip-gram with negative sampling. https://aegis4048.github.io/optimize_computational_efficiency_of_skip-gram_with_negative_sampling.
- Kim, S. M. and E. H. Hovy. 2004. Determining the sentiment of opinions. *COLING*.
- King, S. 2020. From African American Vernacular English to African American Language: Rethinking the study of race and language in African Americans’ speech. *Annual Review of Linguistics*, 6:285–300.
- Kingma, D. and J. Ba. 2015. Adam: A method for stochastic optimization. *ICLR 2015*.
- Kintsch, W. and T. A. Van Dijk. 1978. Toward a model of text comprehension and production. *Psychological review*, 85(5):363–394.
- Kiperwasser, E. and Y. Goldberg. 2016. Simple and accurate dependency parsing using bidirectional LSTM feature representations. *TACL*, 4:313–327.
- Kipper, K., H. T. Dang, and M. Palmer. 2000. Class-based construction of a verb lexicon. *AAAI*.
- Kiritchenko, S. and S. M. Mohammad. 2017. Best-worst scaling more reliable than rating scales: A case study on sentiment intensity annotation. *ACL*.
- Kiritchenko, S. and S. M. Mohammad. 2018. Examining gender and race bias in two hundred sentiment analysis systems. **SEM*.
- Kirk, H. R., B. Vidgen, P. Röttger, and S. A. Hale. 2023. The empty signifier problem: Towards clearer paradigms for operationalising “alignment” in large language models. *Proceedings of the Socially Responsible Language Modelling Research Workshop (NeurIPS 2023)*.
- Kiss, T. and J. Strunk. 2006. Unsupervised multilingual sentence boundary detection. *Computational Linguistics*, 32(4):485–525.
- Kitaev, N., S. Cao, and D. Klein. 2019. Multilingual constituency parsing with self-attention and pre-training. *ACL*.
- Kitaev, N. and D. Klein. 2018. Constituency parsing with a self-attentive encoder. *ACL*.
- Klatt, D. H. 1975. Voice onset time, friction, and aspiration in word-initial consonant clusters. *Journal of Speech and Hearing Research*, 18:686–706.
- Klatt, D. H. 1977. Review of the ARPA speech understanding project. *JASA*, 62(6):1345–1366.
- Klatt, D. H. 1982. The Klattalk text-to-speech conversion system. *ICASSP*.
- Kleene, S. C. 1951. Representation of events in nerve nets and finite automata. Technical Report RM-704, RAND Corporation. RAND Research Memorandum.
- Kleene, S. C. 1956. Representation of events in nerve nets and finite automata. In C. Shannon and J. McCarthy, eds, *Automata Studies*, 3–41. Princeton University Press.
- Klein, S. and R. F. Simmons. 1963. A computational approach to grammatical coding of English words. *Journal of the ACM*, 10(3):334–347.
- Knott, A. and R. Dale. 1994. Using linguistic phenomena to motivate a set of coherence relations. *Discourse Processes*, 18(1):35–62.
- Kocijan, V., A.-M. Cretu, O.-M. Camburu, Y. Yordanov, and T. Lukasiewicz. 2019. A surprisingly robust trick for the Winograd Schema Challenge. *ACL*.
- Kocmi, T., C. Federmann, R. Grundkiewicz, M. Junczys-Dowmunt, H. Matsushita, and A. Menezes. 2021. To ship or not to ship: An extensive evaluation of automatic metrics for machine translation. *ArXiv*.
- Koehn, P. 2005. Europarl: A parallel corpus for statistical machine translation. *MT summit, vol. 5*.
- Koehn, P., H. Hoang, A. Birch, C. Callison-Burch, M. Federico, N. Bertoldi, B. Cowan, W. Shen, C. Moran, R. Zens, C. Dyer, O. Bojar, A. Constantin, and E. Herbst. 2006. Moses: Open source toolkit for statistical machine translation. *ACL*.
- Koehn, P., F. J. Och, and D. Marcu. 2003. Statistical phrase-based translation. *HLT-NAACL*.
- Koenecke, A., A. Nam, E. Lake, J. Nudell, M. Quartey, Z. Mengesha, C. Touns, J. R. Rickford, D. Jurafsky, and S. Goel. 2020. Racial disparities in automated speech recognition. *PNAS*, 117(14):7684–7689.
- Koenig, W., H. K. Dunn, and L. Y. Lacy. 1946. The sound spectrograph. *JASA*, 18:19–49.

- Kolhatkar, V., A. Roussel, S. Dipper, and H. Zinsmeister. 2018. [Anaphora with non-nominal antecedents in computational linguistics: A survey](#). *Computational Linguistics*, 44(3):547–612.
- Kreutzer, J., I. Caswell, L. Wang, A. Wahab, D. van Esch, N. Ulzii-Orshikh, A. Tapo, N. Subramani, A. Sokolov, C. Sikasote, M. Setyawan, S. Sarin, S. Samb, B. Sagot, C. Rivera, A. Rios, I. Papadimitriou, S. Osei, P. O. Suarez, I. Orife, K. Ogueji, A. N. Rubungo, T. Q. Nguyen, M. Müller, A. Müller, S. H. Muhammad, N. Muhammad, A. Mnyakeni, J. Mirzakhlov, T. Matangira, C. Leong, N. Lawson, S. Kudugunta, Y. Jernite, M. Jenny, O. Firat, B. F. P. Dossou, S. Dlamini, N. de Silva, S. Çabuk Balı, S. Biderman, A. Battisti, A. Barua, A. Bapna, P. Baljekar, I. A. Azime, A. Awokoya, D. Ataman, O. Ahia, O. Ahia, S. Agrawal, and M. Adeyemi. 2022. [Quality at a glance: An audit of web-crawled multilingual datasets](#). *TACL*, 10:50–72.
- Kruskal, J. B. 1983. An overview of sequence comparison. In D. Sankoff and J. B. Kruskal, eds, *Time Warps, String Edits, and Macromolecules: The Theory and Practice of Sequence Comparison*, 1–44. Addison-Wesley.
- Kudo, T. 2018. [Subword regularization: Improving neural network translation models with multiple subword candidates](#). *ACL*.
- Kudo, T. and Y. Matsumoto. 2002. [Japanese dependency analysis using cascaded chunking](#). *CoNLL*.
- Kudo, T. and J. Richardson. 2018. [SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing](#). *EMNLP*.
- Kullback, S. and R. A. Leibler. 1951. On information and sufficiency. *Annals of Mathematical Statistics*, 22:79–86.
- Kulmizev, A., M. de Lhoneux, J. Gontrum, E. Fano, and J. Nivre. 2019. [Deep contextualized word embeddings in transition-based and graph-based dependency parsing - a tale of two parsers revisited](#). *EMNLP*.
- Kumar, S. and W. Byrne. 2004. [Minimum Bayes-risk decoding for statistical machine translation](#). *HLT-NAACL*.
- Kummerfeld, J. K. and D. Klein. 2013. [Error-driven analysis of challenges in coreference resolution](#). *EMNLP*.
- Kuno, S. 1965. The predictive analyzer and a path elimination technique. *CACM*, 8(7):453–462.
- Kupiec, J. 1992. Robust part-of-speech tagging using a hidden Markov model. *Computer Speech and Language*, 6:225–242.
- Kurebito, M. 2017. Koryak. In M. Fortescue, M. Mithun, and N. Evans, eds, *Oxford Handbook of Polysynthesis*. Oxford.
- Kurita, K., N. Vyas, A. Pareek, A. W. Black, and Y. Tsvetkov. 2019. [Quantifying social biases in contextual word representations](#). *1st ACL Workshop on Gender Bias for Natural Language Processing*.
- Kwiatkowski, T., J. Palomaki, O. Redfield, M. Collins, A. Parikh, C. Alberti, D. Epstein, I. Polosukhin, J. Devlin, K. Lee, K. Toutanova, L. Jones, M. Kelcey, M.-W. Chang, A. M. Dai, J. Uszkoreit, Q. Le, and S. Petrov. 2019. [Natural questions: A benchmark for question answering research](#). *TACL*, 7:452–466.
- Ladefoged, P. 1993. *A Course in Phonetics*. Harcourt Brace Jovanovich. (3rd ed.).
- Ladefoged, P. 1996. *Elements of Acoustic Phonetics*, 2nd edition. University of Chicago.
- Lafferty, J. D., A. McCallum, and F. C. N. Pereira. 2001. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. *ICML*.
- Lai, A. and J. Tetreault. 2018. [Discourse coherence in the wild: A dataset, evaluation and methods](#). *SIGDIAL*.
- Lake, B. M. and G. L. Murphy. 2021. Word meaning in minds and machines. *Psychological Review*. In press.
- Lakoff, G. 1965. *On the Nature of Syntactic Irregularity*. Ph.D. thesis, Indiana University. Published as *Irregularity in Syntax*. Holt, Rinehart, and Winston, New York, 1970.
- Lakoff, G. 1972. Structural complexity in fairy tales. In *The Study of Man*, 128–50. School of Social Sciences, University of California, Irvine, CA.
- Lakoff, G. and M. Johnson. 1980. *Metaphors We Live By*. University of Chicago Press, Chicago, IL.
- Lambert, N., L. Tunstall, N. Rajani, and T. Thrush. 2023. [Huggingface h4 stack exchange preference dataset](#).
- Lample, G., M. Ballesteros, S. Subramanian, K. Kawakami, and C. Dyer. 2016. [Neural architectures for named entity recognition](#). *NAACL HLT*.
- Lample, G. and A. Conneau. 2019. Cross-lingual language model pre-training. *NeurIPS*, volume 32.
- Lan, Z., M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut. 2020. ALBERT: A lite BERT for self-supervised learning of language representations. *ICLR*.
- Landauer, T. K. and S. T. Dumais. 1997. A solution to Plato’s problem: The Latent Semantic Analysis theory of acquisition, induction, and representation of knowledge. *Psychological Review*, 104:211–240.
- Landauer, T. K., D. Laham, B. Rehder, and M. E. Schreiner. 1997. How well can passage meaning be derived without using word order? A comparison of Latent Semantic Analysis and humans. *COGSCI*.
- Lang, J. and M. Lapata. 2014. [Similarity-driven semantic role induction via graph partitioning](#). *Computational Linguistics*, 40(3):633–669.
- Lang, K. J., A. H. Waibel, and G. E. Hinton. 1990. A time-delay neural network architecture for isolated word recognition. *Neural networks*, 3(1):23–43.
- Lapata, M. 2003. [Probabilistic text structuring: Experiments with sentence ordering](#). *ACL*.
- Lapesa, G. and S. Evert. 2014. [A large scale evaluation of distributional semantic models: Parameters, interactions and model selection](#). *TACL*, 2:531–545.
- Lappin, S. and H. Leass. 1994. [An algorithm for pronominal anaphora resolution](#). *Computational Linguistics*, 20(4):535–561.
- Lascarides, A. and N. Asher. 1993. Temporal interpretation, discourse relations, and common sense entailment. *Linguistics and Philosophy*, 16(5):437–493.
- Lawrence, W. 1953. The synthesis of speech from signals which have a low information rate. In W. Jackson, ed., *Communication Theory*, 460–469. Butterworth.
- LDC. 1998. *LDC Catalog: Hub4 project*. University of Pennsylvania. www.ldc.upenn.edu/Catalog/LDC98S71.html.
- LeCun, Y., B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel. 1989. Backpropagation applied to handwritten zip code recognition. *Neural computation*, 1(4):541–551.
- Lee, D. D. and H. S. Seung. 1999. Learning the parts of objects by non-negative matrix factorization. *Nature*, 401(6755):788–791.
- Lee, H.-P., A. Sarkar, L. Tankelevitch, I. Drosos, S. Rintel, R. Banks, and

- N. Wilson. 2025. The impact of generative AI on critical thinking: Self-reported reductions in cognitive effort and confidence effects from a survey of knowledge workers. *CHI*.
- Lee, H., A. Chang, Y. Peirsman, N. Chambers, M. Surdeanu, and D. Jurafsky. 2013. Deterministic coreference resolution based on entity-centric, precision-ranked rules. *Computational Linguistics*, 39(4):885–916.
- Lee, H., Y. Peirsman, A. Chang, N. Chambers, M. Surdeanu, and D. Jurafsky. 2011. Stanford’s multi-pass sieve coreference resolution system at the CoNLL-2011 shared task. *CoNLL*.
- Lee, H., M. Surdeanu, and D. Jurafsky. 2017a. A scaffolding approach to coreference resolution integrating statistical and rule-based models. *Natural Language Engineering*, 23(5):733–762.
- Lee, K., M.-W. Chang, and K. Toutanova. 2019. Latent retrieval for weakly supervised open domain question answering. *ACL*.
- Lee, K., L. He, M. Lewis, and L. Zettlemoyer. 2017b. End-to-end neural coreference resolution. *EMNLP*.
- Lee, K., L. He, and L. Zettlemoyer. 2018. Higher-order coreference resolution with coarse-to-fine inference. *NAACL HLT*.
- Lehiste, I., ed. 1967. *Readings in Acoustic Phonetics*. MIT Press.
- Lehnert, W. G., C. Cardie, D. Fisher, E. Riloff, and R. Williams. 1991. Description of the CIRCUS system as used for MUC-3. *MUC-3*.
- Levenshtein, V. I. 1966. Binary codes capable of correcting deletions, insertions, and reversals. *Cybernetics and Control Theory*, 10(8):707–710. Original in *Doklady Akademii Nauk SSSR* 163(4): 845–848 (1965).
- Levesque, H. 2011. The Winograd Schema Challenge. *Logical Formalizations of Commonsense Reasoning — Papers from the AAAI 2011 Spring Symposium (SS-11-06)*.
- Levesque, H., E. Davis, and L. Morgenstern. 2012. The Winograd Schema Challenge. *KR-12*.
- Levin, B. 1977. Mapping sentences to case frames. Technical Report 167, MIT AI Laboratory. AI Working Paper 143.
- Levin, B. 1993. *English Verb Classes and Alternations: A Preliminary Investigation*. University of Chicago Press.
- Levin, B. and M. Rappaport Hovav. 2005. *Argument Realization*. Cambridge University Press.
- Levy, O. and Y. Goldberg. 2014a. Dependency-based word embeddings. *ACL*.
- Levy, O. and Y. Goldberg. 2014b. Linguistic regularities in sparse and explicit word representations. *CoNLL*.
- Levy, O. and Y. Goldberg. 2014c. Neural word embedding as implicit matrix factorization. *NeurIPS*.
- Levy, O., Y. Goldberg, and I. Dagan. 2015. Improving distributional similarity with lessons learned from word embeddings. *TACL*, 3:211–225.
- Li, A., F. Zheng, W. Byrne, P. Fung, T. Kamm, L. Yi, Z. Song, U. Ruhi, V. Venkataramani, and X. Chen. 2000. CASS: A phonetically transcribed corpus of Mandarin spontaneous speech. *ICSLP*.
- Li, B. Z., S. Min, S. Iyer, Y. Mehdad, and W.-t. Yih. 2020. Efficient one-pass end-to-end entity linking for questions. *EMNLP*.
- Li, J. and D. Jurafsky. 2017. Neural net models of open-domain discourse coherence. *EMNLP*.
- Li, J., R. Li, and E. H. Hovy. 2014. Recursive deep models for discourse parsing. *EMNLP*.
- Li, Q., T. Li, and B. Chang. 2016. Discourse parsing with attention-based hierarchical neural networks. *EMNLP*.
- Li, X., Y. Meng, X. Sun, Q. Han, A. Yuan, and J. Li. 2019. Is word segmentation necessary for deep learning of Chinese representations? *ACL*.
- Liang, P., R. Bommasani, T. Lee, D. Tsipras, D. Soylu, M. Yasunaga, Y. Zhang, D. Narayanan, Y. Wu, A. Kumar, B. Newman, B. Yuan, B. Yan, C. Zhang, C. Cosgrove, C. D. Manning, C. Ré, D. Acosta-Navas, D. A. Hudson, E. Zelikman, E. Durmus, F. Ladhak, F. Rong, H. Ren, H. Yao, J. Wang, K. Santhanam, L. Orr, L. Zheng, M. Yuksekgonul, M. Suzgun, N. Kim, N. Guha, N. Chatterji, O. Khattab, P. Henderson, Q. Huang, R. Chi, S. M. Xie, S. Santurkar, S. Ganguli, T. Hashimoto, T. Icard, T. Zhang, V. Chaudhary, W. Wang, X. Li, Y. Mai, Y. Zhang, and Y. Koreeda. 2023. Holistic evaluation of language models. *Transactions on Machine Learning Research*.
- Lieberman, A. M., P. C. Delattre, and F. S. Cooper. 1952. The role of selected stimulus variables in the perception of the unvoiced stop consonants. *American Journal of Psychology*, 65:497–516.
- Lin, D. 2003. Dependency-based evaluation of minipar. *Workshop on the Evaluation of Parsing Systems*.
- Lin, Y., J.-B. Michel, E. Aiden Lieberman, J. Orwant, W. Brockman, and S. Petrov. 2012a. Syntactic annotations for the Google books Ngram corpus. *ACL*.
- Lin, Y., J.-B. Michel, E. Lieberman Aiden, J. Orwant, W. Brockman, and S. Petrov. 2012b. Syntactic annotations for the Google Books Ngram corpus. *ACL*.
- Lin, Z., M.-Y. Kan, and H. T. Ng. 2009. Recognizing implicit discourse relations in the Penn Discourse Treebank. *EMNLP*.
- Lin, Z., H. T. Ng, and M.-Y. Kan. 2011. Automatically evaluating text coherence using discourse relations. *ACL*.
- Lin, Z., H. T. Ng, and M.-Y. Kan. 2014. A pdtb-styled end-to-end discourse parser. *Natural Language Engineering*, 20(2):151–184.
- Ling, W., C. Dyer, A. W. Black, I. Trancoso, R. Fernandez, S. Amir, L. Marujo, and T. Luís. 2015. Finding function in form: Compositional character models for open vocabulary word representation. *EMNLP*.
- Linzen, T. 2016. Issues in evaluating semantic spaces using word analogies. *1st Workshop on Evaluating Vector-Space Representations for NLP*.
- Lison, P. and J. Tiedemann. 2016. Opensubtitles2016: Extracting large parallel corpora from movie and tv subtitles. *LREC*.
- Litman, D. J. 1985. *Plan Recognition and Discourse Analysis: An Integrated Approach for Understanding Dialogues*. Ph.D. thesis, University of Rochester, Rochester, NY.
- Litman, D. J. and J. Allen. 1987. A plan recognition model for subdialogues in conversation. *Cognitive Science*, 11:163–200.
- Liu, A., J. Hayase, V. Hofmann, S. Oh, N. A. Smith, and Y. Choi. 2025. SuperBPE: Space travel for language models. ArXiv preprint.
- Liu, B. and L. Zhang. 2012. A survey of opinion mining and sentiment analysis. In C. C. Aggarwal and C. Zhai, eds, *Mining text data*, 415–464. Springer.
- Liu, G., B. Christian, T. Dumbalska, M. A. Bakker, and R. Dubey. 2026. AI assistance reduces persistence and hurts independent performance. ArXiv preprint.
- Liu, J., S. Min, L. Zettlemoyer, Y. Choi, and H. Hajishirzi. 2024. Infi-gram: Scaling unbounded n-gram language models to a trillion tokens. ArXiv preprint.
- Liu, Y., C. Sun, L. Lin, and X. Wang. 2016. Learning natural language inference using bidirectional LSTM model and inner-attention. ArXiv.

- Liu, Y., P. Fung, Y. Yang, C. Cieri, S. Huang, and D. Graff. 2006. HKUST/MTS: A very large scale Mandarin telephone speech corpus. *International Conference on Chinese Spoken Language Processing*.
- Liu, Y., M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov. 2019. [RoBERTa: A robustly optimized BERT pre-training approach](#). ArXiv preprint arXiv:1907.11692.
- Llama Team. 2024. [The llama 3 herd of models](#).
- Logeswaran, L., H. Lee, and D. Radev. 2018. Sentence ordering and coherence modeling using recurrent neural networks. *AAAI*.
- Longpre, S., L. Hou, T. Vu, A. Webson, H. W. Chung, Y. Tay, D. Zhou, Q. V. Le, B. Zoph, J. Wei, and A. Roberts. 2023. The Flan collection: Designing data and methods for effective instruction tuning. *ICML*.
- Longpre, S., G. Yauney, E. Reif, K. Lee, A. Roberts, B. Zoph, D. Zhou, J. Wei, K. Robinson, D. Mimno, and D. Ippolito. 2024. [A pretrainer’s guide to training data: Measuring the effects of data age, domain coverage, quality, & toxicity](#). *NAACL HLT*.
- Louis, A. and A. Nenkova. 2012. [A coherence model based on syntactic patterns](#). *EMNLP*.
- Loureiro, D. and A. Jorge. 2019. [Language modelling makes sense: Propagating representations through WordNet for full-coverage word sense disambiguation](#). *ACL*.
- Louviere, J. J., T. N. Flynn, and A. A. J. Marley. 2015. *Best-worst scaling: Theory, methods and applications*. Cambridge University Press.
- Lovins, J. B. 1968. Development of a stemming algorithm. *Mechanical Translation and Computational Linguistics*, 11(1–2):9–13.
- Lowerre, B. T. 1976. *The Harpy Speech Recognition System*. Ph.D. thesis, Carnegie Mellon University, Pittsburgh, PA.
- Lukasik, M., B. Dadachev, K. Papineni, and G. Simões. 2020. [Text segmentation by cross segment attention](#). *EMNLP*.
- Luo, X. 2005. [On coreference resolution performance metrics](#). *EMNLP*.
- Luo, X. and S. Pradhan. 2016. Evaluation metrics. In M. Poesio, R. Stuckardt, and Y. Versley, eds, *Anaphora resolution: Algorithms, resources, and applications*, 141–163. Springer.
- Luo, X., S. Pradhan, M. Recasens, and E. H. Hovy. 2014. [An extension of BLANC to system mentions](#). *ACL*.
- Ma, X. and E. H. Hovy. 2016. [End-to-end sequence labeling via bi-directional LSTM-CNNs-CRF](#). *ACL*.
- Maas, A., Z. Xie, D. Jurafsky, and A. Y. Ng. 2015. [Lexicon-free conversational speech recognition with neural networks](#). *NAACL HLT*.
- Maas, A. L., A. Y. Hannun, and A. Y. Ng. 2013. Rectifier nonlinearities improve neural network acoustic models. *ICML*.
- Maas, A. L., P. Qi, Z. Xie, A. Y. Hannun, C. T. Lengerich, D. Jurafsky, and A. Y. Ng. 2017. Building dnn acoustic models for large vocabulary speech recognition. *Computer Speech & Language*, 41:195–213.
- Magerman, D. M. 1995. [Statistical decision-tree models for parsing](#). *ACL*.
- Mairesse, F. and M. A. Walker. 2008. [Trainable generation of big-five personality styles through data-driven parameter estimation](#). *ACL*.
- Mann, W. C. and S. A. Thompson. 1987. Rhetorical structure theory: A theory of text organization. Technical Report RS-87-190, Information Sciences Institute.
- Manning, C. D. 2011. Part-of-speech tagging from 97% to 100%: Is it time for some linguistics? *CICLing 2011*.
- Manning, C. D., K. Clark, J. Hewitt, U. Khandelwal, and O. Levy. 2020. [Emergent linguistic structure in artificial neural networks trained by self-supervision](#). *PNAS*, 117(48):30046–30054.
- Manning, C. D., P. Raghavan, and H. Schütze. 2008. *Introduction to Information Retrieval*. Cambridge.
- Manning, C. D., M. Surdeanu, J. Bauer, J. Finkel, S. Bethard, and D. McClosky. 2014. [The Stanford CoreNLP natural language processing toolkit](#). *ACL*.
- Marcu, D. 1997. [The rhetorical parsing of natural language texts](#). *ACL*.
- Marcu, D. 1999. [A decision-based approach to rhetorical parsing](#). *ACL*.
- Marcu, D. 2000a. [The rhetorical parsing of unrestricted texts: A surface-based approach](#). *Computational Linguistics*, 26(3):395–448.
- Marcu, D., ed. 2000b. *The Theory and Practice of Discourse Parsing and Summarization*. MIT Press.
- Marcu, D. and A. Echihiabi. 2002. [An unsupervised approach to recognizing discourse relations](#). *ACL*.
- Marcu, D. and W. Wong. 2002. [A phrase-based, joint probability model for statistical machine translation](#). *EMNLP*.
- Marcus, M. P. 1980. *A Theory of Syntactic Recognition for Natural Language*. MIT Press.
- Marcus, M. P., B. Santorini, and M. A. Marcinkiewicz. 1993. [Building a large annotated corpus of English: The Penn treebank](#). *Computational Linguistics*, 19(2):313–330.
- Marie, B., A. Fujita, and R. Rubino. 2021. [Scientific credibility of machine translation research: A meta-evaluation of 769 papers](#). *ACL*.
- Markov, A. A. 1913. Essai d’une recherche statistique sur le texte du roman “Eugene Onegin” illustrant la liaison des épreuve en chain (‘Example of a statistical investigation of the text of “Eugene Onegin” illustrating the dependence between samples in chain’). *Izvestia Imperatorskoi Akademii Nauk (Bulletin de l’Académie Impériale des Sciences de St.-Petersbourg)*, 7:153–162.
- de Marneffe, M.-C., T. Dozat, N. Silveira, K. Haverinen, F. Ginter, J. Nivre, and C. D. Manning. 2014. [Universal Stanford dependencies: A cross-linguistic typology](#). *LREC*.
- de Marneffe, M.-C., B. MacCartney, and C. D. Manning. 2006. [Generating typed dependency parses from phrase structure parses](#). *LREC*.
- de Marneffe, M.-C. and C. D. Manning. 2008. [The Stanford typed dependencies representation](#). *COLING Workshop on Cross-Framework and Cross-Domain Parser Evaluation*.
- de Marneffe, M.-C., C. D. Manning, J. Nivre, and D. Zeman. 2021. [Universal Dependencies](#). *Computational Linguistics*, 47(2):255–308.
- de Marneffe, M.-C., M. Recasens, and C. Potts. 2015. Modeling the lifespan of discourse entities with application to coreference resolution. *JAIR*, 52:445–475.
- Màrquez, L., X. Carreras, K. C. Litkowski, and S. Stevenson. 2008. [Semantic role labeling: An introduction to the special issue](#). *Computational linguistics*, 34(2):145–159.
- Marshall, I. 1983. Choice of grammatical word-class without global syntactic analysis: Tagging words in the LOB corpus. *Computers and the Humanities*, 17:139–150.
- Marshall, I. 1987. Tag selection using probabilistic methods. In R. Garside, G. Leech, and G. Sampson, eds, *The Computational Analysis of English*, 42–56. Longman.
- Martschat, S. and M. Strube. 2014. [Recall error analysis for coreference resolution](#). *EMNLP*.

- Martschat, S. and M. Strube. 2015. [Latent structures for coreference resolution](#). *TACL*, 3:405–418.
- Mathis, D. A. and M. C. Mozer. 1995. [On the computational utility of consciousness](#). *NeurIPS*. MIT Press.
- McCallum, A., D. Freitag, and F. C. N. Pereira. 2000. Maximum entropy Markov models for information extraction and segmentation. *ICML*.
- McCallum, A. and W. Li. 2003. [Early results for named entity recognition with conditional random fields, feature induction and web-enhanced lexicons](#). *CoNLL*.
- McCarthy, J. F. and W. G. Lehnert. 1995. Using decision trees for coreference resolution. *IJCAI-95*.
- McClelland, J. L. and J. L. Elman. 1986. The TRACE model of speech perception. *Cognitive Psychology*, 18:1–86.
- McClelland, J. L. and D. E. Rumelhart, eds. 1986. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, volume 2: *Psychological and Biological Models*. MIT Press.
- McCreery, Z. and M. Naser. 2026. [LLM-assisted writing in engineering is associated with higher citations as evidenced from 1.17 million papers](#). *Journal of Informetrics*, 20(3):101842.
- McCulloch, W. S. and W. Pitts. 1943. A logical calculus of ideas immanent in nervous activity. *Bulletin of Mathematical Biophysics*, 5:115–133.
- McDonald, R., K. Crammer, and F. C. N. Pereira. 2005a. [Online large-margin training of dependency parsers](#). *ACL*.
- McDonald, R. and J. Nivre. 2011. [Analyzing and integrating dependency parsers](#). *Computational Linguistics*, 37(1):197–230.
- McDonald, R., F. C. N. Pereira, K. Ribarov, and J. Hajič. 2005b. [Non-projective dependency parsing using spanning tree algorithms](#). *HLT-EMNLP*.
- McLuhan, M. 1964. *Understanding Media: The Extensions of Man*. New American Library.
- Melamud, O., J. Goldberger, and I. Dagan. 2016. [context2vec: Learning generic context embedding with bidirectional LSTM](#). *CoNLL*.
- Meng, K., D. Bau, A. Andonian, and Y. Belinkov. 2022. [Locating and editing factual associations in GPT](#). *NeurIPS*, volume 36.
- Mengesha, Z., C. Heldreth, M. Lahav, J. Sublewski, and E. Tuennenman. 2021. [“I don’t think these devices are very culturally sensitive.”—Impact of automated speech recognition errors on African Americans](#). *Frontiers in Artificial Intelligence*, 4:725911.
- Merialdo, B. 1994. [Tagging English text with a probabilistic model](#). *Computational Linguistics*, 20(2):155–172.
- Mesgar, M. and M. Strube. 2016. [Lexical coherence graph modeling using word embeddings](#). *ACL*.
- Meyers, A., R. Reeves, C. Macleod, R. Szekely, V. Zielinska, B. Young, and R. Grishman. 2004. [The nombank project: An interim report](#). *NAACL/HLT Workshop: Frontiers in Corpus Annotation*.
- Mihalcea, R. and A. Csomai. 2007. Wikify!: Linking documents to encyclopedic knowledge. *CIKM 2007*.
- Mikheev, A., M. Moens, and C. Grover. 1999. [Named entity recognition without gazetteers](#). *EACL*.
- Mikolov, T. 2012. *Statistical language models based on neural networks*. Ph.D. thesis, Brno University of Technology.
- Mikolov, T., K. Chen, G. S. Corrado, and J. Dean. 2013a. Efficient estimation of word representations in vector space. *ICLR 2013*.
- Mikolov, T., M. Karafiát, L. Burget, J. Černocký, and S. Khudanpur. 2010. [Recurrent neural network based language model](#). *INTERSPEECH*.
- Mikolov, T., S. Kombrink, L. Burget, J. H. Černocký, and S. Khudanpur. 2011. Extensions of recurrent neural network language model. *ICASSP*.
- Mikolov, T., I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. 2013b. [Distributed representations of words and phrases and their compositionality](#). *NeurIPS*.
- Mikolov, T., W.-t. Yih, and G. Zweig. 2013c. [Linguistic regularities in continuous space word representations](#). *NAACL HLT*.
- Miller, G. A. and P. E. Nicely. 1955. An analysis of perceptual confusions among some English consonants. *JASA*, 27:338–352.
- Miller, G. A. and J. G. Beebe-Center. 1956. [Some psychological methods for evaluating the quality of translations](#). *Mechanical Translation*, 3:73–80.
- Miller, G. A. and W. G. Charles. 1991. Contextual correlates of semantics similarity. *Language and Cognitive Processes*, 6(1):1–28.
- Miller, G. A. and N. Chomsky. 1963. Finitary models of language users. In R. D. Luce, R. R. Bush, and E. Galanter, eds, *Handbook of Mathematical Psychology*, volume II, 419–491. John Wiley.
- Miller, G. A. and J. A. Selfridge. 1950. Verbal context and the recall of meaningful material. *American Journal of Psychology*, 63:176–185.
- Milne, D. and I. H. Witten. 2008. Learning to link with wikipedia. *CIKM 2008*.
- Miltsakaki, E., R. Prasad, A. K. Joshi, and B. L. Webber. 2004. [The Penn Discourse Treebank](#). *LREC*.
- Min, S., X. Lyu, A. Holtzman, M. Artetxe, M. Lewis, H. Hajishirzi, and L. Zettlemoyer. 2022. [Rethinking the role of demonstrations: What makes in-context learning work?](#) *EMNLP*.
- Minsky, M. 1974. A framework for representing knowledge. Technical Report 306, MIT AI Laboratory. Memo 306.
- Minsky, M. and S. Papert. 1969. *Perceptrons*. MIT Press.
- Mintz, M., S. Bills, R. Snow, and D. Jurafsky. 2009. [Distant supervision for relation extraction without labeled data](#). *ACL IJCNLP*.
- Mirza, P. and S. Tonelli. 2016. [CATENA: Causal and Temporal relation extraction from Natural language texts](#). *COLING*.
- Mishra, S., D. Khashabi, C. Baral, and H. Hajishirzi. 2022. Cross-task generalization via natural language crowdsourcing instructions. *ACL*.
- Mitchell, M., S. Wu, A. Zaldívar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru. 2019. [Model cards for model reporting](#). *ACM FAccT*.
- Mitkov, R. 2002. *Anaphora Resolution*. Longman.
- Mohamed, A., G. E. Dahl, and G. E. Hinton. 2009. Deep Belief Networks for phone recognition. *NIPS Workshop on Deep Learning for Speech Recognition and Related Applications*.
- Mohamed, A., M. Assi, and M. Guizani. 2026. [The impact of llm-assistants on software developer productivity: A systematic review and mapping study](#). *ACM Trans. Softw. Eng. Methodol.*
- Mohammad, S. M. 2018a. [Obtaining reliable human ratings of valence, arousal, and dominance for 20,000 English words](#). *ACL*.
- Mohammad, S. M. 2018b. [Word affect intensities](#). *LREC*.
- Mohammad, S. M. and P. D. Turney. 2013. Crowdsourcing a word-emotion association lexicon. *Computational Intelligence*, 29(3):436–465.

- Monroe, B. L., M. P. Colaresi, and K. M. Quinn. 2008. Fightin' words: Lexical feature selection and evaluation for identifying the content of political conflict. *Political Analysis*, 16(4):372–403.
- Moors, A., P. C. Ellsworth, K. R. Scherer, and N. H. Frijda. 2013. Appraisal theories of emotion: State of the art and future development. *Emotion Review*, 5(2):119–124.
- Moosavi, N. S. and M. Strube. 2016. Which coreference evaluation metric do you trust? A proposal for a link-based entity aware metric. *ACL*.
- Morey, M., P. Muller, and N. Asher. 2017. How much progress have we made on RST discourse parsing? a replication study of recent results on the rst-dt. *EMNLP*.
- Morgan, A. A., L. Hirschman, M. Colosimo, A. S. Yeh, and J. B. Colombe. 2004. Gene name identification and normalization using a model organism database. *Journal of Biomedical Informatics*, 37(6):396–410.
- Morgan, N. and H. Bourlard. 1990. Continuous speech recognition using multilayer perceptrons with hidden markov models. *ICASSP*.
- Morgan, N. and H. A. Bourlard. 1995. Neural networks for statistical recognition of continuous speech. *Proceedings of the IEEE*, 83(5):742–772.
- Morris, J. and G. Hirst. 1991. Lexical cohesion computed by thesaural relations as an indicator of the structure of text. *Computational Linguistics*, 17(1):21–48.
- Mousavi, P., G. Maimon, A. Moumen, D. Petermann, J. Shi, H. Wu, H. Yang, A. Kuznetsova, A. Ploujnikov, R. Marxer, B. Ramabhadran, B. Elizalde, L. Lugosch, J. Li, C. Subakan, P. Woodland, M. Kim, H. yi Lee, S. Watanabe, Y. Adi, and M. Ravanelli. 2025. Discrete audio tokens: More than a survey! ArXiv preprint.
- Muller, P., C. Braud, and M. Morey. 2019. ToNy: Contextual embeddings for accurate multilingual discourse segmentation of full documents. *Workshop on Discourse Relation Parsing and Treebanking*.
- Murphy, K. P. 2012. *Machine learning: A probabilistic perspective*. MIT Press.
- Musi, E., M. Stede, L. Kriesche, S. Muresan, and A. Rocci. 2018. A multi-layer annotated corpus of argumentative text: From argument schemes to discourse relations. *LREC*.
- Myers, G. 1992. "In this paper we report...": Speech acts and scientific facts. *Journal of Pragmatics*, 17(4):295–313.
- Nádas, A. 1984. Estimation of probabilities in the language model of the IBM speech recognition system. *IEEE Transactions on ASSP*, 32(4):859–861.
- Nadeem, M., A. Bethke, and S. Reddy. 2021. StereoSet: Measuring stereotypical bias in pretrained language models. *ACL*.
- Namvarpour, M., B. Brofsky, J. Y. Medina, M. Akter, and A. Razi. 2026. Understanding teen overreliance on AI companion chatbots through self-reported Reddit narratives. *CHI 2026*.
- Nash-Webber, B. L. 1975. The role of semantics in automatic speech understanding. In D. G. Bobrow and A. Collins, eds, *Representation and Understanding*, 351–382. Academic Press.
- Naur, P., J. W. Backus, F. L. Bauer, J. Green, C. Katz, J. McCarthy, A. J. Perlis, H. Rutishauser, K. Samelson, B. Vauquois, J. H. Wegstein, A. van Wijnagaarden, and M. Woodger. 1960. Report on the algorithmic language ALGOL 60. *CACM*, 3(5):299–314. Revised in *CACM* 6:1, 1-17, 1963.
- Nekoto, W., V. Marivate, T. Matsila, T. Fasubaa, T. Kolawole, T. Fagbohunge, S. O. Akinola, S. H. Muhammad, S. Kabongo, S. Osei, S. Freshia, R. A. Niyongabo, R. M. P. Ogayo, O. Ahia, M. Meressa, M. Adeyemi, M. Mokgesi-Seling, L. Okegbemi, L. J. Martinus, K. Tajudeen, K. Degila, K. Ogueji, K. Siminyu, J. Kreutzer, J. Webster, J. T. Ali, J. A. I. Orife, I. Ezeani, I. A. Dangana, H. Kamper, H. Elshahar, G. Duru, G. Kioko, E. Murhabazi, E. van Biljon, D. Whitenack, C. Onye-fuluchi, C. Emezue, B. Dossou, B. Sibanda, B. I. Bassey, A. Olabiyi, A. Ramkilowan, A. Öktem, A. Akinfaderin, and A. Bashir. 2020. Participatory research for low-resourced machine translation: A case study in African languages. *Findings of EMNLP*. The authors use the forall symbol to represent the whole Masakhane community.
- Ng, H. T., L. H. Teo, and J. L. P. Kwan. 2000. A machine learning approach to answering questions for reading comprehension tests. *EMNLP*.
- Ng, V. 2004. Learning noun phrase anaphoricity to improve coreference resolution: Issues in representation and optimization. *ACL*.
- Ng, V. 2005a. Machine learning for coreference resolution: From local classification to global ranking. *ACL*.
- Ng, V. 2005b. Supervised ranking for pronoun resolution: Some recent improvements. *AAAI*.
- Ng, V. 2010. Supervised noun phrase coreference research: The first fifteen years. *ACL*.
- Ng, V. 2017. Machine learning for entity coreference resolution: A retrospective look at two decades of research. *AAAI*.
- Ng, V. and C. Cardie. 2002a. Identifying anaphoric and non-anaphoric noun phrases to improve coreference resolution. *COLING*.
- Ng, V. and C. Cardie. 2002b. Improving machine learning approaches to coreference resolution. *ACL*.
- Nguyen, D. T. and S. Joty. 2017. A neural local coherence model. *ACL*.
- Nickerson, R. S. 1976. On conversational interaction with computers. *Proceedings of the ACM/SIGGRAPH workshop on User-oriented design of interactive graphics systems*.
- Nie, A., E. Bennett, and N. Goodman. 2019. DisSent: Learning sentence representations from explicit discourse relations. *ACL*.
- Nielsen, M. A. 2015. *Neural networks and Deep learning*. Determination Press USA.
- Nigam, K., J. D. Lafferty, and A. McCallum. 1999. Using maximum entropy for text classification. *IJCAI-99 workshop on machine learning for information filtering*.
- Nirenburg, S., H. L. Somers, and Y. Wilks, eds. 2002. *Readings in Machine Translation*. MIT Press.
- Nissim, M., S. Dingare, J. Carletta, and M. Steedman. 2004. An annotation scheme for information status in dialogue. *LREC*.
- NIST. 1990. TIMIT Acoustic-Phonetic Continuous Speech Corpus. National Institute of Standards and Technology Speech Disc 1-1.1. NIST Order No. PB91-505065.
- NIST. 2005. Speech recognition scoring toolkit (sctk) version 2.1. <http://www.nist.gov/speech/tools/>.
- NIST. 2007. Matched Pairs Sentence-Segment Word Error (MAPSSWE) Test.
- Nivre, J. 2007. Incremental non-projective dependency parsing. *NAACL-HLT*.
- Nivre, J. 2003. An efficient algorithm for projective dependency parsing. *IWPT-03*.
- Nivre, J. 2006. *Inductive Dependency Parsing*. Springer.
- Nivre, J. 2009. Non-projective dependency parsing in expected linear time. *ACL IJCNLP*.

- Nivre, J., J. Hall, S. Kübler, R. McDonald, J. Nilsson, S. Riedel, and D. Yuret. 2007a. [The conll 2007 shared task on dependency parsing](#). *EMNLP/CoNLL*.
- Nivre, J., J. Hall, J. Nilsson, A. Chanev, G. Eryigit, S. Kübler, S. Marinov, and E. Marsi. 2007b. Malt-parser: A language-independent system for data-driven dependency parsing. *Natural Language Engineering*, 13(02):95–135.
- Nivre, J. and J. Nilsson. 2005. [Pseudo-projective dependency parsing](#). *ACL*.
- Nivre, J. and M. Scholz. 2004. [Deterministic dependency parsing of english text](#). *COLING*.
- Noreen, E. W. 1989. *Computer Intensive Methods for Testing Hypothesis*. Wiley.
- Norman, D. A. 1988. *The Design of Everyday Things*. Basic Books.
- Norvig, P. 1991. [Techniques for automatic memoization with applications to context-free parsing](#). *Computational Linguistics*, 17(1):91–98.
- Nosek, B. A., M. R. Banaji, and A. G. Greenwald. 2002a. Harvesting implicit group attitudes and beliefs from a demonstration web site. *Group Dynamics: Theory, Research, and Practice*, 6(1):101.
- Nosek, B. A., M. R. Banaji, and A. G. Greenwald. 2002b. Math=male, me=female, therefore math≠ me. *Journal of personality and social psychology*, 83(1):44.
- Nostalgebraist. 2020. [Interpreting gpt: the logit lens](#). White paper.
- Ocal, M., A. Perez, A. Radas, and M. Finlayson. 2022. [Holistic evaluation of automatic TimeML annotators](#). *REC*.
- Och, F. J. 1998. *Ein beispieles-basierter und statistischer Ansatz zum maschinellen Lernen von natürlichsprachlicher Übersetzung*. Ph.D. thesis, Universität Erlangen-Nürnberg, Germany. Diplomarbeit (diploma thesis).
- Och, F. J. 2003. [Minimum error rate training in statistical machine translation](#). *ACL*.
- Och, F. J. and H. Ney. 2002. [Discriminative training and maximum entropy models for statistical machine translation](#). *ACL*.
- Och, F. J. and H. Ney. 2003. [A systematic comparison of various statistical alignment models](#). *Computational Linguistics*, 29(1):19–51.
- Och, F. J. and H. Ney. 2004. [The alignment template approach to statistical machine translation](#). *Computational Linguistics*, 30(4):417–449.
- Olive, J. P. 1977. Rule synthesis of speech from dyadic units. *ICASSP77*.
- Olson, K. D., D. Meeker, M. Troup, T. D. Barker, V. H. Nguyen, J. B. Manders, C. D. Stults, V. G. Jones, S. D. Shah, T. Shah, and L. H. Schwamm. 2025. [Use of ambient ai scribes to reduce administrative burden and professional burnout](#). *JAMA Network Open*, 8(10):e2534976.
- Olsson, C., N. Elhage, N. Nanda, N. Joseph, N. DasSarma, T. Henighan, B. Mann, A. Askell, Y. Bai, A. Chen, et al. 2022. [In-context learning and induction heads](#). ArXiv preprint.
- OpenAI. 2022. [Aligning language models to follow instructions](#). Blog post, January 27, 2022.
- Oppenheim, A. V., R. W. Schaffer, and T. G. J. Stockham. 1968. Nonlinear filtering of multiplexed and convolved signals. *Proceedings of the IEEE*, 56(8):1264–1291.
- Oravecz, C. and P. Dienes. 2002. [Efficient stochastic part-of-speech tagging for Hungarian](#). *REC*.
- Osgood, C. E., G. J. Suci, and P. H. Tannenbaum. 1957. *The Measurement of Meaning*. University of Illinois Press.
- Ostendorf, M., P. Price, and S. Shattuck-Hufnagel. 1995. The Boston University Radio News Corpus. Technical Report ECS-95-001, Boston University.
- Ouyang, L., J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelton, L. Miller, M. Simens, A. Askell, P. Welinder, P. Christiano, J. Leike, and R. Lowe. 2022. Training language models to follow instructions with human feedback. *NeurIPS*, volume 35.
- Packard, D. W. 1973. [Computer-assisted morphological analysis of ancient Greek](#). *COLING*.
- Palmer, M., D. Gildea, and N. Xue. 2010. Semantic role labeling. *Synthesis Lectures on Human Language Technologies*, 3(1):1–103.
- Palmer, M., P. Kingsbury, and D. Gildea. 2005. [The proposition bank: An annotated corpus of semantic roles](#). *Computational Linguistics*, 31(1):71–106.
- Panayotov, V., G. Chen, D. Povey, and S. Khudanpur. 2015. Librispeech: an ASR corpus based on public domain audio books. *ICASSP*.
- Pang, B. and L. Lee. 2008. Opinion mining and sentiment analysis. *Foundations and trends in information retrieval*, 2(1-2):1–135.
- Pang, B., L. Lee, and S. Vaithyanathan. 2002. [Thumbs up? Sentiment classification using machine learning techniques](#). *EMNLP*.
- Papadimitriou, I. and D. Jurafsky. 2023. [Injecting structural hints: Using language models to study inductive biases in language learning](#). *Findings of EMNLP*.
- Papadimitriou, I., K. Lopez, and D. Jurafsky. 2023. [Multilingual BERT has an accent: Evaluating English influences on fluency in multilingual models](#). *EACL Findings*.
- Papineni, K., S. Roukos, T. Ward, and W.-J. Zhu. 2002. [Bleu: A method for automatic evaluation of machine translation](#). *ACL*.
- Park, J. H., J. Shin, and P. Fung. 2018. [Reducing gender bias in abusive language detection](#). *EMNLP*.
- Park, J. and C. Cardie. 2014. [Identifying appropriate support for propositions in online user comments](#). *First workshop on argumentation mining*.
- Parrish, A., A. Chen, N. Nangia, V. Padmakumar, J. Phang, J. Thompson, P. M. Htut, and S. Bowman. 2022. [BBQ: A hand-built bias benchmark for question answering](#). *Findings of ACL 2022*.
- Paszke, A., S. Gross, S. Chintala, G. Chanan, E. Yang, Z. DeVito, Z. Lin, A. Desmaison, L. Antiga, and A. Lerer. 2017. Automatic differentiation in pytorch. *NIPS-W*.
- Peldszus, A. and M. Stede. 2013. From argument diagrams to argumentation mining in texts: A survey. *International Journal of Cognitive Informatics and Natural Intelligence (IJCINI)*, 7(1):1–31.
- Peldszus, A. and M. Stede. 2016. An annotated corpus of argumentative microtexts. *1st European Conference on Argumentation*.
- Peng, Y., J. Tian, B. Yan, D. Berrebbi, X. Chang, X. Li, J. Shi, S. Arora, W. Chen, R. Sharma, W. Zhang, Y. Sudo, M. Shakee, J. weon Jung, S. Maiti, and S. Watanabe. 2023. Reproducing whisper-style training using an open-source toolkit and publicly available data. *ASRU*.
- Penn, G. and P. Kiparsky. 2012. [On Pāṇini and the generative capacity of contextualized replacement systems](#). *COLING*.
- Pennebaker, J. W., R. J. Booth, and M. E. Francis. 2007. *Linguistic Inquiry and Word Count: LIWC 2007*. Austin, TX.
- Pennington, J., R. Socher, and C. D. Manning. 2014. [GloVe: Global vectors for word representation](#). *EMNLP*.

- Percival, W. K. 1976. On the historical source of immediate constituent analysis. In J. D. McCawley, ed., *Syntax and Semantics Volume 7, Notes from the Linguistic Underground*, 229–242. Academic Press.
- Peters, M., M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, and L. Zettlemoyer. 2018. [Deep contextualized word representations](#). *NAACL HLT*.
- Peterson, G. E. and H. L. Barney. 1952. Control methods used in a study of the vowels. *JASA*, 24:175–184.
- Peterson, G. E., W. S.-Y. Wang, and E. Sivertsen. 1958. Segmentation techniques in speech synthesis. *JASA*, 30(8):739–742.
- Peterson, J. C., D. Chen, and T. L. Griffiths. 2020. Parallelograms revisited: Exploring the limitations of vector space models for simple analogies. *Cognition*, 205.
- Petroni, F., T. Rocktäschel, S. Riedel, P. Lewis, A. Bakhtin, Y. Wu, and A. Miller. 2019. [Language models as knowledge bases? EMNLP](#).
- Petrov, S., D. Das, and R. McDonald. 2012. [A universal part-of-speech tagset](#). *LREC*.
- Petrov, S. and R. McDonald. 2012. Overview of the 2012 shared task on parsing the web. *Notes of the First Workshop on Syntactic Analysis of Non-Canonical Language (SANCL)*, volume 59.
- Picard, R. W. 1995. Affective computing. Technical Report 321, MIT Media Lab Perceptual Computing Technical Report. Revised November 26, 1995.
- Pierce, J. R., J. B. Carroll, E. P. Hamp, D. G. Hays, C. F. Hockett, A. G. Oettinger, and A. J. Perlis. 1966. *Language and Machines: Computers in Translation and Linguistics*. ALPAC report. National Academy of Sciences, National Research Council, Washington, DC.
- Pilehvar, M. T. and J. Camacho-Collados. 2019. [WiC: the word-in-context dataset for evaluating context-sensitive meaning representations](#). *NAACL HLT*.
- Pitler, E., A. Louis, and A. Nenkova. 2009. [Automatic sense prediction for implicit discourse relations in text](#). *ACL IJCNLP*.
- Pitler, E. and A. Nenkova. 2009. [Using syntax to disambiguate explicit discourse connectives in text](#). *ACL IJCNLP*.
- Pitt, M. A., L. Dilley, K. Johnson, S. Kiesling, W. D. Raymond, E. Hume, and E. Fosler-Lussier. 2007. Buckeye corpus of conversational speech (2nd release). Department of Psychology, Ohio State University (Distributor).
- Pitt, M. A., K. Johnson, E. Hume, S. Kiesling, and W. D. Raymond. 2005. The buckeye corpus of conversational speech: Labeling conventions and a test of transcriber reliability. *Speech Communication*, 45:90–95.
- Plutchik, R. 1962. *The emotions: Facts, theories, and a new model*. Random House.
- Plutchik, R. 1980. A general psycho-evolutionary theory of emotion. In R. Plutchik and H. Kellerman, eds, *Emotion: Theory, Research, and Experience, Volume 1*, 3–33. Academic Press.
- Poesio, M., R. Stevenson, B. Di Eugenio, and J. Hitzeman. 2004. [Centering: A parametric theory and its instantiations](#). *Computational Linguistics*, 30(3):309–363.
- Poesio, M., R. Stuckardt, and Y. Versley. 2016. *Anaphora resolution: Algorithms, resources, and applications*. Springer.
- Poesio, M., P. Sturt, R. Artstein, and R. Filik. 2006. Underspecification and anaphora: Theoretical issues and preliminary evidence. *Discourse processes*, 42(2):157–175.
- Poesio, M. and R. Vieira. 1998. [A corpus-based investigation of definite description use](#). *Computational Linguistics*, 24(2):183–216.
- Polanyi, L. 1988. A formal model of the structure of discourse. *Journal of Pragmatics*, 12.
- Polanyi, L., C. Culy, M. van den Berg, G. L. Thione, and D. Ahn. 2004. [A rule based approach to discourse parsing](#). *Proceedings of SIGDIAL*.
- Pollard, C. and I. A. Sag. 1994. *Head-Driven Phrase Structure Grammar*. University of Chicago Press.
- Ponzetto, S. P. and M. Strube. 2006. [Exploiting semantic role labeling, WordNet and Wikipedia for coreference resolution](#). *HLT-NAACL*.
- Ponzetto, S. P. and M. Strube. 2007. Knowledge derived from Wikipedia for computing semantic relatedness. *JAIR*, 30:181–212.
- Popović, M. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). *Proceedings of the Tenth Workshop on Statistical Machine Translation*.
- Popp, D., R. A. Donovan, M. Crawford, K. L. Marsh, and M. Peele. 2003. Gender, race, and speech style stereotypes. *Sex Roles*, 48(7-8):317–325.
- Post, M. 2018. [A call for clarity in reporting BLEU scores](#). *WMT 2018*.
- Potts, C. 2011. On the negativity of negation. In N. Li and D. Lutz, eds, *Proceedings of Semantics and Linguistic Theory 20*, 636–659. CLC Publications, Ithaca, NY.
- Povey, D., A. Ghoshal, G. Boulianne, L. Burget, O. Glembek, N. Goel, M. Hannemann, P. Motlicek, Y. Qian, P. Schwarz, J. Silovsky, G. Stemmer, and K. Vesely. 2011. The Kaldi speech recognition toolkit. *ASRU*.
- Pradhan, S., E. H. Hovy, M. P. Marcus, M. Palmer, L. Ramshaw, and R. Weischedel. 2007a. OntoNotes: A unified relational semantic representation. *Proceedings of ICSC*.
- Pradhan, S., E. H. Hovy, M. P. Marcus, M. Palmer, L. A. Ramshaw, and R. M. Weischedel. 2007b. Ontonotes: a unified relational semantic representation. *Int. J. Semantic Computing*, 1(4):405–419.
- Pradhan, S., X. Luo, M. Recasens, E. H. Hovy, V. Ng, and M. Strube. 2014. [Scoring coreference partitions of predicted mentions: A reference implementation](#). *ACL*.
- Pradhan, S., A. Moschitti, N. Xue, H. T. Ng, A. Björkelund, O. Uryupina, Y. Zhang, and Z. Zhong. 2013. [Towards robust linguistic analysis using OntoNotes](#). *CoNLL*.
- Pradhan, S., A. Moschitti, N. Xue, O. Uryupina, and Y. Zhang. 2012a. [CoNLL-2012 shared task: Modeling multilingual unrestricted coreference in OntoNotes](#). *CoNLL*.
- Pradhan, S., A. Moschitti, N. Xue, O. Uryupina, and Y. Zhang. 2012b. [Conll-2012 shared task: Modeling multilingual unrestricted coreference in OntoNotes](#). *CoNLL*.
- Pradhan, S., L. Ramshaw, M. P. Marcus, M. Palmer, R. Weischedel, and N. Xue. 2011. [CoNLL-2011 shared task: Modeling unrestricted coreference in OntoNotes](#). *CoNLL*.
- Pradhan, S., L. Ramshaw, R. Weischedel, J. MacBride, and L. Micciulla. 2007c. Unrestricted coreference: Identifying entities and events in OntoNotes. *Proceedings of ICSC 2007*.
- Pradhan, S., W. Ward, K. Hacioglu, J. H. Martin, and D. Jurafsky. 2005. [Semantic role labeling using different syntactic views](#). *ACL*.
- Prakash, A., S. Aggarwal, J. J. Varghese, and J. J. Varghese. 2025. Writing without borders: AI and cross-cultural convergence in academic writing quality. *Humanities and Social Sciences Communications*, 12(1):1058.

- Prasad, R., N. Dinesh, A. Lee, E. Milt-sakaki, L. Robaldo, A. K. Joshi, and B. L. Webber. 2008. [The Penn Discourse TreeBank 2.0](#). *LREC*.
- Prasad, R., B. L. Webber, and A. Joshi. 2014. [Reflections on the Penn Discourse Treebank, comparable corpora, and complementary annotation](#). *Computational Linguistics*, 40(4):921–950.
- Prates, M. O. R., P. H. Avelar, and L. C. Lamb. 2019. Assessing gender bias in machine translation: a case study with Google Translate. *Neural Computing and Applications*, 32:6363–6381.
- Price, P. J., W. Fisher, J. Bernstein, and D. Pallet. 1988. The DARPA 1000-word resource management database for continuous speech recognition. *ICASSP*.
- Price, P. J., M. Ostendorf, S. Shattuck-Hufnagel, and C. Fong. 1991. The use of prosody in syntactic disambiguation. *JASA*, 90(6).
- Prince, E. 1981. Toward a taxonomy of given-new information. In P. Cole, ed., *Radical Pragmatics*, 223–255. Academic Press.
- Propp, V. 1968. *Morphology of the Folktale*, 2nd edition. University of Texas Press. Original Russian 1928. Translated by Laurence Scott.
- Pundak, G. and T. N. Sainath. 2016. Lower frame rate neural network acoustic models. *INTERSPEECH*.
- Pustejovsky, J. 1991. [The generative lexicon](#). *Computational Linguistics*, 17(4).
- Pustejovsky, J., P. Hanks, R. Saurí, A. See, R. Gaizauskas, A. Setzer, D. Radev, B. Sundheim, D. S. Day, L. Ferro, and M. Lazo. 2003. [The TIMEBANK corpus](#). *Proceedings of Corpus Linguistics 2003 Conference*. UCREL Technical Paper number 16.
- Pustejovsky, J., R. Ingria, R. Saurí, J. Castaño, J. Littman, R. Gaizauskas, A. Setzer, G. Katz, and I. Mani. 2005. *The Specification Language TimeML*, chapter 27. Oxford.
- Qin, L., Z. Zhang, and H. Zhao. 2016. [A stacking gated neural architecture for implicit discourse relation classification](#). *EMNLP*.
- Qin, L., Z. Zhang, H. Zhao, Z. Hu, and E. Xing. 2017. [Adversarial connective-exploiting networks for implicit discourse relation classification](#). *ACL*.
- Qwen Team. 2025. [Qwen3 technical report](#). ArXiv preprint.
- Radford, A., J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever. 2023. [Robust speech recognition via large-scale weak supervision](#). *ICML*.
- Radford, A., J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever. 2019. Language models are unsupervised multitask learners. OpenAI tech report.
- Rafailov, R., A. Sharma, E. Mitchell, S. Ermon, C. D. Manning, and C. Finn. 2023. [Direct preference optimization: Your language model is secretly a reward model](#). *NeurIPS*.
- Raffel, C., N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, and P. J. Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *JMLR*, 21(140):1–67.
- Raghunathan, K., H. Lee, S. Rangarajan, N. Chambers, M. Surdeanu, D. Jurafsky, and C. D. Manning. 2010. [A multi-pass sieve for coreference resolution](#). *EMNLP*.
- Rahman, A. and V. Ng. 2009. [Supervised models for coreference resolution](#). *EMNLP*.
- Rahman, A. and V. Ng. 2012. [Resolving complex cases of definite pronouns: the Winograd Schema challenge](#). *EMNLP*.
- Rajpurkar, P., R. Jia, and P. Liang. 2018. [Know what you don’t know: Unanswerable questions for SQuAD](#). *ACL*.
- Rajpurkar, P., J. Zhang, K. Lopyrev, and P. Liang. 2016. [SQuAD: 100,000+ questions for machine comprehension of text](#). *EMNLP*.
- Ram, O., Y. Levine, I. Dalmedigos, D. Muhlgay, A. Shashua, K. Leyton-Brown, and Y. Shoham. 2023. [In-context retrieval-augmented language models](#). ArXiv preprint.
- Ramachandran, P., B. Zoph, and Q. V. Le. 2017. [Searching for activation functions](#). ArXiv preprint.
- Ramshaw, L. A. and M. P. Marcus. 1995. [Text chunking using transformation-based learning](#). *Proceedings of the 3rd Annual Workshop on Very Large Corpora*.
- Rashkin, H., E. Bell, Y. Choi, and S. Volkova. 2017. [Multilingual connotation frames: A case study on social media for targeted sentiment analysis and forecast](#). *ACL*.
- Rashkin, H., S. Singh, and Y. Choi. 2016. [Connotation frames: A data-driven investigation](#). *ACL*.
- Ratinov, L. and D. Roth. 2012. [Learning-based multi-sieve coreference resolution with knowledge](#). *EMNLP*.
- Ratnaparkhi, A. 1996. [A maximum entropy part-of-speech tagger](#). *EMNLP*.
- Ratnaparkhi, A. 1997. [A linear observed time statistical parser based on maximum entropy models](#). *EMNLP*.
- Rawls, J. 2001. *Justice as fairness: A restatement*. Harvard University Press.
- Recasens, M. and E. H. Hovy. 2011. BLANC: Implementing the Rand index for coreference evaluation. *Natural Language Engineering*, 17(4):485–510.
- Recasens, M., E. H. Hovy, and M. A. Martí. 2011. Identity, non-identity, and near-identity: Addressing the complexity of coreference. *Lingua*, 121(6):1138–1152.
- Recasens, M. and M. A. Martí. 2010. AnCora-CO: Coreferentially annotated corpora for Spanish and Catalan. *Language Resources and Evaluation*, 44(4):315–345.
- Reed, C., R. Mochales Palau, G. Rowe, and M.-F. Moens. 2008. [Language resources for studying argument](#). *LREC*.
- Rehder, B., M. E. Schreiner, M. B. W. Wolfe, D. Laham, T. K. Landauer, and W. Kintsch. 1998. [Using Latent Semantic Analysis to assess knowledge: Some technical considerations](#). *Discourse Processes*, 25(2-3):337–354.
- Rei, R., C. Stewart, A. C. Farinha, and A. Lavie. 2020. [COMET: A neural framework for MT evaluation](#). *EMNLP*.
- Reichenbach, H. 1947. *Elements of Symbolic Logic*. Macmillan, New York.
- Reichman, R. 1985. *Getting Computers to Talk Like You and Me*. MIT Press.
- Reitmaier, T., D. K. Raju, O. Klejch, E. Wallington, N. Markl, J. Pearson, M. Jones, P. Bell, and S. Robinson. 2024. [Cultivating spoken language technologies for unwritten languages](#). *CHI*.
- Renals, S., T. Hain, and H. Bourlard. 2007. Recognition and understanding of meetings: The AMI and AMIDA projects. *ASRU*.
- Resnik, P. 1993. [Semantic classes and syntactic ambiguity](#). *HLT*.
- Resnik, P. 1996. Selectional constraints: An information-theoretic model and its computational realization. *Cognition*, 61:127–159.
- Riedel, S., L. Yao, and A. McCallum. 2010. Modeling relations and their mentions without labeled text. In *Machine Learning and Knowledge Discovery in Databases*, 148–163. Springer.

- Riedel, S., L. Yao, A. McCallum, and B. M. Marlin. 2013. [Relation extraction with matrix factorization and universal schemas](#). *NAACL HLT*.
- Riloff, E. 1993. Automatically constructing a dictionary for information extraction tasks. *AAAI*.
- Riloff, E. 1996. Automatically generating extraction patterns from untagged text. *AAAI*.
- Riloff, E. and R. Jones. 1999. [Learning dictionaries for information extraction by multi-level bootstrapping](#). *AAAI*.
- Riloff, E. and M. Schmelzenbach. 1998. An empirical approach to conceptual case frame acquisition. *Proceedings of the Sixth Workshop on Very Large Corpora*.
- Riloff, E. and J. Shepherd. 1997. [A corpus-based approach for building semantic lexicons](#). *EMNLP*.
- Riloff, E. and M. Thelen. 2000. [A rule-based question answering system for reading comprehension tests](#). *ANLP/NAACL workshop on reading comprehension tests*.
- Riloff, E. and J. Wiebe. 2003. [Learning extraction patterns for subjective expressions](#). *EMNLP*.
- Ritter, A., O. Etzioni, and Mausam. 2010. [A latent dirichlet allocation method for selectional preferences](#). *ACL*.
- Ritter, A., L. Zettlemoyer, Mausam, and O. Etzioni. 2013. [Modeling missing data in distant supervision for information extraction](#). *TACL*, 1:367–378.
- Roberts, A., C. Raffel, and N. Shazeer. 2020. [How much knowledge can you pack into the parameters of a language model?](#) *EMNLP*.
- Robertson, S., S. Walker, S. Jones, M. M. Hancock-Beaulieu, and M. Gatford. 1995. [Okapi at TREC-3](#). *Overview of the Third Text REtrieval Conference (TREC-3)*.
- Robins, R. H. 1967. *A Short History of Linguistics*. Indiana University Press, Bloomington.
- Robinson, T. and F. Fallside. 1991. A recurrent error propagation network speech recognition system. *Computer Speech & Language*, 5(3):259–274.
- Robinson, T., M. Hochberg, and S. Renals. 1996. The use of recurrent neural networks in continuous speech recognition. In C.-H. Lee, F. K. Soong, and K. K. Paliwal, eds, *Automatic speech and speaker recognition*, 233–258. Springer.
- Rogers, A., M. Gardner, and I. Augenstein. 2023. QA dataset explosion: A taxonomy of NLP resources for question answering and reading comprehension. *ACM Computing Surveys*, 55(10):1–45.
- Rohde, D. L. T., L. M. Gonnerman, and D. C. Plaut. 2006. An improved model of semantic similarity based on lexical co-occurrence. *CACM*, 8:627–633.
- Rooth, M., S. Riezler, D. Prescher, G. Carroll, and F. Beil. 1999. [Inducing a semantically annotated lexicon via EM-based clustering](#). *ACL*.
- Rosenblatt, F. 1958. The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological review*, 65(6):386–408.
- Rosenfeld, R. 1994. *Adaptive Statistical Language Modeling: A Maximum Entropy Approach*. Ph.D. thesis, Carnegie Mellon University.
- Rosenfeld, R. 1996. A maximum entropy approach to adaptive statistical language modeling. *Computer Speech and Language*, 10:187–228.
- Rosenthal, S. and K. McKeown. 2017. Detecting influencers in multiple online genres. *ACM Transactions on Internet Technology (TOIT)*, 17(2).
- Rothe, S., S. Ebert, and H. Schütze. 2016. [Ultradense Word Embeddings by Orthogonal Transformation](#). *NAACL HLT*.
- Rudinger, R., J. Naradowsky, B. Leonard, and B. Van Durme. 2018. [Gender bias in coreference resolution](#). *NAACL HLT*.
- Rumelhart, D. E., G. E. Hinton, and R. J. Williams. 1986. Learning internal representations by error propagation. In D. E. Rumelhart and J. L. McClelland, eds, *Parallel Distributed Processing*, volume 2, 318–362. MIT Press.
- Rumelhart, D. E. and J. L. McClelland. 1986a. On learning the past tense of English verbs. In D. E. Rumelhart and J. L. McClelland, eds, *Parallel Distributed Processing*, volume 2, 216–271. MIT Press.
- Rumelhart, D. E. and J. L. McClelland, eds. 1986b. *Parallel Distributed Processing*. MIT Press.
- Rumelhart, D. E. and A. A. Abrahamson. 1973. A model for analogical reasoning. *Cognitive Psychology*, 5(1):1–28.
- Rumelhart, D. E. and J. L. McClelland, eds. 1986c. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, volume 1: *Foundations*. MIT Press.
- Ruppenhofer, J., M. Ellsworth, M. R. L. Petrucci, C. R. Johnson, C. F. Baker, and J. Scheffczyk. 2016. *FrameNet II: Extended theory and practice*. Ruppenhofer, J., C. Sporleder, R. Morante, C. F. Baker, and M. Palmer. 2010. [Semeval-2010 task 10: Linking events and their participants in discourse](#). *5th International Workshop on Semantic Evaluation*.
- Russell, J. A. 1980. A circumplex model of affect. *Journal of personality and social psychology*, 39(6):1161–1178.
- Russell, S. and P. Norvig. 2002. *Artificial Intelligence: A Modern Approach*, 2nd edition. Prentice Hall.
- Rust, P., J. Pfeiffer, I. Vulić, S. Ruder, and I. Gurevych. 2021. [How good is your tokenizer? on the monolingual performance of multilingual language models](#). *ACL*.
- Rutherford, A. and N. Xue. 2015. [Improving the inference of implicit discourse relations via classifying explicit discourse connectives](#). *NAACL HLT*.
- Sachan, D. S., M. Lewis, D. Yogatama, L. Zettlemoyer, J. Pineau, and M. Zaheer. 2023. [Questions are all you need to train a dense passage retriever](#). *TACL*, 11:600–616.
- Sacks, H., E. A. Schegloff, and G. Jefferson. 1974. A simplest systematics for the organization of turn-taking for conversation. *Language*, 50(4):696–735.
- Sagae, K. 2009. [Analysis of discourse structure with syntactic dependencies and data-driven shift-reduce parsing](#). *IWPT-09*.
- Sagawa, S., P. W. Koh, T. B. Hashimoto, and P. Liang. 2020. Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization. *ICLR*.
- Sagisaka, Y. 1988. Speech synthesis by rule using an optimal selection of non-uniform synthesis units. *ICASSP*.
- Sagisaka, Y., N. Kaiki, N. Iwahashi, and K. Mimura. 1992. *Atr – v-talk* speech synthesis system. *ICSLP*.
- Sakoe, H. and S. Chiba. 1971. A dynamic programming approach to continuous speech recognition. *Proceedings of the Seventh International Congress on Acoustics*, volume 3. Akadémiai Kiadó.
- Sakoe, H. and S. Chiba. 1984. Dynamic programming algorithm optimization for spoken word recognition. *IEEE Transactions on ASSP*, ASSP-26(1):43–49.
- Salomaa, A. 1969. Probabilistic and weighted grammars. *Information and Control*, 15:529–544.

- Salton, G. 1971. *The SMART Retrieval System: Experiments in Automatic Document Processing*. Prentice Hall.
- Sampson, G. 1987. Alternative grammatical coding systems. In R. Gar-side, G. Leech, and G. Sampson, eds, *The Computational Analysis of English*, 165–183. Longman.
- Sankoff, D. and W. Labov. 1979. On the uses of variable rules. *Language in society*, 8(2-3):189–222.
- Sap, M., D. Card, S. Gabriel, Y. Choi, and N. A. Smith. 2019. [The risk of racial bias in hate speech detection](#). *ACL*.
- Sap, M., M. C. Prasettio, A. Holtzman, H. Rashkin, and Y. Choi. 2017. [Connotation frames of power and agency in modern films](#). *EMNLP*.
- Saurí, R., J. Littman, B. Knippen, R. Gaizauskas, A. Setzer, and J. Pustejovsky. 2006. [TimeML annotation guidelines version 1.2.1](#). Manuscript.
- Scha, R. and L. Polanyi. 1988. [An augmented context free grammar for discourse](#). *COLING*.
- Schank, R. C. and R. P. Abelson. 1975. Scripts, plans, and knowledge. *Proceedings of IJCAI-75*.
- Schank, R. C. and R. P. Abelson. 1977. *Scripts, Plans, Goals and Understanding*. Lawrence Erlbaum.
- Schegloff, E. A. 1968. Sequencing in conversational openings. *American Anthropologist*, 70:1075–1095.
- Schegloff, E. A. 1982. Discourse as an interactional achievement: Some uses of ‘uh huh’ and other things that come between sentences. In D. Tannen, ed., *Analyzing Discourse: Text and Talk*, 71–93. Georgetown University Press, Washington, D.C.
- Scherer, K. R. 2000. Psychological models of emotion. In J. C. Borod, ed., *The neuropsychology of emotion*, 137–162. Oxford.
- Schiebinger, L. 2013. Machine translation: Analyzing gender. <http://genderedinnovations.stanford.edu/case-studies/nlp.html#tabs-2>.
- Schiebinger, L. 2014. Scientific research must take gender into account. *Nature*, 507(7490):9.
- Schluter, N. 2018. [The word analogy testing caveat](#). *NAACL HLT*.
- Schmidt, C. W., V. Reddy, C. Tanner, and Y. Pinter. 2025. [Boundless byte pair encoding: Breaking the pre-tokenization barrier](#). *COLM*.
- Schone, P. and D. Jurafsky. 2000. [Knowledge-free induction of morphology using latent semantic analysis](#). *CoNLL*.
- Schone, P. and D. Jurafsky. 2001a. [Is knowledge-free induction of multi-word unit dictionary headwords a solved problem?](#) *EMNLP*.
- Schone, P. and D. Jurafsky. 2001b. [Knowledge-free induction of inflectional morphologies](#). *NAACL*.
- Schuster, M. and K. Nakajima. 2012. [Japanese and Korean voice search](#). *ICASSP*.
- Schuster, M. and K. K. Paliwal. 1997. Bidirectional recurrent neural networks. *IEEE Transactions on Signal Processing*, 45:2673–2681.
- Schütze, H. 1992a. Context space. *AAAI Fall Symposium on Probabilistic Approaches to Natural Language*.
- Schütze, H. 1992b. [Dimensions of meaning](#). *Proceedings of Supercomputing ’92*. IEEE Press.
- Schütze, H. 1997. *Ambiguity Resolution in Language Learning – Computational and Cognitive Models*. CSLI, Stanford, CA.
- Schütze, H., D. A. Hull, and J. Pedersen. 1995. [A comparison of classifiers and document representations for the routing problem](#). *SIGIR-95*.
- Schütze, H. and J. Pedersen. 1993. A vector model for syntagmatic and paradigmatic relatedness. *9th Annual Conference of the UW Centre for the New OED and Text Research*.
- Schütze, H. and Y. Singer. 1994. [Part-of-speech tagging using a variable memory Markov model](#). *ACL*.
- Schwartz, H. A., J. C. Eichstaedt, M. L. Kern, L. Dziurzynski, S. M. Ramones, M. Agrawal, A. Shah, M. Kosinski, D. Stillwell, M. E. P. Seligman, and L. H. Ungar. 2013. Personality, gender, and age in the language of social media: The open-vocabulary approach. *PloS one*, 8(9):e73791.
- Schwenk, H. 2007. [Continuous space language models](#). *Computer Speech & Language*, 21(3):492–518.
- Schwenk, H. 2018. [Filtering and mining parallel data in a joint multilingual space](#). *ACL*.
- Schwenk, H., D. Dechelotte, and J.-L. Gauvain. 2006. [Continuous space language models for statistical machine translation](#). *COLING/ACL*.
- Schwenk, H., G. Wenzek, S. Edunov, E. Grave, A. Joulin, and A. Fan. 2021. [CCMatrix: Mining billions of high-quality parallel sentences on the web](#). *ACL*.
- Séaghdha, D. O. 2010. [Latent variable models of selectional preference](#). *ACL*.
- Seddah, D., R. Tsarfaty, S. Kübler, M. Candito, J. D. Choi, R. Farkas, J. Foster, I. Goenaga, K. Gojenola, Y. Goldberg, S. Green, N. Habash, M. Kuhlmann, W. Maier, J. Nivre, A. Przyiórkowski, R. Roth, W. Seeker, Y. Versley, V. Vincze, M. Woliński, A. Wróblewska, and E. Villemonte de la Clérgerie. 2013. [Overview of the SPMRL 2013 shared task: cross-framework evaluation of parsing morphologically rich languages](#). *4th Workshop on Statistical Parsing of Morphologically-Rich Languages*.
- Sekine, S. and M. Collins. 1997. The evalb software. <http://cs.nyu.edu/cs/projects/proteus/evalb>.
- Sellam, T., D. Das, and A. Parikh. 2020. [BLEURT: Learning robust metrics for text generation](#). *ACL*.
- Seneff, S. and V. W. Zue. 1988. Transcription and alignment of the TIMIT database. *Proceedings of the Second Symposium on Advanced Man-Machine Interface through Spoken Language*.
- Sennrich, R., B. Haddow, and A. Birch. 2016. [Neural machine translation of rare words with subword units](#). *ACL*.
- Seo, M., A. Kembhavi, A. Farhadi, and H. Hajishirzi. 2017. Bidirectional attention flow for machine comprehension. *ICLR*.
- Shaikh, O., K. Gligorić, A. Khetan, M. Gerstgrasser, D. Yang, and D. Jurafsky. 2024. [Grounding gaps in language model generations](#). *NAACL HLT*.
- Shanahan, M. 2024. [Talking about large language models](#). *Communications of the ACM*, 67(2):68–79.
- Shannon, C. E. 1948. [A mathematical theory of communication](#). *Bell System Technical Journal*, 27(3):379–423. Continued in the following volume.
- Shannon, C. E. 1951. [Prediction and entropy of printed English](#). *Bell System Technical Journal*, 30:50–64.
- Sharma, M., M. Tong, T. Korbak, D. Duvenaud, A. Askell, S. R. Bowman, N. Cheng, E. Durmus, Z. Hatfield-Dodds, S. R. Johnston, S. Kravec, T. Maxwell, S. McCandlish, K. Ndousse, O. Rausch, N. Schiefer, D. Yan, M. Zhang, and E. Perez. 2024. [Towards understanding sycophancy in language models](#). *ICLR*.
- Shazeer, N. 2020. [Glu variants improve transformer](#). ArXiv preprint.
- Sheil, B. A. 1976. Observations on context free parsing. *SMIL: Statistical Methods in Linguistics*, 1:71–109.

- Shen, J. H. and A. Tamkin. 2026. [How AI impacts skill formation](#). ArXiv preprint.
- Sheng, E., K.-W. Chang, P. Natarajan, and N. Peng. 2019. [The woman worked as a babysitter: On biases in language generation](#). *EMNLP*.
- Shi, P. and J. Lin. 2019. [Simple BERT models for relation extraction and semantic role labeling](#). ArXiv.
- Shi, W., S. Min, M. Yasunaga, M. Seo, R. James, M. Lewis, L. Zettlemoyer, and W.-t. Yih. 2023. [REPLUG: Retrieval-augmented black-box language models](#). ArXiv preprint.
- Shoup, J. E. 1980. Phonological aspects of speech recognition. In W. A. Lea, ed., *Trends in Speech Recognition*, 125–138. Prentice Hall.
- Sidner, C. L. 1979. Towards a computational theory of definite anaphora comprehension in English discourse. Technical Report 537, MIT Artificial Intelligence Laboratory, Cambridge, MA.
- Sidner, C. L. 1983. Focusing in the comprehension of definite anaphora. In M. Brady and R. C. Berwick, eds, *Computational Models of Discourse*, 267–330. MIT Press.
- Silverman, K., M. E. Beckman, J. F. Pitrelli, M. Ostendorf, C. W. Wightman, P. J. Price, J. B. Pierrehumbert, and J. Hirschberg. 1992. ToBI: A standard for labelling English prosody. *ICSLP*.
- Simmons, R. F. 1965. Answering English questions by computer: A survey. *CACM*, 8(1):53–70.
- Simmons, R. F. 1973. Semantic networks: Their computation and use for understanding English sentences. In R. C. Schank and K. M. Colby, eds, *Computer Models of Thought and Language*, 61–113. W.H. Freeman & Co.
- Simmons, R. F., S. Klein, and K. McConlogie. 1964. Indexing and dependency logic for answering English questions. *American Documentation*, 15(3):196–204.
- Simons, G. F. and C. D. Fennig. 2018. *Ethnologue: Languages of the world*, 21st edition. SIL International.
- Singh, S., F. Vargus, D. D’souza, B. F. Karlsson, A. Mahendiran, W.-Y. Ko, H. Shandilya, J. Patel, D. Mataciunas, L. O’Mahony, M. Zhang, R. Hettiarachchi, J. Wilson, M. Machado, L. S. Moura, D. Krzemiński, H. Fadaei, I. Ergün, I. Okoh, A. Alaagib, O. Mudanayake, Z. Alyafeai, V. M. Chien, S. Ruder, S. Guthikonda, E. A. Alghamdi, S. Gehrmann, N. Muenighoff, M. Bartolo, J. Kreutzer, A. Üstün, M. Fadaee, and S. Hooker. 2024. [Aya dataset: An open-access collection for multilingual instruction tuning](#). ArXiv preprint.
- Sleator, D. and D. Temperley. 1993. [Parsing English with a link grammar](#). *IWPT-93*.
- Sloan, M. C. 2010. Aristotle’s Nicomachean Ethics as the original locus for the Septem Circumstantiae. *Classical Philology*, 105(3):236–251.
- Slobin, D. I. 1996. Two ways to travel. In M. Shibatani and S. A. Thompson, eds, *Grammatical Constructions: Their Form and Meaning*, 195–220. Clarendon Press.
- Smolensky, P. 1988. On the proper treatment of connectionism. *Behavioral and brain sciences*, 11(1):1–23.
- Smolensky, P. 1990. Tensor product variable binding and the representation of symbolic structures in connectionist systems. *Artificial intelligence*, 46(1-2):159–216.
- Snover, M., B. Dorr, R. Schwartz, L. Micciulla, and J. Makhoul. 2006. A study of translation edit rate with targeted human annotation. *AMTA-2006*.
- Snow, R., D. Jurafsky, and A. Y. Ng. 2005. Learning syntactic patterns for automatic hypernym discovery. *NeurIPS*.
- Socher, R., J. Bauer, C. D. Manning, and A. Y. Ng. 2013. [Parsing with compositional vector grammars](#). *ACL*.
- Socher, R., C. C.-Y. Lin, A. Y. Ng, and C. D. Manning. 2011. Parsing natural scenes and natural language with recursive neural networks. *ICML*.
- Soderland, S., D. Fisher, J. Aseltine, and W. G. Lehnert. 1995. CRYSTAL: Inducing a conceptual dictionary. *IJCAI-95*.
- Søgaard, A. 2010. [Simple semi-supervised training of part-of-speech taggers](#). *ACL*.
- Søgaard, A. and Y. Goldberg. 2016. [Deep multi-task learning with low level tasks supervised at lower layers](#). *ACL*.
- Søgaard, A., A. Johannsen, B. Plank, D. Hovy, and H. M. Alonso. 2014. [What’s in a p-value in NLP?](#) *CoNLL*.
- Solorio, T., E. Blair, S. Maharjan, S. Bethard, M. Diab, M. Ghoneim, A. Hawwari, F. AlGhamdi, J. Hirschberg, A. Chang, and P. Fung. 2014. [Overview for the first shared task on language identification in code-switched data](#). *Workshop on Computational Approaches to Code Switching*.
- Somasundaran, S., J. Burstein, and M. Chodorow. 2014. [Lexical chaining for measuring discourse coherence quality in test-taker essays](#). *COLING*.
- Soon, W. M., H. T. Ng, and D. C. Y. Lim. 2001. [A machine learning approach to coreference resolution of noun phrases](#). *Computational Linguistics*, 27(4):521–544.
- Soricut, R. and D. Marcu. 2003. [Sentence level discourse parsing using syntactic and lexical information](#). *HLT-NAACL*.
- Soricut, R. and D. Marcu. 2006. [Discourse generation using utility-trained coherence models](#). *COLING/ACL*.
- Sorokin, D. and I. Gurevych. 2018. [Mixing context granularities for improved entity linking on question answering data across entity categories](#). **SEM*.
- Sparck Jones, K. 1972. A statistical interpretation of term specificity and its application in retrieval. *Journal of Documentation*, 28(1):11–21.
- Sparck Jones, K. 1986. *Synonymy and Semantic Classification*. Edinburgh University Press, Edinburgh. Republication of 1964 PhD Thesis.
- Sporleder, C. and A. Lascarides. 2005. Exploiting linguistic cues to classify rhetorical relations. *RANLP-05*.
- Sporleder, C. and M. Lapata. 2005. [Discourse chunking and its application to sentence compression](#). *EMNLP*.
- Srivastava, N., G. E. Hinton, A. Krizhevsky, I. Sutskever, and R. R. Salakhutdinov. 2014. Dropout: a simple way to prevent neural networks from overfitting. *JMLR*, 15(1):1929–1958.
- Stab, C. and I. Gurevych. 2014a. [Annotating argument components and relations in persuasive essays](#). *COLING*.
- Stab, C. and I. Gurevych. 2014b. [Identifying argumentative discourse structures in persuasive essays](#). *EMNLP*.
- Stab, C. and I. Gurevych. 2017. [Parsing argumentation structures in persuasive essays](#). *Computational Linguistics*, 43(3):619–659.
- Stalnaker, R. C. 1978. Assertion. In P. Cole, ed., *Pragmatics: Syntax and Semantics Volume 9*, 315–332. Academic Press.
- Stamatatos, E. 2009. A survey of modern authorship attribution methods. *JASIST*, 60(3):538–556.
- Stanovsky, G., N. A. Smith, and L. Zettlemoyer. 2019. [Evaluating gender bias in machine translation](#). *ACL*.
- Stede, M. 2011. *Discourse processing*. Morgan & Claypool.

- Stede, M. and J. Schneider. 2018. *Argumentation Mining*. Morgan & Claypool.
- Stern, M., J. Andreas, and D. Klein. 2017. *A minimal span-based neural constituency parser*. *ACL*.
- Stevens, K. N. 1998. *Acoustic Phonetics*. MIT Press.
- Stevens, K. N. and A. S. House. 1955. Development of a quantitative description of vowel articulation. *JASA*, 27:484–493.
- Stevens, K. N. and A. S. House. 1961. An acoustical theory of vowel production and some of its implications. *Journal of Speech and Hearing Research*, 4:303–320.
- Stevens, K. N., S. Kasowski, and G. M. Fant. 1953. An electrical analog of the vocal tract. *JASA*, 25(4):734–742.
- Stevens, S. S. and J. Volkmann. 1940. The relation of pitch to frequency: A revised scale. *The American Journal of Psychology*, 53(3):329–353.
- Stevens, S. S., J. Volkmann, and E. B. Newman. 1937. A scale for the measurement of the psychological magnitude pitch. *JASA*, 8:185–190.
- Stolcke, A. 1998. Entropy-based pruning of backoff language models. *Proc. DARPA Broadcast News Transcription and Understanding Workshop*.
- Stolcke, A. 2002. SRILM – an extensible language modeling toolkit. *ICSLP*.
- Stolcke, A., Y. Konig, and M. Weintraub. 1997. *Explicit word error minimization in N-best list rescoring*. *EUROSPEECH*, volume 1.
- Stolz, W. S., P. H. Tannenbaum, and F. V. Carstensen. 1965. A stochastic approach to the grammatical coding of English. *CACM*, 8(6):399–405.
- Stone, P., D. Dunphy, M. Smith, and D. Ogilvie. 1966. *The General Inquirer: A Computer Approach to Content Analysis*. MIT Press.
- Strötgen, J. and M. Gertz. 2013. *Multilingual and cross-domain temporal tagging*. *Language Resources and Evaluation*, 47(2):269–298.
- Strube, M. and U. Hahn. 1996. *Functional centering*. *ACL*.
- Strubell, E., A. Ganesh, and A. McCallum. 2019. *Energy and policy considerations for deep learning in NLP*. *ACL*.
- Su, J., M. Ahmed, Y. Lu, S. Pan, B. Wen, and Y. Liu. 2024. *Roformer: Enhanced transformer with Rotary Position Embedding*. *Neurocomputing*, 568:127063.
- Su, Y., H. Sun, B. Sadler, M. Srivatsa, I. Gür, Z. Yan, and X. Yan. 2016. *On generating characteristic-rich question sets for QA evaluation*. *EMNLP*.
- Subba, R. and B. Di Eugenio. 2009. *An effective discourse parser that uses rich linguistic information*. *NAACL HLT*.
- Sukhbaatar, S., A. Szlam, J. Weston, and R. Fergus. 2015. End-to-end memory networks. *NeurIPS*.
- Sundheim, B., ed. 1991. *Proceedings of MUC-3*.
- Sundheim, B., ed. 1992. *Proceedings of MUC-4*.
- Sundheim, B., ed. 1993. *Proceedings of MUC-5*. Baltimore, MD.
- Sundheim, B., ed. 1995. *Proceedings of MUC-6*.
- Surdeanu, M. 2013. Overview of the TAC2013 Knowledge Base Population evaluation: English slot filling and temporal slot filling. *TAC-13*.
- Surdeanu, M., S. Harabagiu, J. Williams, and P. Aarseth. 2003. *Using predicate-argument structures for information extraction*. *ACL*.
- Surdeanu, M., T. Hicks, and M. A. Valenzuela-Escarcega. 2015. *Two practical rhetorical structure theory parsers*. *NAACL HLT*.
- Surdeanu, M., R. Johansson, A. Meyers, L. Márquez, and J. Nivre. 2008. *The CoNLL 2008 shared task on joint parsing of syntactic and semantic dependencies*. *CoNLL*.
- Sutskever, I., O. Vinyals, and Q. V. Le. 2014. Sequence to sequence learning with neural networks. *NeurIPS*.
- Sutton, R. S. and A. G. Barto. 1998. *Reinforcement Learning: An Introduction*. MIT Press.
- Suzgun, M., L. Melas-Kyriazi, and D. Jurafsky. 2023a. *Follow the wisdom of the crowd: Effective text generation via minimum Bayes risk decoding*. *Findings of ACL 2023*.
- Suzgun, M., N. Scales, N. Schärli, S. Gehrmann, Y. Tay, H. W. Chung, A. Chowdhery, Q. Le, E. Chi, D. Zhou, and J. Wei. 2023b. *Challenging BIG-bench tasks and whether chain-of-thought can solve them*. *ACL Findings*.
- Swanson, K., W. Wu, N. L. Bulaong, J. E. Pak, and J. Zou. 2025. *The virtual lab of AI agents designs new SARS-CoV-2 nanobodies*. *Nature*, 646:716–723.
- Sweet, H. 1877. *A Handbook of Phonetics*. Clarendon Press.
- Swier, R. and S. Stevenson. 2004. *Unsupervised semantic role labelling*. *EMNLP*.
- Switzer, P. 1965. Vector images in document retrieval. *Statistical Association Methods For Mechanized Documentation. Symposium Proceedings. Washington, D.C., USA, March 17, 1964*. <https://nvlpubs.nist.gov/nistpubs/Legacy/MP/nbsmiscellaneouspub269.pdf>.
- Syrdal, A. K., C. W. Wightman, A. Conkie, Y. Stylianou, M. Beutnagel, J. Schroeter, V. Strom, and K.-S. Lee. 2000. Corpus-based techniques in the AT&T NEXTGEN synthesis system. *ICSLP*.
- Talmy, L. 1985. Lexicalization patterns: Semantic structure in lexical forms. In T. Shopen, ed., *Language Typology and Syntactic Description, Volume 3*. Cambridge University Press. Originally appeared as UC Berkeley Cognitive Science Program Report No. 30, 1980.
- Talmy, L. 1991. *Path to realization: A typology of event conflation*. *BLS-91*.
- Tan, C., V. Niculae, C. Danescu-Niculescu-Mizil, and L. Lee. 2016. Winning arguments: Interaction dynamics and persuasion strategies in good-faith online discussions. *WWW-16*.
- Tannen, D. 1979. What’s in a frame? Surface evidence for underlying expectations. In R. Freedle, ed., *New Directions in Discourse Processing*, 137–181. Ablex.
- Taori, R., I. Gulrajani, T. Zhang, Y. Dubois, X. Li, C. Guestrin, P. Liang, and T. B. Hashimoto. 2023. *Alpaca: A strong, replicable instruction-following model*. Stanford Center for Research on Foundation Models.
- Taylor, P. 2009. *Text-to-Speech Synthesis*. Cambridge University Press.
- Taylor, W. L. 1953. Cloze procedure: A new tool for measuring readability. *Journalism Quarterly*, 30:415–433.
- Tenney, I., D. Das, and E. Pavlick. 2019. *BERT rediscovered the classical NLP pipeline*. *ACL*.
- Teranishi, R. and N. Umeda. 1968. Use of pronouncing dictionary in speech synthesis experiments. *6th International Congress on Acoustics*.
- Tesnière, L. 1959. *Éléments de Syntaxe Structurale*. Librairie C. Klincksieck, Paris.
- Tetreault, J. R. 2001. *A corpus-based evaluation of centering and pronoun resolution*. *Computational Linguistics*, 27(4):507–520.
- Teufel, S., J. Carletta, and M. Moens. 1999. *An annotation scheme for discourse-level argumentation in research articles*. *EACL*.

- Teufel, S., A. Siddharthan, and C. Batchelor. 2009. [Towards domain-independent argumentative zoning: Evidence from chemistry and computational linguistics](#). *EMNLP*.
- Theide, S. M. and M. P. Harper. 1999. [A second-order hidden Markov model for part-of-speech tagging](#). *ACL*.
- Thompson, B. and P. Koehn. 2019. [Vecalign: Improved sentence alignment in linear time and space](#). *EMNLP*.
- Thompson, K. 1968. [Regular expression search algorithm](#). *CACM*, 11(6):419–422.
- Tian, Y., V. Kulkarni, B. Perozzi, and S. Skiena. 2016. On the convergent properties of word embedding methods. ArXiv preprint arXiv:1605.03956.
- Tibshirani, R. J. 1996. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society. Series B (Methodological)*, 58(1):267–288.
- Timkey, W. and M. van Schijndel. 2021. [All bark and no bite: Rogue dimensions in transformer language models obscure representational quality](#). *EMNLP*.
- Titov, I. and E. Khoddam. 2014. [Unsupervised induction of semantic roles within a reconstruction-error minimization framework](#). *NAACL HLT*.
- Titov, I. and A. Klementiev. 2012. [A Bayesian approach to unsupervised semantic role induction](#). *EACL*.
- Tomkins, S. S. 1962. *Affect, imagery, consciousness: Vol. I. The positive affects*. Springer.
- Toutanova, K., D. Klein, C. D. Manning, and Y. Singer. 2003. [Feature-rich part-of-speech tagging with a cyclic dependency network](#). *HLT-NAACL*.
- Tria, F., V. Loreto, and V. D. Servadio. 2018. Zipf’s, heaps’ and taylor’s laws are determined by the expansion into the adjacent possible. *Entropy*, 20(10):752.
- Trichelair, P., A. Emami, J. C. K. Cheung, A. Trischler, K. Suleman, and F. Diaz. 2018. On the evaluation of common-sense reasoning in natural language understanding. *NeurIPS 2018 Workshop on Critiquing and Correcting Trends in Machine Learning*.
- Trnka, K., D. Yarrington, J. McCaw, K. F. McCoy, and C. Pennington. 2007. [The effects of word prediction on communication rate for AAC](#). *NAACL-HLT*.
- Turian, J. P., L. Shen, and I. D. D. Melamed. 2003. Evaluation of machine translation and its evaluation. *Proceedings of MT Summit IX*.
- Turian, J., L. Ratnoff, and Y. Bengio. 2010. [Word representations: a simple and general method for semi-supervised learning](#). *ACL*.
- Turney, P. D. 2002. [Thumbs up or thumbs down? Semantic orientation applied to unsupervised classification of reviews](#). *ACL*.
- Turney, P. D. and M. Littman. 2003. Measuring praise and criticism: Inference of semantic orientation from association. *ACM Transactions on Information Systems (TOIS)*, 21:315–346.
- Turney, P. D. and M. L. Littman. 2005. Corpus-based learning of analogies and semantic relations. *Machine Learning*, 60(1-3):251–278.
- Umeda, N. 1976. Linguistic rules for text-to-speech synthesis. *Proceedings of the IEEE*, 64(4):443–451.
- Umeda, N., E. Matui, T. Suzuki, and H. Omura. 1968. Synthesis of fairy tale using an analog vocal tract. *6th International Congress on Acoustics*.
- Uryupina, O., R. Artstein, A. Bristot, F. Cavicchio, F. Delogu, K. J. Rodriguez, and M. Poesio. 2020. Annotating a broad range of anaphoric phenomena, in a variety of genres: The ARRAU corpus. *Natural Language Engineering*, 26(1):1–34.
- Uszkoreit, J. 2017. [Transformer: A novel neural network architecture for language understanding](#). Google Research blog post, Thursday August 31, 2017.
- van Deemter, K. and R. Kibble. 2000. [On coreferring: coreference in MUC and related annotation schemes](#). *Computational Linguistics*, 26(4):629–637.
- Van Den Oord, A., O. Vinyals, and K. Kavukcuoglu. 2017. Neural discrete representation learning. *NeurIPS*.
- van der Maaten, L. and G. E. Hinton. 2008. [Visualizing high-dimensional data using t-SNE](#). *JMLR*, 9:2579–2605.
- van Rijsbergen, C. J. 1975. *Information Retrieval*. Butterworths.
- Vaswani, A., N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. 2017. [Attention is all you need](#). *NeurIPS*.
- Vauquois, B. 1968. A survey of formal grammars and algorithms for recognition and transformation in machine translation. *IFIP Congress 1968*.
- Velichko, V. M. and N. G. Zagoruyko. 1970. Automatic recognition of 200 words. *International Journal of Man-Machine Studies*, 2:223–234.
- Velikovich, L., S. Blair-Goldensohn, K. Hannan, and R. McDonald. 2010. [The viability of web-derived polarity lexicons](#). *NAACL HLT*.
- Vendler, Z. 1967. *Linguistics in Philosophy*. Cornell University Press.
- Verhagen, M., R. Gaizauskas, F. Schilder, M. Hepple, J. Moszkowicz, and J. Pustejovsky. 2009. The TempEval challenge: Identifying temporal relations in text. *Language Resources and Evaluation*, 43(2):161–179.
- Verhagen, M., I. Mani, R. Sauri, R. Knippen, S. B. Jang, J. Littman, A. Rumshisky, J. Phillips, and J. Pustejovsky. 2005. [Automating temporal annotation with TARSQI](#). *ACL*.
- Versley, Y. 2008. Vagueness and referential ambiguity in a large-scale annotated corpus. *Research on Language and Computation*, 6(3-4):333–353.
- Vieira, R. and M. Poesio. 2000. [An empirically based system for processing definite descriptions](#). *Computational Linguistics*, 26(4):539–593.
- Vilain, M., J. D. Burger, J. Aberdeen, D. Connolly, and L. Hirschman. 1995. [A model-theoretic coreference scoring scheme](#). *MUC-6*.
- Vintsyuk, T. K. 1968. Speech discrimination by dynamic programming. *Cybernetics*, 4(1):52–57. Original Russian: Kibernetika 4(1):81–88. 1968.
- Vinyals, O., L. Kaiser, T. Koo, S. Petrov, I. Sutskever, and G. Hinton. 2015. Grammar as a foreign language. *NeurIPS*.
- Voorhees, E. M. and D. K. Harman. 2005. *TREC: Experiment and Evaluation in Information Retrieval*. MIT Press.
- Voutilainen, A. 1999. Handcrafted rules. In H. van Halteren, ed., *Syntactic Wordclass Tagging*, 217–246. Kluwer.
- Vrandečić, D. and M. Krötzsch. 2014. Wikidata: a free collaborative knowledge base. *CACM*, 57(10):78–85.
- Wagner, R. A. and M. J. Fischer. 1974. [The string-to-string correction problem](#). *Journal of the ACM*, 21:168–173.
- Waibel, A., T. Hanazawa, G. Hinton, K. Shikano, and K. J. Lang. 1989. Phoneme recognition using time-delay neural networks. *IEEE Transactions on ASSP*, 37(3):328–339.

- Walker, M. A., M. Iida, and S. Cote. 1994. [Japanese discourse and the process of centering](#). *Computational Linguistics*, 20(2):193–232.
- Walker, M. A., A. K. Joshi, and E. Prince, eds. 1998. *Centering in Discourse*. Oxford University Press.
- Walker, M. A., E. Maier, J. Allen, J. Carletta, S. Condon, G. Flammia, J. Hirschberg, S. Isard, M. Ishizaki, L. Levin, S. Luperfoy, D. R. Traum, and S. Whittaker. 1996. Penn multiparty standard coding scheme: Draft annotation manual. www.cis.upenn.edu/~ircs/discourse-tagging/newcoding.html.
- Wang, A., A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman. 2018a. Glue: A multi-task benchmark and analysis platform for natural language understanding. *ICLR*.
- Wang, S. and C. D. Manning. 2012. [Baselines and bigrams: Simple, good sentiment and topic classification](#). *ACL*.
- Wang, W. and B. Chang. 2016. [Graph-based dependency parsing with bidirectional LSTM](#). *ACL*.
- Wang, Y., S. Li, and J. Yang. 2018b. [Toward fast and accurate neural discourse segmentation](#). *EMNLP*.
- Wang, Y., S. Mishra, P. Alipoormolabashi, Y. Kordi, A. Mirzaei, A. Naik, A. Ashok, A. S. Dhanasekaran, A. Arunkumar, D. Stap, E. Pathak, G. Karamanolakis, H. Lai, I. Purohit, I. Mondal, J. Anderson, K. Kuznia, K. Doshi, K. K. Pal, M. Patel, M. Moradshahi, M. Parmar, M. Purohit, N. Varshney, P. R. Kaza, P. Verma, R. S. Puri, R. Karia, S. Doshi, S. K. Sampat, S. Mishra, S. Reddy A., S. Patro, T. Dixit, and X. Shen. 2022. SuperNaturalInstructions: Generalization via declarative instructions on 1600+ NLP tasks. *EMNLP*.
- Wang, Z., Y. Dong, J. Zeng, V. Adams, M. N. Sreedhar, D. Egert, O. Delalleau, J. Scowcroft, N. Kant, A. Swope, and O. Kuchaiev. 2024. HelpSteer: Multi-attribute helpfulness dataset for SteerLM. *NAACL HLT*.
- Ward, N. and W. Tsukahara. 2000. Prosodic features which cue back-channel feedback in English and Japanese. *Journal of Pragmatics*, 32:1177–1207.
- Watanabe, S., T. Hori, S. Karita, T. Hayashi, J. Nishitoba, Y. Unno, N. E. Y. Soplin, J. Heymann, M. Wiesner, N. Chen, A. Renduchintala, and T. Ochiai. 2018. ESPnet: End-to-end speech processing toolkit. *INTERSPEECH*.
- Weaver, W. 1949/1955. Translation. In W. N. Locke and A. D. Boothe, eds, *Machine Translation of Languages*, 15–23. MIT Press. Reprinted from a memorandum written by Weaver in 1949.
- Webber, B. L. 1978. *A Formal Approach to Discourse Anaphora*. Ph.D. thesis, Harvard University.
- Webber, B. L. 1983. So what can we talk about now? In M. Brady and R. C. Berwick, eds, *Computational Models of Discourse*, 331–371. The MIT Press.
- Webber, B. L. 1991. Structure and tension in the interpretation of discourse deixis. *Language and Cognitive Processes*, 6(2):107–135.
- Webber, B. L. and B. Baldwin. 1992. [Accommodating context change](#). *ACL*.
- Webber, B. L., M. Egg, and V. Kordoni. 2012. Discourse structure and language technology. *Natural Language Engineering*, 18(4):437–490.
- Webber, B. L. 1988. [Discourse deixis: Reference to discourse segments](#). *ACL*.
- Webson, A. and E. Pavlick. 2022. Do prompt-based models really understand the meaning of their prompts? *NAACL HLT*.
- Webster, K., M. Recasens, V. Axelrod, and J. Baldridge. 2018. [Mind the GAP: A balanced corpus of gendered ambiguous pronouns](#). *TACL*, 6:605–617.
- Wei, J., X. Wang, D. Schuurmans, M. Bosma, F. Xia, E. Chi, Q. V. Le, D. Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *NeurIPS*, volume 35.
- Weischedel, R., M. Meteor, R. Schwartz, L. A. Ramshaw, and J. Palmucci. 1993. [Coping with ambiguity and unknown words through probabilistic models](#). *Computational Linguistics*, 19(2):359–382.
- Weizenbaum, J. 1966. ELIZA – A computer program for the study of natural language communication between man and machine. *CACM*, 9(1):36–45.
- Weizenbaum, J. 1976. *Computer Power and Human Reason: From Judgment to Calculation*. W.H. Freeman & Co.
- Wells, J. C. 1982. *Accents of English*. Cambridge University Press.
- Werbos, P. 1974. *Beyond regression: new tools for prediction and analysis in the behavioral sciences*. Ph.D. thesis, Harvard University.
- Werbos, P. J. 1990. Backpropagation through time: what it does and how to do it. *Proceedings of the IEEE*, 78(10):1550–1560.
- Weston, J., S. Chopra, and A. Bordes. 2015. [Memory networks](#). *ICLR 2015*.
- Widrow, B. and M. E. Hoff. 1960. Adaptive switching circuits. *IRE WESCON Convention Record*, volume 4.
- Wiebe, J. 1994. [Tracking point of view in narrative](#). *Computational Linguistics*, 20(2):233–287.
- Wiebe, J. 2000. Learning subjective adjectives from corpora. *AAAI*.
- Wiebe, J., R. F. Bruce, and T. P. O’Hara. 1999. [Development and use of a gold-standard data set for subjectivity classifications](#). *ACL*.
- Wierzbicka, A. 1992. *Semantics, Culture, and Cognition: University Human Concepts in Culture-Specific Configurations*. Oxford University Press.
- Wierzbicka, A. 1996. *Semantics: Primes and Universals*. Oxford University Press.
- Wilks, Y. 1973. An artificial intelligence approach to machine translation. In R. C. Schank and K. M. Colby, eds, *Computer Models of Thought and Language*, 114–151. W.H. Freeman.
- Wilks, Y. 1975a. Preference semantics. In E. L. Keenan, ed., *The Formal Semantics of Natural Language*, 329–350. Cambridge Univ. Press.
- Wilks, Y. 1975b. A preferential, pattern-seeking, semantics for natural language inference. *Artificial Intelligence*, 6(1):53–74.
- Williams, A., N. Nangia, and S. Bowman. 2018. [A broad-coverage challenge corpus for sentence understanding through inference](#). *NAACL HLT*.
- Wilson, T., J. Wiebe, and P. Hoffmann. 2005. [Recognizing contextual polarity in phrase-level sentiment analysis](#). *EMNLP*.
- Winograd, T. 1972. *Understanding Natural Language*. Academic Press.
- Winston, P. H. 1977. *Artificial Intelligence*. Addison Wesley.
- Wiseman, S., A. M. Rush, and S. M. Shieber. 2016. [Learning global features for coreference resolution](#). *NAACL HLT*.
- Wiseman, S., A. M. Rush, S. M. Shieber, and J. Weston. 2015. [Learning anaphoricity and antecedent ranking features for coreference resolution](#). *ACL*.
- Witten, I. H. and T. C. Bell. 1991. The zero-frequency problem: Estimating the probabilities of novel

- events in adaptive text compression. *IEEE Transactions on Information Theory*, 37(4):1085–1094.
- Witten, I. H. and E. Frank. 2005. *Data Mining: Practical Machine Learning Tools and Techniques*, 2nd edition. Morgan Kaufmann.
- Wittgenstein, L. 1953. *Philosophical Investigations. (Translated by Anscombe, G.E.M.)*. Blackwell.
- Wolf, F. and E. Gibson. 2005. [Representing discourse coherence: A corpus-based analysis](#). *Computational Linguistics*, 31(2):249–287.
- Woods, W. A. 1978. Semantics and quantification in natural language question answering. In M. Yovits, ed., *Advances in Computers*, 2–64. Academic.
- Woods, W. A., R. M. Kaplan, and B. L. Nash-Webber. 1972. The lunar sciences natural language information system: Final report. Technical Report 2378, BBN.
- Woodsend, K. and M. Lapata. 2015. [Distributed representations for unsupervised semantic role labeling](#). *EMNLP*.
- Wu, D. 1996. [A polynomial-time algorithm for statistical machine translation](#). *ACL*.
- Wu, F. and D. S. Weld. 2007. Autonomously semantifying Wikipedia. *CIKM-07*.
- Wu, F. and D. S. Weld. 2010. [Open information extraction using Wikipedia](#). *ACL*.
- Wu, L., F. Petroni, M. Josifoski, S. Riedel, and L. Zettlemoyer. 2020. [Scalable zero-shot entity linking with dense entity retrieval](#). *EMNLP*.
- Wu, S. and M. Dredze. 2019. [Beto, Bentz, Becas: The surprising cross-lingual effectiveness of BERT](#). *EMNLP*.
- Wu, Y., M. Schuster, Z. Chen, Q. V. Le, M. Norouzi, W. Macherey, M. Krikun, Y. Cao, Q. Gao, K. Macherey, J. Klingner, A. Shah, M. Johnson, X. Liu, Ł. Kaiser, S. Gouws, Y. Kato, T. Kudo, H. Kazawa, K. Stevens, G. Kurian, N. Patil, W. Wang, C. Young, J. Smith, J. Riesa, A. Rudnick, O. Vinyals, G. S. Corrado, M. Hughes, and J. Dean. 2016. Google’s neural machine translation system: Bridging the gap between human and machine translation. ArXiv preprint arXiv:1609.08144.
- Wundt, W. 1900. *Völkerpsychologie: eine Untersuchung der Entwicklungsgesetze von Sprache, Mythos, und Sitte*. W. Engelmann, Leipzig. Band II: Die Sprache, Zweiter Teil.
- Xiao, Y. 2025. [The impact of ai-driven speech recognition on EFL listening comprehension, flow experience, and anxiety: a randomized controlled trial](#). *Humanities and Social Sciences Communications*, 12(425).
- Xu, A., E. Pathak, E. Wallace, S. Gururangan, M. Sap, and D. Klein. 2021. [Detoxifying language models risks marginalizing minority voices](#). *NAACL HLT*.
- Xu, P., H. Saghir, J. S. Kang, T. Long, A. J. Bose, Y. Cao, and J. C. K. Cheung. 2019. [A cross-domain transferable neural coherence model](#). *ACL*.
- Xu, Y. 2005. Speech melody as articulatorily implemented communicative functions. *Speech communication*, 46(3-4):220–251.
- Xue, N., H. T. Ng, S. Pradhan, A. Rutherford, B. L. Webber, C. Wang, and H. Wang. 2016. [CoNLL 2016 shared task on multilingual shallow discourse parsing](#). *CoNLL-16 shared task*.
- Xue, N. and M. Palmer. 2004. [Calibrating features for semantic role labeling](#). *EMNLP*.
- Yamada, H. and Y. Matsumoto. 2003. [Statistical dependency analysis with support vector machines](#). *IWPT-03*.
- Yang, D., J. Chen, Z. Yang, D. Jurafsky, and E. H. Hovy. 2019. [Let’s make your request more persuasive: Modeling persuasive strategies via semi-supervised neural nets on crowd-funding platforms](#). *NAACL HLT*.
- Yang, X., G. Zhou, J. Su, and C. L. Tan. 2003. [Coreference resolution using competition learning approach](#). *ACL*.
- Yang, Y. and J. Pedersen. 1997. A comparative study on feature selection in text categorization. *ICML*.
- Yao, S., J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. 2023. [ReAct: Synergizing reasoning and acting in language models](#). *ICLR*.
- Yih, W.-t., M. Richardson, C. Meek, M.-W. Chang, and J. Suh. 2016. [The value of semantic parse labeling for knowledge base question answering](#). *ACL*.
- Yngve, V. H. 1970. On getting a word in edgewise. *CLS-70*. University of Chicago.
- Young, S. J., M. Gašić, S. Keizer, F. Mairesse, J. Schatzmann, B. Thomson, and K. Yu. 2010. The Hidden Information State model: A practical framework for POMDP-based spoken dialogue management. *Computer Speech & Language*, 24(2):150–174.
- Younger, D. H. 1967. Recognition and parsing of context-free languages in time n^3 . *Information and Control*, 10:189–208.
- Yu, N., M. Zhang, and G. Fu. 2018. [Transition-based neural RST parsing with implicit syntax features](#). *COLING*.
- Yu, Y., Y. Zhu, Y. Liu, Y. Liu, S. Peng, M. Gong, and A. Zeldes. 2019. [GumDrop at the DISRPT2019 shared task: A model stacking approach to discourse unit segmentation and connective detection](#). *Workshop on Discourse Relation Parsing and Treebanking 2019*.
- Yuan, J., M. Liberman, and C. Cieri. 2006. [Towards an integrated understanding of speaking rate in conversation](#). *Interspeech*.
- Zapirain, B., E. Agirre, L. Màrquez, and M. Surdeanu. 2013. [Selectional preferences for semantic role classification](#). *Computational Linguistics*, 39(3):631–663.
- Zelle, J. M. and R. J. Mooney. 1996. Learning to parse database queries using inductive logic programming. *AAAI*.
- Zeman, D. 2008. [Reusable tagset conversion using tagset drivers](#). *LREC*.
- Zens, R. and H. Ney. 2007. [Efficient phrase-table representation for machine translation with applications to online MT and speech translation](#). *NAACL-HLT*.
- Zettlemoyer, L. and M. Collins. 2005. Learning to map sentences to logical form: Structured classification with probabilistic categorical grammars. *Uncertainty in Artificial Intelligence, UAI’05*.
- Zettlemoyer, L. and M. Collins. 2007. [Online learning of relaxed CCG grammars for parsing to logical form](#). *EMNLP/CoNLL*.
- Zhang, A. K., K. Klyman, Y. Mai, Y. Levine, Y. Zhang, R. Bommasani, and P. Liang. 2025. Language model developers should report train-test overlap. *ICML*.
- Zhang, B. and R. Sennrich. 2019. Root mean square layer normalization. *Advances in Neural Information Processing Systems*, volume 32.
- Zhang, R., C. N. dos Santos, M. Yasunaga, B. Xiang, and D. Radev. 2018. [Neural coreference resolution with deep biaffine attention by joint mention detection and mention clustering](#). *ACL*.
- Zhang, T., V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi. 2020. [BERTscore: Evaluating text generation with BERT](#). *ICLR 2020*.

- Zhang, Y., V. Zhong, D. Chen, G. Angeli, and C. D. Manning. 2017. [Position-aware attention and supervised data improve slot filling](#). *EMNLP*.
- Zhang, Y., D. Zhao, J. T. Hancock, R. Kraut, and D. Yang. 2026. [Interaction with ai companions and psychological well-being](#). *Nature Human Behaviour*.
- Zhao, H., W. Chen, C. Kit, and G. Zhou. 2009. [Multilingual dependency learning: A huge feature engineering method to semantic dependency parsing](#). *CoNLL*.
- Zhao, J., T. Wang, M. Yatskar, R. Cotterell, V. Ordonez, and K.-W. Chang. 2019. [Gender bias in contextualized word embeddings](#). *NAACL HLT*.
- Zhao, J., T. Wang, M. Yatskar, V. Ordonez, and K.-W. Chang. 2017. [Men also like shopping: Reducing gender bias amplification using corpus-level constraints](#). *EMNLP*.
- Zhao, J., T. Wang, M. Yatskar, V. Ordonez, and K.-W. Chang. 2018a. [Gender bias in coreference resolution: Evaluation and debiasing methods](#). *NAACL HLT*.
- Zhao, J., Y. Zhou, Z. Li, W. Wang, and K.-W. Chang. 2018b. [Learning gender-neutral word embeddings](#). *EMNLP*.
- Zheng, J., L. Vilnis, S. Singh, J. D. Choi, and A. McCallum. 2013. [Dynamic knowledge-base alignment for coreference resolution](#). *CoNLL*.
- Zhou, D., O. Bousquet, T. N. Lal, J. Weston, and B. Schölkopf. 2004a. Learning with local and global consistency. *NeurIPS*.
- Zhou, G., J. Su, J. Zhang, and M. Zhang. 2005. [Exploring various knowledge in relation extraction](#). *ACL*.
- Zhou, J., T. Lu, S. Mishra, S. Brahma, S. Basu, Y. Luan, D. Zhou, and L. Hou. 2023. [Instruction-following evaluation for large language models](#). *ArXiv*, abs/2311.07911.
- Zhou, J. and W. Xu. 2015a. [End-to-end learning of semantic role labeling using recurrent neural networks](#). *ACL*.
- Zhou, J. and W. Xu. 2015b. [End-to-end learning of semantic role labeling using recurrent neural networks](#). *ACL*.
- Zhou, K., K. Ethayarajh, D. Card, and D. Jurafsky. 2022. [Problems with cosine as a measure of embedding similarity for high frequency words](#). *ACL*.
- Zhou, K., J. Hwang, X. Ren, and M. Sap. 2024a. [Relying on the unreliable: The impact of language models' reluctance to express uncertainty](#). *ACL*.
- Zhou, K., J. D. Hwang, X. Ren, and M. Sap. 2024b. [Relying on the unreliable: The impact of language models' reluctance to express uncertainty](#). *ACL*.
- Zhou, L., M. Ticea, and E. H. Hovy. 2004b. [Multi-document biography summarization](#). *EMNLP*.
- Zhou, Y. and N. Xue. 2015. The Chinese Discourse TreeBank: a Chinese corpus annotated with discourse relations. *Language Resources and Evaluation*, 49(2):397–431.
- Zhu, X. and Z. Ghahramani. 2002. Learning from labeled and unlabeled data with label propagation. Technical Report CMU-CALD-02, CMU.
- Zhu, X., Z. Ghahramani, and J. Lafferty. 2003. Semi-supervised learning using gaussian fields and harmonic functions. *ICML*.
- Zhu, Y., R. Kiros, R. Zemel, R. Salakhutdinov, R. Urtasun, A. Torralba, and S. Fidler. 2015. Aligning books and movies: Towards story-like visual explanations by watching movies and reading books. *IEEE International Conference on Computer Vision*.
- Ziegler, D. M., N. Stiennon, J. Wu, T. B. Brown, A. Radford, D. Amodei, P. Christiano, and G. Irving. 2019. Fine-tuning language models from human preferences. *ArXiv*, abs/1909.08593.
- Ziemski, M., M. Junczys-Dowmunt, and B. Poulliquen. 2016. [The United Nations parallel corpus v1.0](#). *LREC*.
- Zou, A., Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson. 2023. [Universal and transferable adversarial attacks on aligned language models](#). *ArXiv preprint*.